

# Rossmann's Uncertainty-Aware Sales Prediction

## Local

If you are running the notebook you should put the .csv in the dataset folder and run this cell

- ignore the next cell

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
# Load datasets
train = pd.read_csv('dataset/train.csv', parse_dates=['Date'],
low_memory=False)
store = pd.read_csv('dataset/store.csv')

# Merge train with store information
df = pd.merge(train, store, on='Store', how='left')

# Basic Info
print(df.info())
print(df.head())
```

A module that was compiled using NumPy 1.x cannot be run in NumPy 2.2.6 as it may crash. To support both 1.x and 2.x versions of NumPy, modules must be compiled with NumPy 2.0. Some module may need to rebuild instead e.g. with 'pybind11>=2.12'.

If you are a user of the module, the easiest solution will be to downgrade to 'numpy<2' or try to upgrade the affected module. We expect that some modules will need time to support NumPy 2.

```
Traceback (most recent call last): File "<frozen runpy>", line 198,
in _run_module_as_main
  File "<frozen runpy>", line 88, in _run_code
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel_launcher.py", line 18, in <module>
    app.launch_new_instance()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\traitlets\config\application.py", line 1075, in launch_instance
    app.start()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\kernelapp.py", line 739, in start
```

```
    self.io_loop.start()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\tornado\platform\asyncio.py", line 205, in start
    self.asyncio_loop.run_forever()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\asyncio\base_events.py", line 607, in run_forever
    self._run_once()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\asyncio\base_events.py", line 1922, in _run_once
    handle._run()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\asyncio\events.py", line 80, in _run
    self._context.run(self._callback, *self._args)
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\kernelbase.py", line 545, in dispatch_queue
    await self.process_one()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\kernelbase.py", line 534, in process_one
    await dispatch(*args)
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\kernelbase.py", line 437, in dispatch_shell
    await result
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\ipkernel.py", line 359, in execute_request
    await super().execute_request(stream, ident, parent)
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\kernelbase.py", line 778, in execute_request
    reply_content = await reply_content
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\ipkernel.py", line 446, in do_execute
    res = shell.run_cell()
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\ipykernel\zmqshell.py", line 549, in run_cell
    return super().run_cell(*args, **kwargs)
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\IPython\core\interactiveshell.py", line 3075, in run_cell
    result = self._run_cell(
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\IPython\core\interactiveshell.py", line 3130, in _run_cell
    result = runner(coro)
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\IPython\core\async_helpers.py", line 129, in _pseudo_sync_runner
    coro.send(None)
  File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\IPython\core\interactiveshell.py", line 3334, in run_cell_async
```

```
    has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\IPython\core\interactiveshell.py", line 3517, in
run_ast_nodes
    if await self.run_code(code, result, async_=asy):
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\IPython\core\interactiveshell.py", line 3577, in
run_code
    exec(code_obj, self.user_global_ns, self.user_ns)
File "C:\Users\francois\AppData\Local\Temp\ipykernel_2256\3720819326.py", line 2, in <module>
    import matplotlib.pyplot as plt
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\matplotlib\__init__.py", line 174, in <module>
    from . import _api, _version, cbook, _docstring, rcsetup
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\matplotlib\rcsetup.py", line 27, in <module>
    from matplotlib.colors import Colormap, is_color_like
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\matplotlib\colors.py", line 57, in <module>
    from matplotlib import _api, _cm, cbook, scale
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\matplotlib\scale.py", line 22, in <module>
    from matplotlib.ticker import (
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\matplotlib\ticker.py", line 143, in <module>
    from matplotlib import transforms as mtransforms
File "c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\matplotlib\transforms.py", line 49, in <module>
    from matplotlib._path import (
```

```
-----
-----
ImportError                                     Traceback (most recent call
last)
File c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\site-packages\numpy\core\_multiarray_umath.py:44, in
__getattr__(attr_name)
    39      # Also print the message (with traceback). This is
because old versions
    40      # of NumPy unfortunately set up the import to replace (and
hide) the
    41      # error. The traceback shouldn't be needed, but e.g.
pytest plugins
    42      # seem to swallow it and we should be failing anyway...
    43      sys.stderr.write(msg + tb_msg)
--> 44      raise ImportError(msg)
    46 ret = getattr(_multiarray_umath, attr_name, None)
    47 if ret is None:
```

```
ImportError:  
A module that was compiled using NumPy 1.x cannot be run in  
NumPy 2.2.6 as it may crash. To support both 1.x and 2.x  
versions of NumPy, modules must be compiled with NumPy 2.0.  
Some module may need to rebuild instead e.g. with 'pybind11>=2.12'.
```

```
If you are a user of the module, the easiest solution will be to  
downgrade to 'numpy<2' or try to upgrade the affected module.  
We expect that some modules will need time to support NumPy 2.
```

```
-----  
-----  
ImportError Traceback (most recent call  
last)  
Cell In[3], line 2  
      1 import pandas as pd  
----> 2 import matplotlib.pyplot as plt  
      3 import seaborn as sns  
      5 # Load datasets  
  
File c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\  
site-packages\matplotlib\__init__.py:174  
    170 from packaging.version import parse as parse_version  
    172 # cbook must import matplotlib only within function  
    173 # definitions, so it is safe to import from it here.  
--> 174 from . import _api, _version, cbook, _docstring, rcsetup  
    175 from matplotlib.cbook import sanitize_sequence  
    176 from matplotlib._api import MatplotlibDeprecationWarning  
  
File c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\  
site-packages\matplotlib\rcsetup.py:27  
    25 from matplotlib import _api, cbook  
    26 from matplotlib.cbook import ls_mapper  
---> 27 from matplotlib.colors import Colormap, is_color_like  
    28 from matplotlib._fontconfig_pattern import  
parse_fontconfig_pattern  
    29 from matplotlib._enums import JoinStyle, CapStyle  
  
File c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\  
site-packages\matplotlib\colors.py:57  
    55 import matplotlib as mpl  
    56 import numpy as np  
---> 57 from matplotlib import _api, _cm, cbook, scale  
    58 from ._color_data import BASE_COLORS, TABLEAU_COLORS,  
CSS4_COLORS, XKCD_COLORS  
    61 class _ColorMapping(dict):  
  
File c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\
```

```

site-packages\matplotlib\scale.py:22
    20 import matplotlib as mpl
    21 from matplotlib import _api, _docstring
--> 22 from matplotlib.ticker import (
    23     NullFormatter, ScalarFormatter, LogFormatterSciNotation,
LogitFormatter,
    24     NullLocator, LogLocator, AutoLocator, AutoMinorLocator,
    25     SymmetricalLogLocator, AsinhLocator, LogitLocator)
    26 from matplotlib.transforms import Transform, IdentityTransform
    29 class ScaleBase:

File c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\
site-packages\matplotlib\ticker.py:143
    141 import matplotlib as mpl
    142 from matplotlib import _api, cbook
--> 143 from matplotlib import transforms as mtransforms
    145 _log = logging.getLogger(__name__)
    147 __all__ = ('TickHelper', 'Formatter', 'FixedFormatter',
    148             'NullFormatter', 'FuncFormatter',
'FormatStrFormatter',
    149             'StrMethodFormatter', 'ScalarFormatter',
'LogFormatter',
    (...),
    155             'MultipleLocator', 'MaxNLocator',
'AutoMinorLocator',
    156             'SymmetricalLogLocator', 'AsinhLocator',
'LogitLocator')

File c:\Users\francois\AppData\Local\Programs\Python\Python311\Lib\
site-packages\matplotlib\transforms.py:49
    46 from numpy.linalg import inv
    48 from matplotlib import _api
--> 49 from matplotlib._path import (
    50     affine_transform, count_bboxes_overlapping_bbox,
update_path_extents)
    51 from .path import Path
    53 DEBUG = False

ImportError: numpy.core.multiarray failed to import

```

## Kaggle

If you are using kaggle as your notebook environment:

1. add /kaggle/competitions/rossmann-store-sales to the input list
2. run this cell

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import numpy as np
base_path = "/kaggle/input/competitions/rossmann-store-sales/"

# Load datasets
train = pd.read_csv(f'{base_path}train.csv', parse_dates=['Date'],
low_memory=False)
store = pd.read_csv(f'{base_path}store.csv')
test = pd.read_csv(f'{base_path}test.csv')

df = pd.merge(train, store, on='Store', how='left')

print(df.info())
print(df.head(-1))

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1017209 entries, 0 to 1017208
Data columns (total 18 columns):
 #   Column           Non-Null Count   Dtype  
--- 
 0   Store            1017209 non-null    int64  
 1   DayOfWeek        1017209 non-null    int64  
 2   Date             1017209 non-null    datetime64[ns]
 3   Sales            1017209 non-null    int64  
 4   Customers        1017209 non-null    int64  
 5   Open              1017209 non-null    int64  
 6   Promo             1017209 non-null    int64  
 7   StateHoliday     1017209 non-null    object  
 8   SchoolHoliday    1017209 non-null    int64  
 9   StoreType         1017209 non-null    object  
 10  Assortment       1017209 non-null    object  
 11  CompetitionDistance 1014567 non-null    float64 
 12  CompetitionOpenSinceMonth 693861 non-null    float64 
 13  CompetitionOpenSinceYear 693861 non-null    float64 
 14  Promo2            1017209 non-null    int64  
 15  Promo2SinceWeek   509178 non-null    float64 
 16  Promo2SinceYear   509178 non-null    float64 
 17  PromoInterval     509178 non-null    object  
dtypes: datetime64[ns](1), float64(5), int64(8), object(4)
memory usage: 139.7+ MB
None
      Store  DayOfWeek      Date   Sales  Customers  Open  Promo \
0        1          5 2015-07-31    5263       555      1      1
1        2          5 2015-07-31    6064       625      1      1
2        3          5 2015-07-31    8314       821      1      1
3        4          5 2015-07-31   13995      1498      1      1
4        5          5 2015-07-31    4822       559      1      1

```

...	...	...	...	...	...	...	...
1017203	1110	2	2013-01-01	0	0	0	0
1017204	1111	2	2013-01-01	0	0	0	0
1017205	1112	2	2013-01-01	0	0	0	0
1017206	1113	2	2013-01-01	0	0	0	0
1017207	1114	2	2013-01-01	0	0	0	0
StateHoliday		SchoolHoliday	StoreType	Assortment			
CompetitionDistance		\					
0	0		1	c	a		
1270.0							
1	0		1	a	a		
570.0							
2	0		1	a	a		
14130.0							
3	0		1	c	c		
620.0							
4	0		1	a	a		
29910.0							
...	...	...	...	...	...	...	...
...							
1017203	a		1	c	c		
900.0							
1017204	a		1	a	a		
1900.0							
1017205	a		1	c	c		
1880.0							
1017206	a		1	a	c		
9260.0							
1017207	a		1	a	c		
870.0							
Promo2		CompetitionOpenSinceMonth	CompetitionOpenSinceYear				
\							
0		9.0		2008.0	0		
1		11.0		2007.0	1		
2		12.0		2006.0	1		
3		9.0		2009.0	0		
4		4.0		2015.0	0		
...		...		...	...	...	...
1017203		9.0		2010.0	0		
1017204		6.0		2014.0	1		

1017205	4.0	2006.0	0
1017206	NaN	NaN	0
1017207	NaN	NaN	0
0	Promo2SinceWeek	Promo2SinceYear	PromoInterval
1	NaN	NaN	NaN
2	13.0	2010.0	Jan,Apr,Jul,Oct
3	14.0	2011.0	Jan,Apr,Jul,Oct
4	NaN	NaN	NaN
...	...	...	...
1017203	NaN	NaN	NaN
1017204	31.0	2013.0	Jan,Apr,Jul,Oct
1017205	NaN	NaN	NaN
1017206	NaN	NaN	NaN
1017207	NaN	NaN	NaN
[1017208 rows x 18 columns]			

## 1. Individual Store Trend Analysis

To understand the underlying patterns in the Rossmann dataset, we visualized a sample of stores (\$ 1, 10, 100, 1000 \$). Retail sales data is "noisy" due to weekly cycles and holiday spikes.

Raw Sales (Gray): Displays the high volatility of daily transactions, including zero during Sunday.

7-Day Rolling Mean (Blue): We applied a transformation to smooth out the daily "zig-zags." This allows us to see the actual sales direction and seasonal strength without being distracted by day-of-week fluctuations.

As we can see from the plots:

- store 1000 has a higher average sales compared to others
- The sales data is missing from store 100 and 1000
- one can notice a spike increase before christmass 2014-01 and 2015-01
- for all of the stores the sales pattern is like a sin function and with my stimate the period is every 2 to 3 weeks

```
sample_stores = [1, 10, 100, 1000]
fig, axes = plt.subplots(len(sample_stores), 1, figsize=(15, 12),
sharex=True)

for i, store_id in enumerate(sample_stores):
    store_data = df[df['Store'] == store_id].sort_values('Date')

    axes[i].plot(store_data['Date'], store_data['Sales'], alpha=0.3,
label='Daily Sales', color='gray')
```

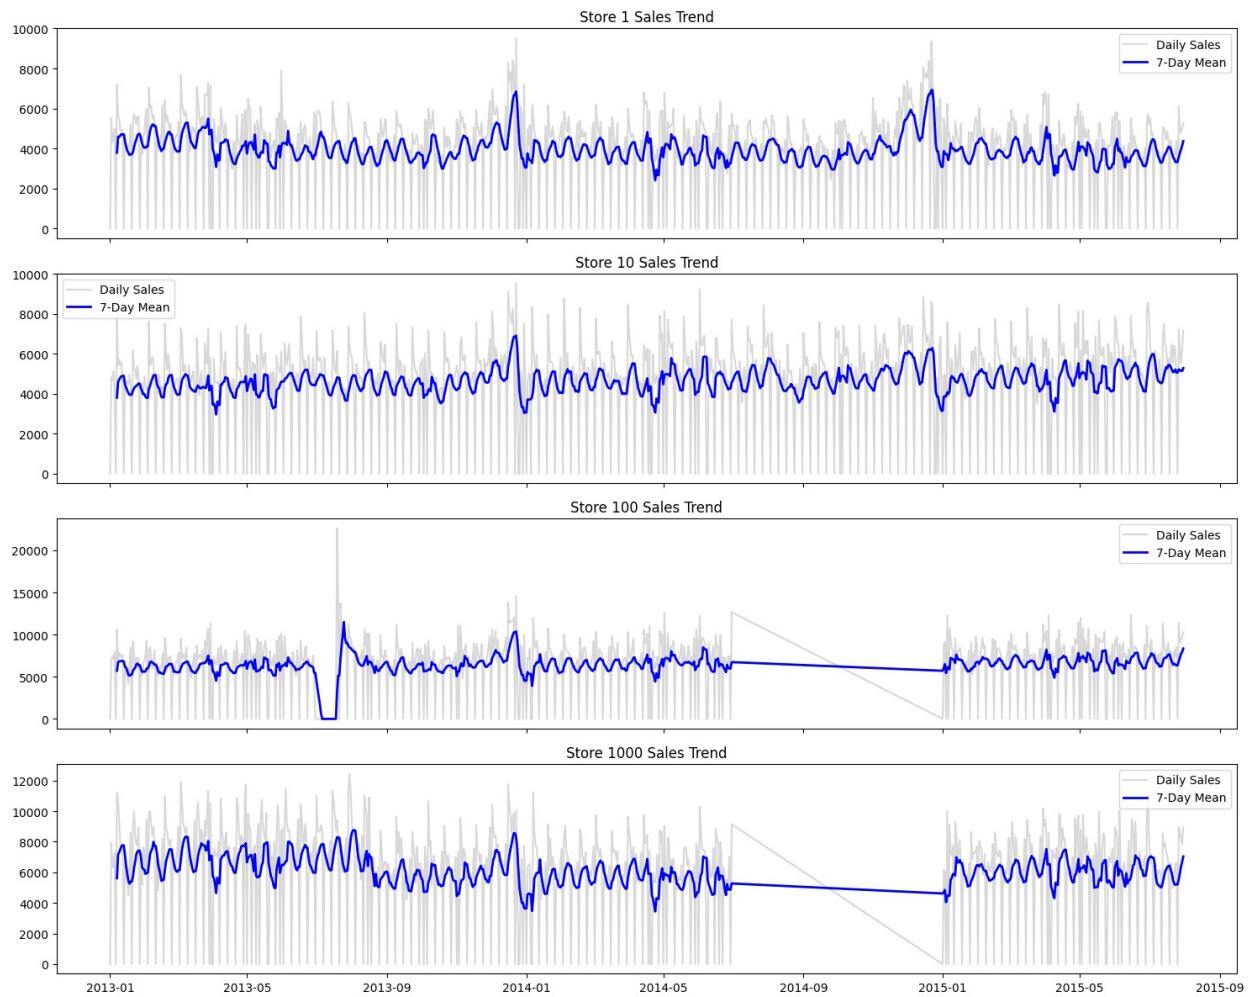
```

rolling_sales = store_data.set_index('Date')
['Sales'].rolling(window=7).mean()
axes[i].plot(rolling_sales.index, rolling_sales.values, label='7-Day Mean', color='blue', linewidth=2)

axes[i].set_title(f'Store {store_id} Sales Trend')
axes[i].legend()

plt.tight_layout()
plt.show()

```



## 2. Phase 6: Categorical Impact Analysis (External Drivers)\*\*

The first plot compares sales on days with and without a temporary promotion.

- There is a significant upward shift in both the median sales and the interquartile range (IQR) when `Promo=1`.

The second plot examines sales during Public holidays (a), Easter (b), and Christmas (c)(for example), compared to normal days (0).

- Most stores show near-zero sales during state holidays. However, the outliers indicate that a small subset of stores remains open.

The final plot compares sales during periods when local schools are closed.

- While the median sales increase slightly during school holidays, the impact is less dramatic than a standard Promo.

## Summary

Feature	Impact Strength	Primary Effect
<b>Promo</b>	<b>High</b>	Increases median sales and overall volume significantly.
<b>State Holiday</b>	<b>Critical</b>	Usually results in zero sales, but marks "exception" days.
<b>School Holiday</b>	<b>Moderate</b>	Provides a subtle lift; likely dependent on store-specific context.

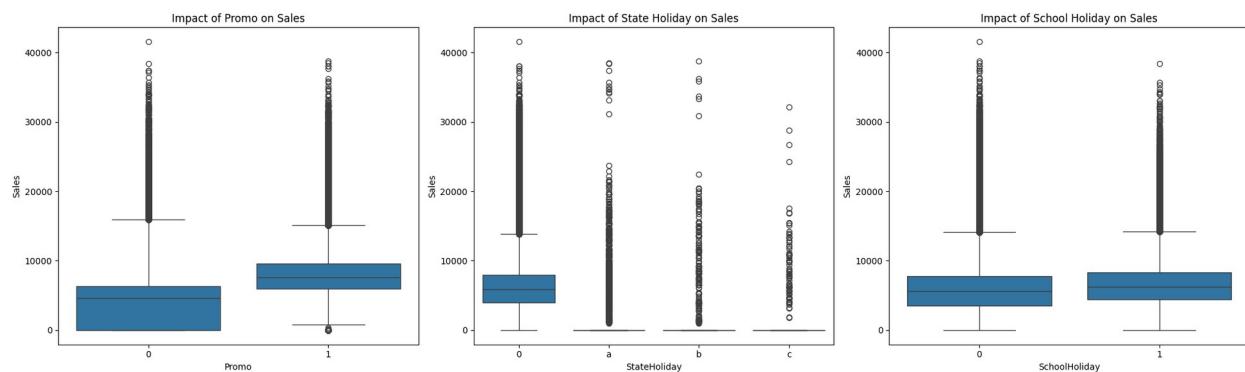
```
fig, axes = plt.subplots(1, 3, figsize=(20, 6))

sns.boxplot(data=df, x='Promo', y='Sales', ax=axes[0])
axes[0].set_title('Impact of Promo on Sales')

sns.boxplot(data=df, x='StateHoliday', y='Sales', ax=axes[1])
axes[1].set_title('Impact of State Holiday on Sales')

sns.boxplot(data=df, x='SchoolHoliday', y='Sales', ax=axes[2])
axes[2].set_title('Impact of School Holiday on Sales')

plt.tight_layout()
plt.show()
```



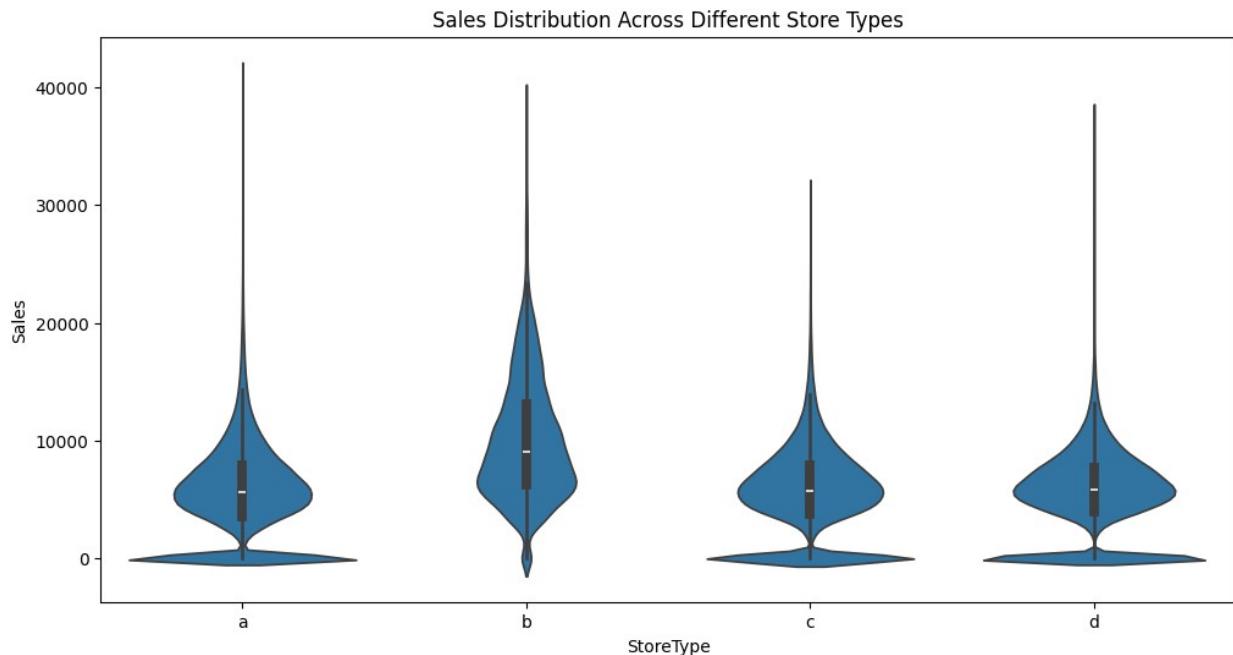
### 3. Sales Density by Store Type (StoreType)

We used a Violin Plot to visualize the distribution and density of sales across the four store categories (a, b, c, and d).

Store Type 'b' : Clearly stands out with a much higher median and a wider distribution at the top. These are likely "flagship" or high-traffic transit stores.

Store Types 'a', 'c', and 'd': These have a more normal distribution. Type 'd' shows a slightly higher and more "stretched" upper distribution.

```
plt.figure(figsize=(12, 6))
sns.violinplot(data=df, x='StoreType', y='Sales', order=['a', 'b',
 'c', 'd'])
plt.title('Sales Distribution Across Different Store Types')
plt.show()
```



### 4. Correlation Heatmap

I filtered for numeric types.

Sales vs. Customers : As expected, these have the highest correlation. But we cannot use "Customers" as an input for future predictions since we won't know the customer count in test faze.

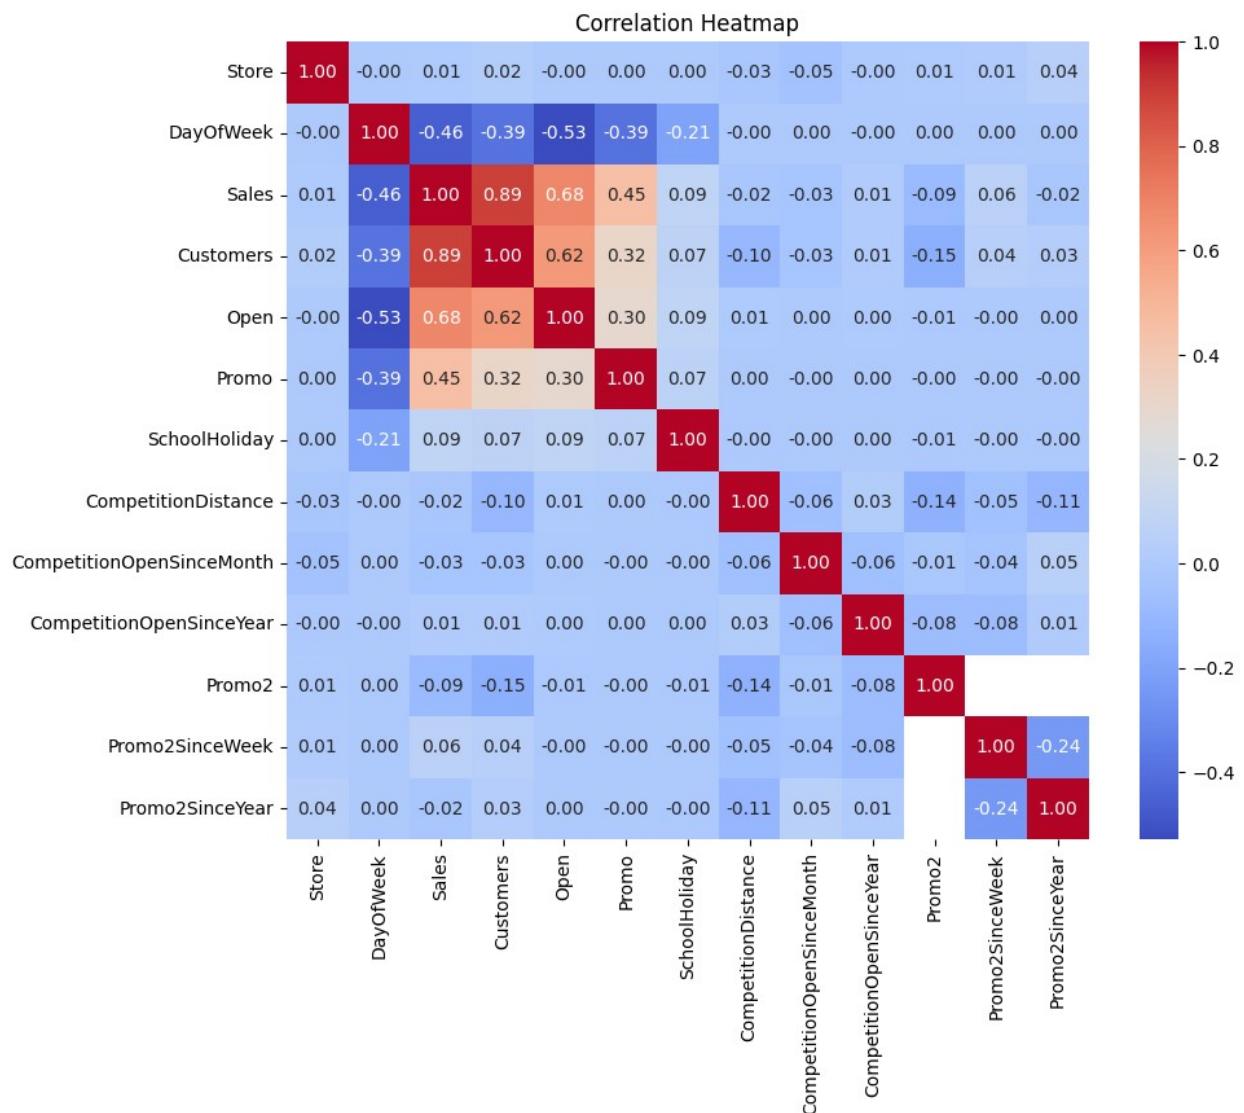
Sales vs. Promo: Shows a clear positive correlation.

DayOfWeek vs. Sales: This reflects the fact that as the Day of the Week index increases (towards 7 = Sunday), sales drop to zero.

CompetitionDistance: A weak correlation here suggests that the distance to a competitor matters less than whether a store is currently running a promotion or if it's a holiday.

Sales vs. open: Selfexplanatory, but less than i expected.

```
plt.figure(figsize=(10, 8))
corr = df.select_dtypes(include=['float64', 'int64']).corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```



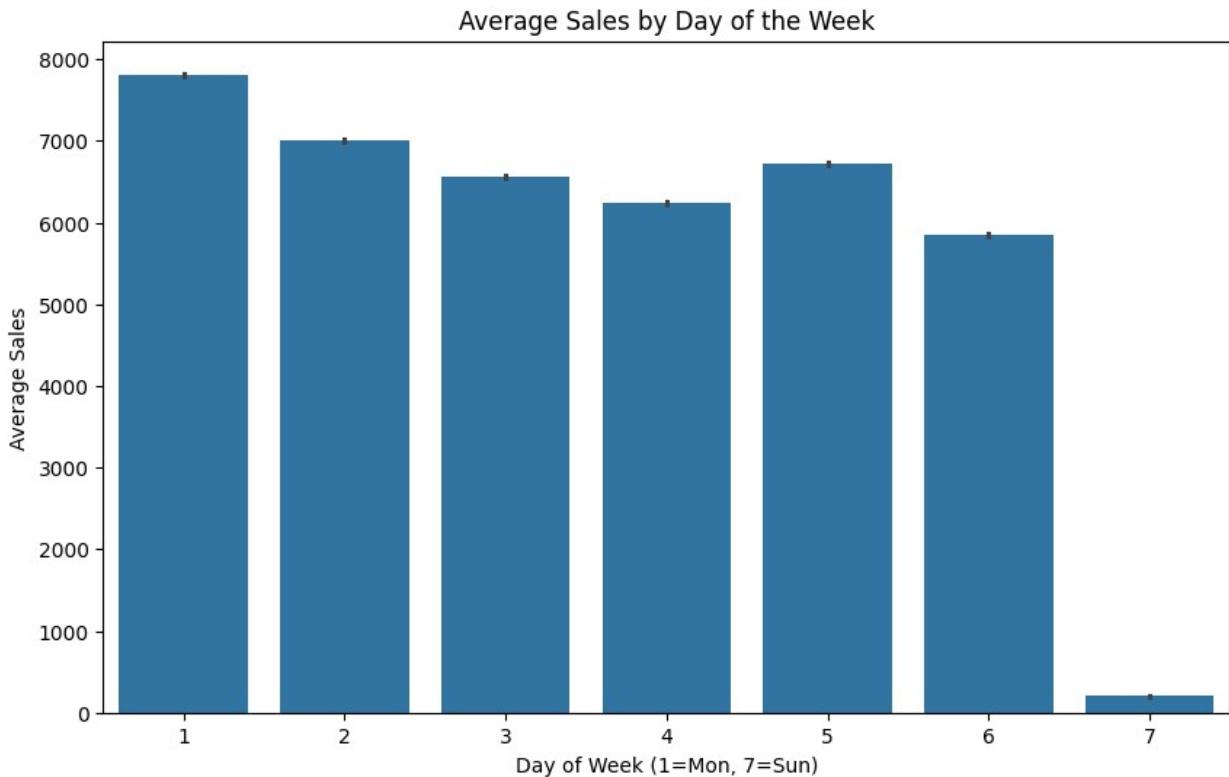
## 5. Average Sales by Day of the Week

As we realized in the previous parts the later in the week the less the sales.

```

plt.figure(figsize=(10, 6))
sns.barplot(x='DayOfWeek', y='Sales', data=df)
plt.title('Average Sales by Day of the Week')
plt.xlabel('Day of Week (1=Mon, 7=Sun)')
plt.ylabel('Average Sales')
plt.show()

```



## 6. Monthly Sales Seasonality

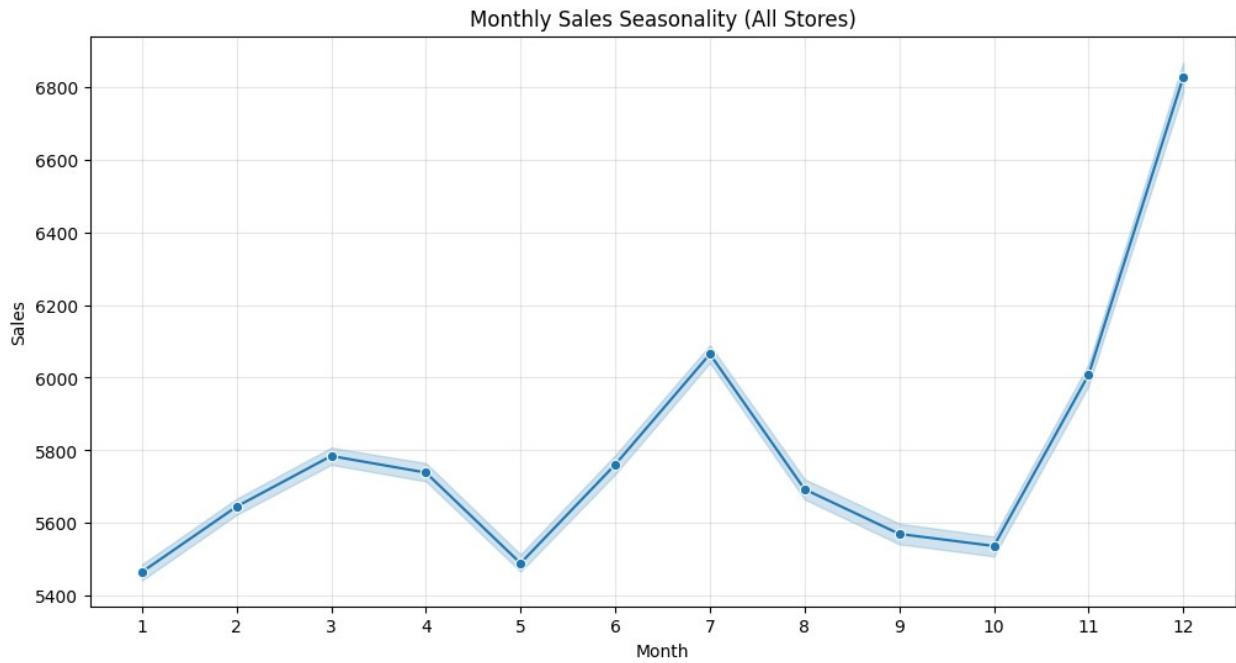
- January marks the lowest sales mainly because of christmas and new year hollidays.
- The highest sales is in December.
- In [1,3] , [5,7] , [10,12] we see a ascending trend.
- In [3,5], [5,10], we see a descending trend.

```

# Extract Month
df['Month'] = df['Date'].dt.month

plt.figure(figsize=(12, 6))
sns.lineplot(x='Month', y='Sales', data=df, marker='o')
plt.title('Monthly Sales Seasonality')
plt.xticks(range(1, 13))
plt.grid(True, alpha=0.3)
plt.show()

```

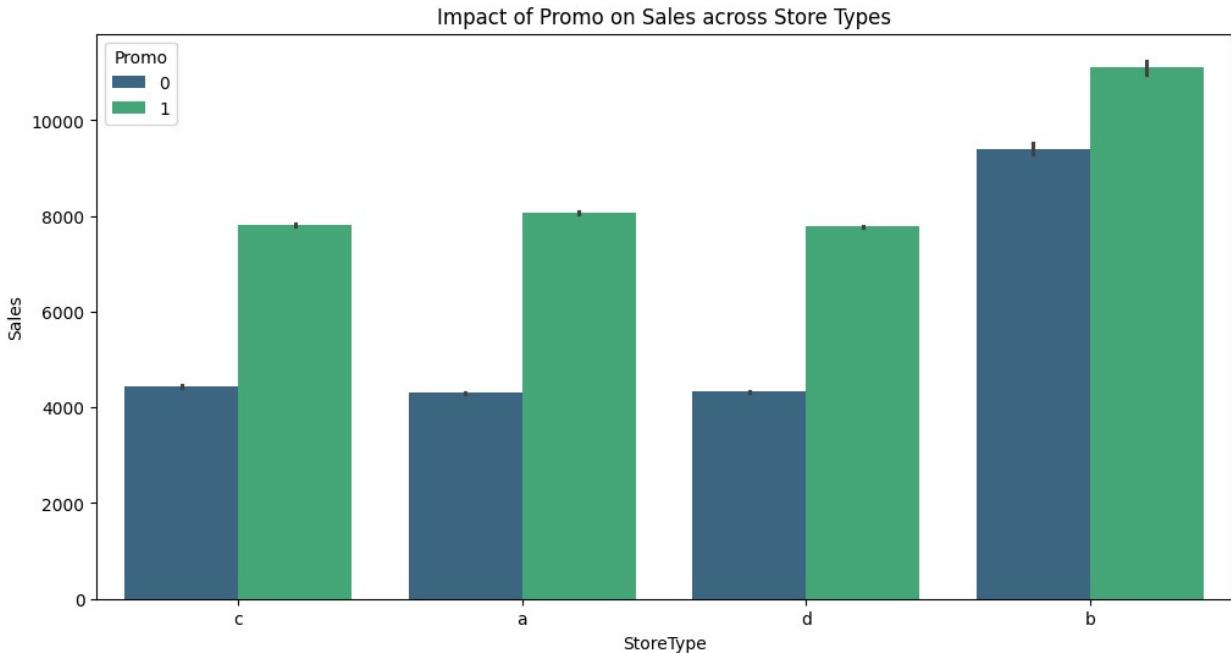


Not all stores react to promotions the same way. This grouped bar chart shows how Promo impacts different StoreType categories (a, b, c, d).

## 7.Impact of Promo on Sales across Store Types

- In type b stores that have higher sales promos yields less increase in sales compare to the rest (In percentage).

```
plt.figure(figsize=(12, 6))
sns.barplot(x='StoreType', y='Sales', hue='Promo', data=df,
palette='viridis')
plt.title('Impact of Promo on Sales across Store Types')
plt.show()
```



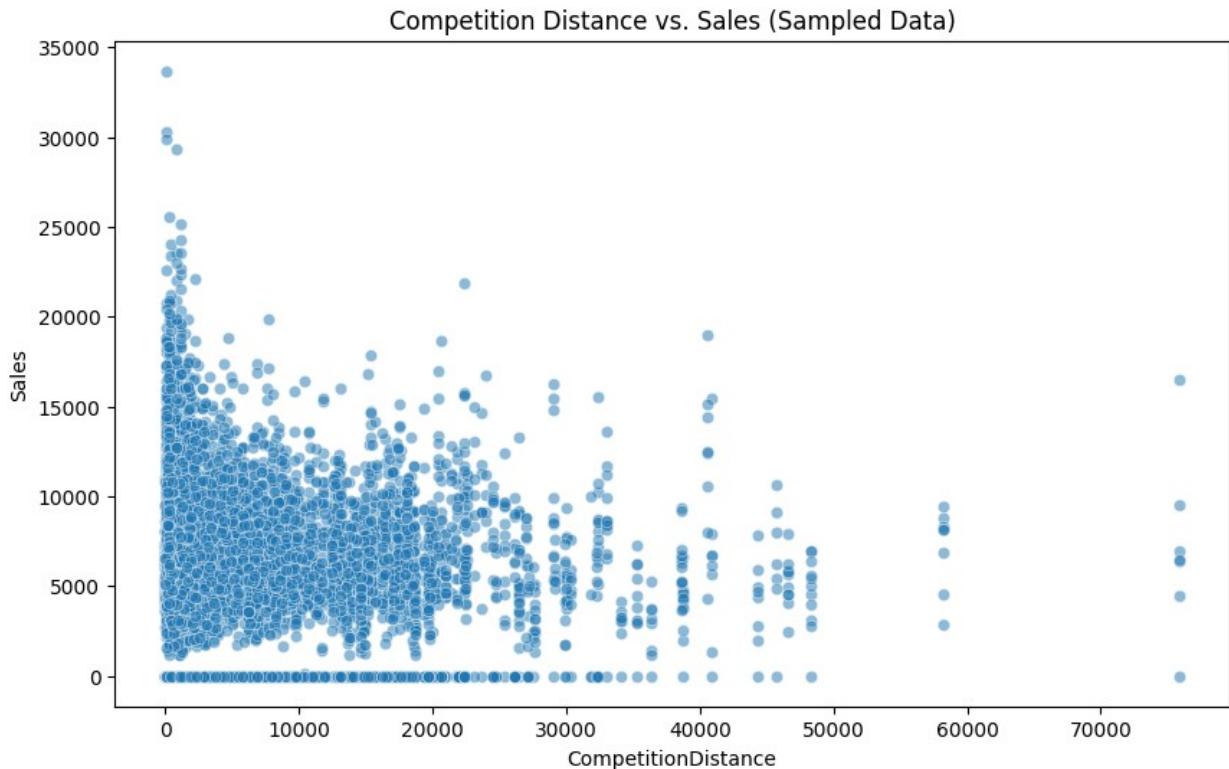
## 8. Competition Distance vs. Sales

Does having a competitor nearby actually hurt sales? We'll use a scatter plot with a regression line to see the correlation.

- This suggests that most Rossmann stores operate in highly competitive environments.
- We do not see a sharp decline in sales for stores with very close competitors. In fact, some of the highest-selling days occur when a competitor is within 1,000 meters.

```
sample_df = df.dropna(subset=['CompetitionDistance']).sample(10000)

plt.figure(figsize=(10, 6))
sns.scatterplot(x='CompetitionDistance', y='Sales', data=sample_df,
alpha=0.5)
plt.title('Competition Distance vs. Sales (Sampled Data)')
plt.show()
```



## 9.Sale distribution

Here we can see the sales distribution that has a mean of approximately 5000.

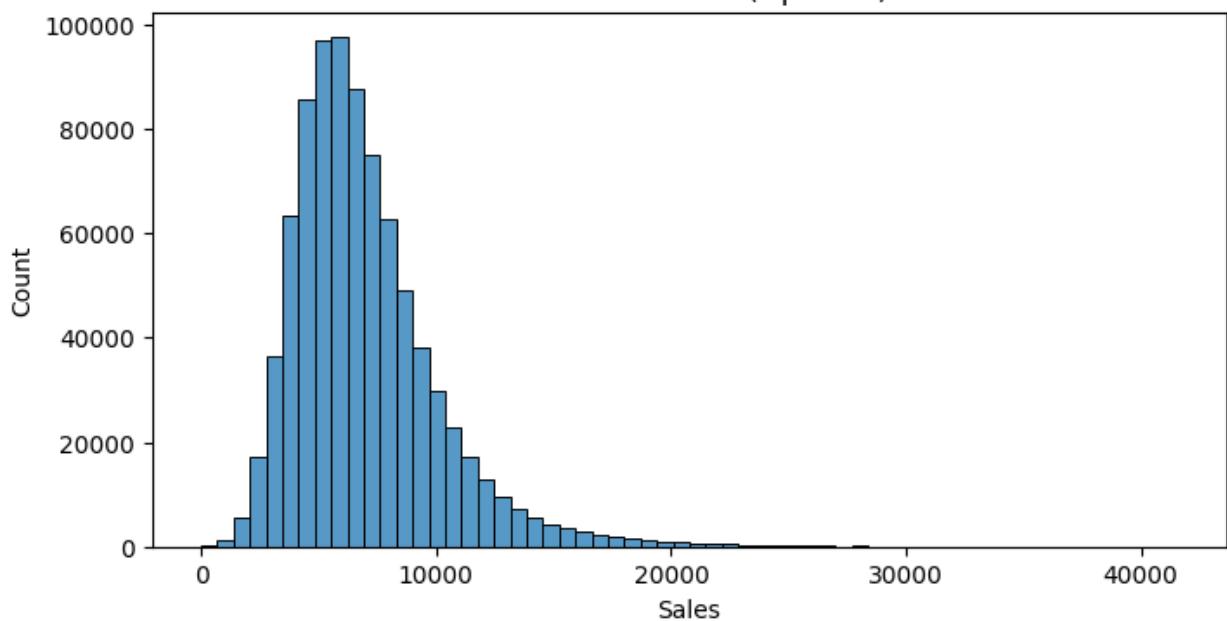
```

tmp = df[df["Open"]==1].copy()
plt.figure(figsize=(8,4))
sns.histplot(tmp["Sales"], bins=60, kde=False)
plt.title("Sales distribution (Open=1)")
plt.show()

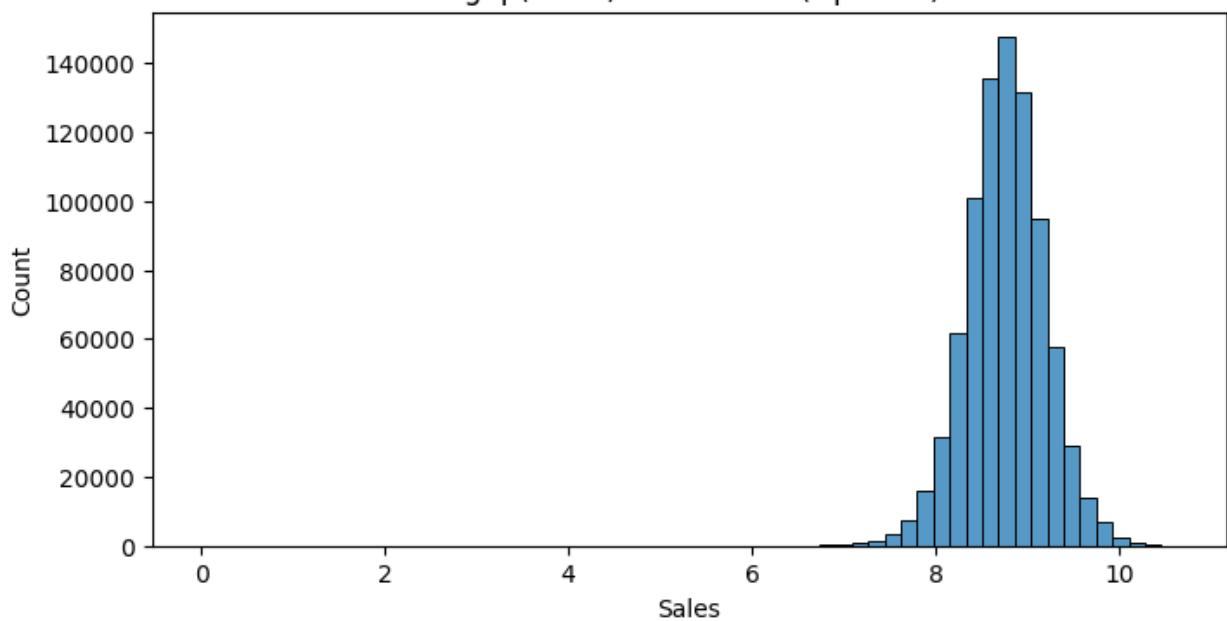
plt.figure(figsize=(8,4))
sns.histplot(np.log1p(tmp["Sales"]), bins=60, kde=False)
plt.title("log1p(Sales) distribution (Open=1)")
plt.show()

```

Sales distribution (Open=1)



log1p(Sales) distribution (Open=1)



```

df = pd.merge(train, store, on='Store', how='left')
print(df.shape)
print("stores in train:", train["Store"].nunique(), "stores in store:", store["Store"].nunique())
print("unmatched stores:", df["StoreType"].isna().sum())
df.head()

(1017209, 18)
stores in train: 1115 stores in store: 1115
unmatched stores: 0

      Store  DayOfWeek     Date   Sales  Customers  Open  Promo  StateHoliday \
0         1          5 2015-07-31    5263       555      1      1          0
1         2          5 2015-07-31    6064       625      1      1          0
2         3          5 2015-07-31    8314       821      1      1          0
3         4          5 2015-07-31  13995      1498      1      1          0
4         5          5 2015-07-31    4822       559      1      1          0

      SchoolHoliday  StoreType Assortment  CompetitionDistance \
0                  1          c            a                 1270.0
1                  1          a            a                  570.0
2                  1          a            a                14130.0
3                  1          c            c                  620.0
4                  1          a            a                29910.0

  CompetitionOpenSinceMonth  CompetitionOpenSinceYear  Promo2 \
0                      9.0                  2008.0      0
1                     11.0                  2007.0      1
2                     12.0                  2006.0      1
3                     9.0                  2009.0      0
4                     4.0                  2015.0      0

  Promo2SinceWeek  Promo2SinceYear  PromoInterval
0           NaN          NaN          NaN
1        13.0        2010.0  Jan,Apr,Jul,Oct
2        14.0        2011.0  Jan,Apr,Jul,Oct
3           NaN          NaN          NaN
4           NaN          NaN          NaN

df[:,]

```

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	\
0	1	5	2015-07-31	5263	555	1	1	
1	2	5	2015-07-31	6064	625	1	1	
2	3	5	2015-07-31	8314	821	1	1	
3	4	5	2015-07-31	13995	1498	1	1	
4	5	5	2015-07-31	4822	559	1	1	
...	...	...	...	...	...	...	...	...
1017204	1111	2	2013-01-01	0	0	0	0	

1017205	1112	2	2013-01-01	0	0	0	0
1017206	1113	2	2013-01-01	0	0	0	0
1017207	1114	2	2013-01-01	0	0	0	0
1017208	1115	2	2013-01-01	0	0	0	0
							\
0		0	1	c	a		1270.0
1		0	1	a	a		570.0
2		0	1	a	a		14130.0
3		0	1	c	c		620.0
4		0	1	a	a		29910.0
...	...	...	...	...	...		...
1017204	a	1	a	a			1900.0
1017205	a	1	c	c			1880.0
1017206	a	1	a	c			9260.0
1017207	a	1	a	c			870.0
1017208	a	1	d	c			5350.0
						\	
		CompetitionOpenSinceMonth	CompetitionOpenSinceYear	Promo2			\
0		9.0	2008.0	0			
1		11.0	2007.0	1			
2		12.0	2006.0	1			
3		9.0	2009.0	0			
4		4.0	2015.0	0			
...		...	...	...			
1017204		6.0	2014.0	1			
1017205		4.0	2006.0	0			
1017206		NaN	NaN	0			
1017207		NaN	NaN	0			
1017208		NaN	NaN	1			
						\	
		Promo2SinceWeek	Promo2SinceYear	PromoInterval			
0		NaN	NaN	NaN			
1		13.0	2010.0	Jan,Apr,Jul,Oct			
2		14.0	2011.0	Jan,Apr,Jul,Oct			
3		NaN	NaN	NaN			
4		NaN	NaN	NaN			
...		...	...	...			
1017204		31.0	2013.0	Jan,Apr,Jul,Oct			
1017205		NaN	NaN	NaN			
1017206		NaN	NaN	NaN			
1017207		NaN	NaN	NaN			
1017208		22.0	2012.0	Mar,Jun,Sept,Dec			

[1017209 rows x 18 columns]

```

nan_counts = df.isna().sum()
nan_total = int(df.isna().sum().sum())

print("Total NaN cells:", nan_total)
nan_counts[nan_counts > 0].sort_values(ascending=False)

Total NaN cells: 2173431

Promo2SinceYear      508031
Promo2SinceWeek      508031
PromoInterval        508031
CompetitionOpenSinceMonth 323348
CompetitionOpenSinceYear   323348
CompetitionDistance     2642
dtype: int64

import pandas as pd

df["Date"] = pd.to_datetime(df["Date"])

df["Year"] = df["Date"].dt.year
df["Month"] = df["Date"].dt.month
df["Day"] = df["Date"].dt.day
df["DayOfWeek0"] = df["Date"].dt.dayofweek
df["WeekOfYear"] = df["Date"].dt.isocalendar().week.astype(int)

df["StateHoliday"] = df["StateHoliday"].astype(str).map({"0": 0, "a": 1, "b": 1, "c": 1}).fillna(0)
df["SchoolHoliday"] = df["SchoolHoliday"].astype(int)

df = df.sort_values(["Store", "Date"]).reset_index(drop=True)

s1 = df.groupby("Store")["Sales"].shift(1)
df["SalesMovingAverage7"] = s1.groupby(df["Store"]).rolling(window=7, min_periods=1).mean()
df["SalesMovingAverage30"] = s1.groupby(df["Store"]).rolling(window=30, min_periods=1).mean()

df.head()

    Store  DayOfWeek      Date  Sales  Customers  Open  Promo  StateHoliday \
0       1          2  2013-01-01      0         0     0     0           1
1       1          3  2013-01-02  5530      668     1     0           0
2       1          4  2013-01-03  4327      578     1     0           0
3       1          5  2013-01-04  4486      619     1     0           0
4       1          6  2013-01-05  4997      635     1     0           0

    SchoolHoliday  StoreType  ...  Promo2SinceWeek  Promo2SinceYear \
0                  1        c  ...            NaN            NaN
1                  1        c  ...            NaN            NaN
2                  1        c  ...            NaN            NaN

```

```

3          1      c ...           NaN           NaN
4          1      c ...           NaN           NaN

   PromoInterval  Year  Month  Day  DayOfWeek0  WeekOfYear \
0            NaN  2013     1    1        1             1
1            NaN  2013     1    2        2             1
2            NaN  2013     1    3        3             1
3            NaN  2013     1    4        4             1
4            NaN  2013     1    5        5             1

   SalesMovingAverage7  SalesMovingAverage30
0            NaN           NaN
1          0.000000       0.000000
2         2765.000000      2765.000000
3        3285.666667      3285.666667
4        3585.750000      3585.750000

```

[5 rows x 25 columns]

In this step, we convert Date to a datetime format and extract calendar-based features such as year, month, day, and week number. Then, the holiday indicators (StateHoliday and SchoolHoliday) are encoded into numeric values. Finally, after sorting by store and date, we compute 7-day and 30-day moving averages of sales using only previous days (shift(1)) to avoid data leakage.

```

df = df.sort_values(["Store", "Date"]).reset_index(drop=True)

g = df.groupby("Store")["Sales"]

df["SalesLag1"] = g.shift(1)
df["SalesLag7"] = g.shift(7)

s1 = g.shift(1)
df["SalesMean7"] = s1.groupby(df["Store"]).rolling(7, min_periods=1).mean().reset_index(level=0)
df["SalesStd7"] = s1.groupby(df["Store"]).rolling(7, min_periods=2).std().reset_index(level=0)

df["SalesMean30"] = s1.groupby(df["Store"]).rolling(30, min_periods=1).mean().reset_index(level=0)
df["SalesStd30"] = s1.groupby(df["Store"]).rolling(30, min_periods=2).std().reset_index(level=0)

df.head()

   Store  DayOfWeek      Date  Sales  Customers  Open  Promo  StateHoliday \
0      1          2 2013-01-01      0        0      0      0           1
1      1          3 2013-01-02  5530      668      1      0           0
2      1          4 2013-01-03  4327      578      1      0           0
3      1          5 2013-01-04  4486      619      1      0           0
4      1          6 2013-01-05  4997      635      1      0           0

```

```

    SchoolHoliday StoreType ... DayOfWeek0 WeekOfYear SalesMovingAverage7 \
0           1      c   ...       1        1               NaN
1           1      c   ...       2        1      0.000000
2           1      c   ...       3        1      2765.000000
3           1      c   ...       4        1      3285.666667
4           1      c   ...       5        1      3585.750000

    SalesMovingAverage30  SalesLag1  SalesLag7  SalesMean7  SalesStd7 \
0            NaN       NaN       NaN       NaN       NaN
1          0.000000      0.0       NaN      0.000000      NaN
2         2765.000000    5530.0     NaN    2765.000000  3910.300500
3         3285.666667    4327.0     NaN    3285.666667  2908.351137
4         3585.750000    4486.0     NaN    3585.750000  2449.327306

    SalesMean30  SalesStd30
0            NaN       NaN
1          0.000000      NaN
2         2765.000000  3910.300500
3         3285.666667  2908.351137
4         3585.750000  2449.327306

```

[5 rows x 31 columns]

To prevent the moving averages from becoming zero in the first days, we computed rolling features with min\_periods=1 (and min\_periods=2 for standard deviation), so the statistics are calculated using whatever history is available instead of producing NaNs that would later be filled with zeros.

```

df[['Store', 'Date', 'Sales', 'SalesLag1', 'SalesLag7', 'SalesMean7', 'SalesStd7', 'SalesMean30', 'SalesStd30']
   Store      Date  Sales  SalesLag1  SalesLag7  SalesMean7  SalesStd7 \
0     1  2013-01-01     0       NaN       NaN       NaN       NaN
1     1  2013-01-02   5530      0.0       NaN      0.000000      NaN
2     1  2013-01-03   4327    5530.0       NaN    2765.000000  3910.300500
3     1  2013-01-04   4486    4327.0       NaN    3285.666667  2908.351137
4     1  2013-01-05   4997    4486.0       NaN    3585.750000  2449.327306
5     1  2013-01-06     0    4997.0       NaN    3868.000000  2213.081223
6     1  2013-01-07   7176      0.0       NaN    3223.333333  2532.144045
7     1  2013-01-08   5580    7176.0      0.0    3788.000000  2752.283961
8     1  2013-01-09   5471    5580.0    5530.0    4585.142857  2231.018633
9     1  2013-01-10   4892    5471.0    4327.0    4576.714286  2226.961885

   SalesMean30  SalesStd30
0            NaN       NaN
1          0.000000      NaN
2         2765.000000  3910.300500

```

```

3 3285.666667 2908.351137
4 3585.750000 2449.327306
5 3868.000000 2213.081223
6 3223.333333 2532.144045
7 3788.000000 2752.283961
8 4012.000000 2625.704205
9 4174.111111 2503.807573

```

We also added 30-day rolling statistics based only on past sales: SalesMean30 (30-day rolling mean) and SalesStd30 (30-day rolling standard deviation) computed per store using shifted sales to avoid leakage.

```

df = train.merge(store, on="Store", how="left")

df[ "Date" ] = pd.to_datetime(df[ "Date" ])

df[ "StateHoliday" ] = df[ "StateHoliday" ].fillna("0").astype(str)
df[ "IsStateHoliday" ] = (df[ "StateHoliday" ] != "0").astype(int)

print(df[ "StateHoliday" ].value_counts(dropna=False).head(10))
print(df[ "IsStateHoliday" ].value_counts(dropna=False))
df_L=df.copy()

StateHoliday
0    986159
a    20260
b    6690
c    4100
Name: count, dtype: int64
IsStateHoliday
0    986159
1    31050
Name: count, dtype: int64

import numpy as np
import pandas as pd

df = df.sort_values(["Date"]).reset_index(drop=True)

df[ "DayOfYear" ] = df[ "Date" ].dt.dayofyear
t = 2*np.pi*df[ "DayOfYear" ]/365.25
df[ "sin_1y" ] = np.sin(t)
df[ "cos_1y" ] = np.cos(t)
df[ "sin_2y" ] = np.sin(2*t)
df[ "cos_2y" ] = np.cos(2*t)

dates = pd.Index(df[ "Date" ].unique()).sort_values()

```

```

is_h = df.groupby("Date")["IsStateHoliday"].max().reindex(dates).fillna(0).astype(int)
h_dates = dates[is_h.values == 1]

prev_h = pd.Series(pd.NaT, index=dates)
next_h = pd.Series(pd.NaT, index=dates)

prev_h.loc[h_dates] = h_dates
next_h.loc[h_dates] = h_dates
prev_h = prev_h.ffill()
next_h = next_h.bfill()

dist_prev = (dates - prev_h.values).days
dist_next = (next_h.values - dates).days

df["DaysSinceHoliday"] = df["Date"].map(pd.Series(dist_prev, index=dates).fillna(0).astype(int))
df["DaysToNextHoliday"] = df["Date"].map(pd.Series(dist_next, index=dates).fillna(0).astype(int))

df.head()

   Store DayOfWeek      Date  Sales  Customers  Open  Promo StateHoliday \
0    1115         2 2013-01-01     0        0       0      0          a
1     379         2 2013-01-01     0        0       0      0          a
2     378         2 2013-01-01     0        0       0      0          a
3     377         2 2013-01-01     0        0       0      0          a
4     376         2 2013-01-01     0        0       0      0          a

   SchoolHoliday StoreType ... Promo2SinceYear  PromoInterval \
0              1        d ...           2012.0  Mar,Jun,Sept,Dec
1              1        d ...             NaN            NaN
2              1        a ...             NaN            NaN
3              1        a ...           2010.0  Feb,May,Aug,Nov
4              1        a ...             NaN            NaN

   IsStateHoliday  DayOfYear  sin_1y  cos_1y  sin_2y  cos_2y \
0                 1         1  0.017202  0.999852  0.034398  0.999408
1                 1         1  0.017202  0.999852  0.034398  0.999408
2                 1         1  0.017202  0.999852  0.034398  0.999408
3                 1         1  0.017202  0.999852  0.034398  0.999408
4                 1         1  0.017202  0.999852  0.034398  0.999408

   DaysSinceHoliday  DaysToNextHoliday
0                  0                  0
1                  0                  0
2                  0                  0
3                  0                  0
4                  0                  0

```

```
[5 rows x 26 columns]
```

We build Fourier seasonal features using DayOfYear. The terms sin\_1y/cos\_1y capture the main annual (1-year) seasonality while sin\_2y/cos\_2y model more complex within-year patterns (e.g., two peaks during the year).

```
df[['Date', 'IsStateHoliday', 'DaysSinceHoliday', 'DaysToNextHoliday']].head(10000)
```

	Date	IsStateHoliday	DaysSinceHoliday	DaysToNextHoliday
0	2013-01-01	1	0	0
1	2013-01-01	1	0	0
2	2013-01-01	1	0	0
3	2013-01-01	1	0	0
4	2013-01-01	1	0	0
...	...	...	...	...
9995	2013-01-09	0	3	79
9996	2013-01-09	0	3	79
9997	2013-01-09	0	3	79
9998	2013-01-09	0	3	79
9999	2013-01-09	0	3	79

```
[10000 rows x 4 columns]
```

```
print(df['IsStateHoliday'].value_counts())
print(df.loc[df['IsStateHoliday']==1, ['Date']].drop_duplicates().head(10))
```

```
IsStateHoliday
```

```
0    986159
```

```
1     31050
```

```
Name: count, dtype: int64
```

```
      Date
```

```
0    2013-01-01
```

```
5575   2013-01-06
```

```
97004  2013-03-29
```

```
100349 2013-04-01
```

```
133799 2013-05-01
```

```
142719 2013-05-09
```

```
154984 2013-05-20
```

```
166134 2013-05-30
```

```
252007 2013-08-15
```

```
306624 2013-10-03
```

```
missing_counts = df.isna().sum()
```

```
missing_pct = (df.isna().mean() * 100).round(3)
```

```
missing_summary = pd.DataFrame({
```

```
    "missing_count": missing_counts,
```

```
    "missing_percent": missing_pct}
```

```

}).sort_values("missing_count", ascending=False)

missing_summary[missing_summary["missing_count"] > 0]

missing_count    missing_percent
PromoInterval      508031        49.944
Promo2SinceYear     508031        49.944
Promo2SinceWeek     508031        49.944
CompetitionOpenSinceYear 323348        31.788
CompetitionOpenSinceMonth 323348        31.788
CompetitionDistance    2642         0.260

df_z = df.copy()

num_cols = df_z.select_dtypes(include=[np.number]).columns
bin_cols = [c for c in num_cols if df_z[c].dropna().isin([0,1]).all()]
z_cols = [c for c in num_cols if c not in bin_cols]

df_z[z_cols] = df_z[z_cols].fillna(df_z[z_cols].median())

mu = df_z[z_cols].mean()
sigma = df_z[z_cols].std().replace(0, 1)

df_z[z_cols] = (df_z[z_cols] - mu) / sigma

df_z.head()

      Store  DayOfWeek      Date   Sales Customers  Open  Promo \
0  1.728970 -1.000475 2013-01-01 -1.499722 -1.363329  0    0
1 -0.557393 -1.000475 2013-01-01 -1.499722 -1.363329  0    0
2 -0.560500 -1.000475 2013-01-01 -1.499722 -1.363329  0    0
3 -0.563606 -1.000475 2013-01-01 -1.499722 -1.363329  0    0
4 -0.566713 -1.000475 2013-01-01 -1.499722 -1.363329  0    0

      StateHoliday SchoolHoliday StoreType ... Promo2SinceYear \
0            a             1       d ...          0.104612
1            a             1       d ...          0.104612
2            a             1       a ...          0.104612
3            a             1       a ...         -1.586055
4            a             1       a ...          0.104612

      PromoInterval IsStateHoliday DayOfYear   sin_1y   cos_1y   sin_2y \
0  Mar,Jun,Sept,Dec           NaN      1 -1.587113 -0.179271  1.442531  0.027059
1                 NaN           1 -1.587113 -0.179271  1.442531  0.027059
2                 NaN           1 -1.587113 -0.179271  1.442531  0.027059
3  Feb,May,Aug,Nov           NaN      1 -1.587113 -0.179271  1.442531  0.027059
4                 NaN           1 -1.587113 -0.179271  1.442531  0.027059

```

```

          cos_2y  DaysSinceHoliday  DaysToNextHoliday
0  1.368905      -1.104201      -1.011865
1  1.368905      -1.104201      -1.011865
2  1.368905      -1.104201      -1.011865
3  1.368905      -1.104201      -1.011865
4  1.368905      -1.104201      -1.011865

[5 rows x 26 columns]

df_mm = df.copy()

num_cols = df_mm.select_dtypes(include=[np.number]).columns
df_mm[num_cols] = df_mm[num_cols].fillna(df_mm[num_cols].median())

mn = df_mm[num_cols].min()
mx = df_mm[num_cols].max()
den = (mx - mn).replace(0, 1)

df_mm[num_cols] = (df_mm[num_cols] - mn) / den

df_mm.head()

      Store  DayOfWeek      Date   Sales  Customers  Open  Promo StateHoliday \
0  0.000000  0.166667 2013-01-01     0.0       0.0     0.0     0.0         a
1  0.000898  0.166667 2013-01-01     0.0       0.0     0.0     0.0         a
2  0.001795  0.166667 2013-01-01     0.0       0.0     0.0     0.0         a
3  0.002693  0.166667 2013-01-01     0.0       0.0     0.0     0.0         a
4  0.003591  0.166667 2013-01-01     0.0       0.0     0.0     0.0         a

      SchoolHoliday  StoreType  ...  Promo2SinceYear  PromoInterval \
0            1.0        c    ...        0.500000           NaN
1            1.0        a    ...        0.166667  Jan,Apr,Jul,Oct
2            1.0        a    ...        0.333333  Jan,Apr,Jul,Oct
3            1.0        c    ...        0.500000           NaN
4            1.0        a    ...        0.500000           NaN

      IsStateHoliday  DayOfYear   sin_1y   cos_1y   sin_2y   cos_2y \
0            1.0        0.0  0.508604  0.999931  0.517211  0.999723
1            1.0        0.0  0.508604  0.999931  0.517211  0.999723
2            1.0        0.0  0.508604  0.999931  0.517211  0.999723
3            1.0        0.0  0.508604  0.999931  0.517211  0.999723
4            1.0        0.0  0.508604  0.999931  0.517211  0.999723

      DaysSinceHoliday  DaysToNextHoliday
0                  0.0              0.0
1                  0.0              0.0

```

```

2          0.0          0.0
3          0.0          0.0
4          0.0          0.0

[5 rows x 26 columns]

df_mm[:]

      Store  DayOfWeek      Date    Sales  Customers  Open  Promo \
0     0.000000  0.166667 2013-01-01  0.000000  0.000000  0.0   0.0
1     0.000898  0.166667 2013-01-01  0.000000  0.000000  0.0   0.0
2     0.001795  0.166667 2013-01-01  0.000000  0.000000  0.0   0.0
3     0.002693  0.166667 2013-01-01  0.000000  0.000000  0.0   0.0
4     0.003591  0.166667 2013-01-01  0.000000  0.000000  0.0   0.0
...
1017204  0.996409  0.666667 2015-07-31  0.137734  0.057120  1.0   1.0
1017205  0.997307  0.666667 2015-07-31  0.231667  0.103817  1.0   1.0
1017206  0.998205  0.666667 2015-07-31  0.175423  0.097455  1.0   1.0
1017207  0.999102  0.666667 2015-07-31  0.662030  0.506903  1.0   1.0
1017208  1.000000  0.666667 2015-07-31  0.208900  0.072821  1.0   1.0

      StateHoliday  SchoolHoliday  StoreType  ...  Promo2SinceYear \
0             a           1.0        c  ...       0.500000
1             a           1.0        a  ...       0.166667
2             a           1.0        a  ...       0.333333
3             a           1.0        c  ...       0.500000
4             a           1.0        a  ...       0.500000
...
1017204         0           1.0        a  ...       0.666667
1017205         0           1.0        c  ...       0.500000
1017206         0           1.0        a  ...       0.500000
1017207         0           1.0        a  ...       0.500000
1017208         0           1.0        d  ...       0.500000

      PromoInterval  IsStateHoliday  DayOfYear    sin_1y    cos_1y \
0            0.0           1.0      0.00000  0.508604  0.999931
1            1.0           1.0      0.00000  0.508604  0.999931
2            1.0           1.0      0.00000  0.508604  0.999931
3            0.0           1.0      0.00000  0.508604  0.999931
4            0.0           1.0      0.00000  0.508604  0.999931
...
1017204         1.0           0.0      0.57967  0.257957  0.062481
1017205         0.0           0.0      0.57967  0.257957  0.062481
1017206         0.0           0.0      0.57967  0.257957  0.062481
1017207         0.0           0.0      0.57967  0.257957  0.062481
1017208         0.0           0.0      0.57967  0.257957  0.062481

```

	sin_2y	cos_2y	DaysSinceHoliday	DaysToNextHoliday
0	0.517211	0.999723	0.000000	0.0
1	0.517211	0.999723	0.000000	0.0
2	0.517211	0.999723	0.000000	0.0
3	0.517211	0.999723	0.000000	0.0
4	0.517211	0.999723	0.000000	0.0
...	...	...	...	...
1017204	0.923608	0.765672	0.542857	0.0
1017205	0.923608	0.765672	0.542857	0.0
1017206	0.923608	0.765672	0.542857	0.0
1017207	0.923608	0.765672	0.542857	0.0
1017208	0.923608	0.765672	0.542857	0.0

[1017209 rows x 26 columns]

The PromoInterval column was transformed from a sparse categorical string into a binary numeric feature by mapping the current month to the scheduled promotion cycles. This process eliminated all NaN values by assigning 1 to active promotion months and 0 to others, allowing the model to effectively capture the impact of recurring discounts during training and scaling.

We implemented two independent normalization strategies in separate cells. First, we created df\_z and applied Z-score standardization only to non-binary numeric features (leaving 0/1 indicators unchanged), so those features have approximately zero mean and unit variance. Second, we created df\_mm and applied Min–Max normalization to scale numeric features into the [0,1] range. Using separate copies ensures the two normalizations do not affect each other.

To handle missing values, we first identified all NaN entries and checked which columns contained missing data. We treated missingness based on the meaning of each feature:

Categorical features (e.g., PromoInterval): NaNs often mean “not applicable” (e.g., stores without Promo2).

Numeric features (e.g., competition and Promo2 date/distance fields): when NaNs represent “unknown” or “not reported,” we filled them with a robust statistic (typically the median) to avoid losing rows, and (when relevant) we also used an indicator meaning “feature exists/active” (e.g., competition exists vs. not).

Before applying normalization, we ensured there were no remaining NaNs in numeric columns by imputing them (median), and we handled edge cases where a column had zero variance to avoid division-by-zero during scaling.

```
df = df_mm.sort_values(["Date", "Store"]).reset_index(drop=True)

dates = np.array(sorted(df["Date"].unique()))
n = len(dates)
```

```

train_end = int(0.70 * n)
val_end   = int(0.85 * n)

train_dates = set(dates[:train_end])
val_dates   = set(dates[train_end:val_end])
test_dates  = set(dates[val_end:])

train_df = df[df["Date"].isin(train_dates)].reset_index(drop=True)
val_df   = df[df["Date"].isin(val_dates)].reset_index(drop=True)
test_df  = df[df["Date"].isin(test_dates)].reset_index(drop=True)

print("Monotonic:", df["Date"].is_monotonic_increasing)
print(f"Train dates: {train_df['Date'].min()} to {train_df['Date'].max()}")
print(f"Val dates:  {val_df['Date'].min()} to {val_df['Date'].max()}")
print(f"Test dates: {test_df['Date'].min()} to {test_df['Date'].max()")

print("No overlap train/val:", train_df["Date"].max() < val_df["Date"].min())
print("No overlap val/test :", val_df["Date"].max() < test_df["Date"].min())

Monotonic: True
Train dates: 2013-01-01 00:00:00 to 2014-10-21 00:00:00
Val dates:  2014-10-22 00:00:00 to 2015-03-11 00:00:00
Test dates: 2015-03-12 00:00:00 to 2015-07-31 00:00:00
No overlap train/val: True
No overlap val/test : True

print("Train rows:", len(train_df))
print("Val rows  :", len(val_df))
print("Test rows :", len(test_df))

print("Train unique dates:", train_df["Date"].nunique())
print("Val unique dates  :", val_df["Date"].nunique())
print("Test unique dates  :", test_df["Date"].nunique())

Train rows: 714444
Val rows  : 144435
Test rows : 158330
Train unique dates: 659
Val unique dates  : 141
Test unique dates  : 142

train_df.PromoInterval.value_counts()

PromoInterval
Jan,Apr,Jul,Oct      206979
Feb,May,Aug,Nov      83297
Mar,Jun,Sept,Dec     68385

```

Name: count, dtype: int64

We split the dataset into train/validation/test while strictly preserving chronological order. To avoid any date overlap (since each date contains many stores/rows), we performed the split using the unique dates rather than row indices. This guarantees monotonic time ordering and ensures that the validation and test sets contain only future dates relative to the training set (no temporal leakage).

```
import numpy as np

mean_pred = np.full(len(val_df), train_df["Sales"].mean(), dtype=float)
mean_pred = np.where(val_df["Open"].values==0, 0.0, mean_pred)
rmse_mean = np.sqrt(np.mean((val_df["Sales"].values.astype(float) - mean_pred) ** 2))

train_open = train_df[train_df["Open"]==1]
store_mean = train_open.groupby("Store")["Sales"].mean()
store_pred = val_df["Store"].map(store_mean).fillna(train_open["Sales"].mean()).values.astype(float)
store_pred = np.where(val_df["Open"].values==0, 0.0, store_pred)
rmse_store = np.sqrt(np.mean((val_df["Sales"].values.astype(float) - store_pred) ** 2))

print("Baseline RMSE (Global mean):", rmse_mean)
print("Baseline RMSE (Store mean): ", rmse_store)

Baseline RMSE (Global mean): 0.07878866888416793
Baseline RMSE (Store mean): 0.046653120646034
```

We report two simple baselines on the validation set: (1) a global-mean predictor that always predicts the average training sales and (2) a stronger store-level baseline that predicts each stores average sales computed from open days in the training set. For both baselines, we force predictions to 0 when Open=0 to match the business logic.

```
df_mm[:]
```

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	\
0	1.000000	0.166667	2013-01-01	0.000000	0.000000	0.0	0.0	
1	0.339318	0.166667	2013-01-01	0.000000	0.000000	0.0	0.0	
2	0.338420	0.166667	2013-01-01	0.000000	0.000000	0.0	0.0	
3	0.337522	0.166667	2013-01-01	0.000000	0.000000	0.0	0.0	
4	0.336625	0.166667	2013-01-01	0.000000	0.000000	0.0	0.0	
...	...	...	...	...	...	...	...	
1017204	0.668761	0.666667	2015-07-31	0.218575	0.086356	1.0	1.0	
1017205	0.669659	0.666667	2015-07-31	0.257707	0.111803	1.0	1.0	
1017206	0.670557	0.666667	2015-07-31	0.180044	0.078235	1.0	1.0	
1017207	0.665171	0.666667	2015-07-31	0.251739	0.137520	1.0	1.0	
1017208	0.000000	0.666667	2015-07-31	0.126664	0.075122	1.0	1.0	
	StateHoliday	SchoolHoliday	StoreType	...	Promo2SinceYear	\		

```

0          a      1.0      d ... 0.500000
1          a      1.0      d ... 0.500000
2          a      1.0      a ... 0.500000
3          a      1.0      a ... 0.166667
4          a      1.0      a ... 0.500000
...
1017204     0      1.0      d ... 0.333333
1017205     0      1.0      c ... 0.500000
1017206     0      1.0      d ... 0.333333
1017207     0      1.0      d ... 0.500000
1017208     0      1.0      c ... 0.500000

          PromoInterval IsStateHoliday DayOfYear   sin_1y   cos_1y \
0    Mar,Jun,Sept,Dec           1.0  0.00000  0.508604  0.999931
1            NaN               1.0  0.00000  0.508604  0.999931
2            NaN               1.0  0.00000  0.508604  0.999931
3  Feb,May,Aug,Nov           1.0  0.00000  0.508604  0.999931
4            NaN               1.0  0.00000  0.508604  0.999931
...
1017204  Mar,Jun,Sept,Dec       0.0  0.57967  0.257957  0.062481
1017205            NaN       0.0  0.57967  0.257957  0.062481
1017206  Jan,Apr,Jul,Oct       0.0  0.57967  0.257957  0.062481
1017207            NaN       0.0  0.57967  0.257957  0.062481
1017208            NaN       0.0  0.57967  0.257957  0.062481

          sin_2y   cos_2y DaysSinceHoliday DaysToNextHoliday
0  0.517211  0.999723        0.000000             0.0
1  0.517211  0.999723        0.000000             0.0
2  0.517211  0.999723        0.000000             0.0
3  0.517211  0.999723        0.000000             0.0
4  0.517211  0.999723        0.000000             0.0
...
1017204  0.923608  0.765672        0.542857             0.0
1017205  0.923608  0.765672        0.542857             0.0
1017206  0.923608  0.765672        0.542857             0.0
1017207  0.923608  0.765672        0.542857             0.0
1017208  0.923608  0.765672        0.542857             0.0

[1017209 rows x 26 columns]

import numpy as np
import pandas as pd

df_mm["Date"] = pd.to_datetime(df_mm["Date"])
df_mm = df_mm.sort_values(["Store", "Date"]).reset_index(drop=True)

```

```

if "Year" not in df_mm.columns: df_mm["Year"] = df_mm["Date"].dt.year
if "Month" not in df_mm.columns: df_mm["Month"] = df_mm["Date"].dt.month
if "Day" not in df_mm.columns: df_mm["Day"] = df_mm["Date"].dt.day
if "WeekOfYear" not in df_mm.columns: df_mm["WeekOfYear"] = df_mm["Date"].dt.isocalendar().week
if "IsWeekend" not in df_mm.columns: df_mm["IsWeekend"] = (df_mm["DayOfWeek"] >= 6).astype(bool)
if "IsMonthStart" not in df_mm.columns: df_mm["IsMonthStart"] = df_mm["Date"].dt.is_month_start
if "IsMonthEnd" not in df_mm.columns: df_mm["IsMonthEnd"] = df_mm["Date"].dt.is_month_end.astype(bool)
if "Quarter" not in df_mm.columns: df_mm["Quarter"] = df_mm["Date"].dt.quarter

g = df_mm.groupby("Store")["Sales"]
if "SalesLag1" not in df_mm.columns: df_mm["SalesLag1"] = g.shift(1)
if "SalesLag7" not in df_mm.columns: df_mm["SalesLag7"] = g.shift(7)

s1 = g.shift(1)
if "SalesMean7" not in df_mm.columns:
    df_mm["SalesMean7"] = s1.groupby(df_mm["Store"]).rolling(7, min_periods=1).mean().reset_index()
if "SalesStd7" not in df_mm.columns:
    df_mm["SalesStd7"] = s1.groupby(df_mm["Store"]).rolling(7, min_periods=2).std().reset_index()
if "SalesMean30" not in df_mm.columns:
    df_mm["SalesMean30"] = s1.groupby(df_mm["Store"]).rolling(30, min_periods=1).mean().reset_index()
if "SalesStd30" not in df_mm.columns:
    df_mm["SalesStd30"] = s1.groupby(df_mm["Store"]).rolling(30, min_periods=2).std().reset_index()

dates = np.array(sorted(df_mm["Date"].unique()))
n = len(dates)
train_end = int(0.70 * n)
val_end = int(0.85 * n)

train_dates = set(dates[:train_end])
val_dates = set(dates[train_end:val_end])
test_dates = set(dates[val_end:])

train_df = df_mm[df_mm["Date"].isin(train_dates)].reset_index(drop=True)
val_df = df_mm[df_mm["Date"].isin(val_dates)].reset_index(drop=True)
test_df = df_mm[df_mm["Date"].isin(test_dates)].reset_index(drop=True)

missing = [c for c in ["Year", "Month", "Day", "WeekOfYear", "IsWeekend", "IsMonthStart", "IsMonthEnd",
                      "SalesLag1", "SalesLag7", "SalesMean7", "SalesStd7", "SalesMean30", "SalesStd30"]]
print("Missing after build:", missing)
print(train_df.shape, val_df.shape, test_df.shape)

Missing after build: []
(714444, 40) (144435, 40) (158330, 40)

import numpy as np

rng = np.random.default_rng(42)

```

```

tr_n = min(50000, Xtr_np.shape[0])
va_n = min(20000, Xva_np.shape[0])

tr_idx = rng.choice(Xtr_np.shape[0], size=tr_n, replace=False)
va_idx = rng.choice(Xva_np.shape[0], size=va_n, replace=False)

Xtr_s = Xtr_np[tr_idx]
ytr_s = ytr[tr_idx]

Xva_s = Xva_np[va_idx]
yva_log_s = yva_log[va_idx]
yva_true_s = yva_true[va_idx]
open_va_s = open_va[va_idx]

print(Xtr_s.shape, Xva_s.shape)
(50000, 31) (20000, 31)

import numpy as np
import pandas as pd

feature_cols = [
    "Store", "DayOfWeek", "Open", "Promo", "SchoolHoliday", "StateHoliday",
    "StoreType", "Assortment",
    "CompetitionDistance", "Promo2",
    "Year", "Month", "Day", "WeekOfYear", "DayOfYear", "IsWeekend", "IsMonthStart", "IsMonthEnd",
    "SalesLag1", "SalesLag7", "SalesMean7", "SalesStd7", "SalesMean30", "SalesStd30",
    "sin_1y", "cos_1y", "sin_2y", "cos_2y",
    "DaysSinceHoliday", "DaysToNextHoliday"
]

cat_cols = ["Store", "DayOfWeek", "StateHoliday", "StoreType", "Assortment"]
cat_cols = [c for c in cat_cols if c in feature_cols]
num_cols = [c for c in feature_cols if c not in cat_cols]

def rmse(y_true, y_pred):
    e = y_true - y_pred
    return float(np.sqrt(np.mean(e * e)))

def prepare_matrices(train_df, val_df, test_df):
    Xtr = train_df[feature_cols].copy()
    Xva = val_df[feature_cols].copy()
    Xte = test_df[feature_cols].copy()

    Xtr[num_cols] = Xtr[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")
    Xva[num_cols] = Xva[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")

```

```

Xte[num_cols] = Xte[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")

med = Xtr[num_cols].median()
Xtr[num_cols] = Xtr[num_cols].fillna(med)
Xva[num_cols] = Xva[num_cols].fillna(med)
Xte[num_cols] = Xte[num_cols].fillna(med)

for c in cat_cols:
    Xtr[c] = Xtr[c].astype("category")
    cats = Xtr[c].cat.categories.astype(str)
    Xva[c] = pd.Categorical(Xva[c].astype(str), categories=cats)
    Xte[c] = pd.Categorical(Xte[c].astype(str), categories=cats)
    Xtr[c] = Xtr[c].cat.codes.astype("int32")
    Xva[c] = Xva[c].cat.codes.astype("int32")
    Xte[c] = Xte[c].cat.codes.astype("int32")

Xtr_np = Xtr.to_numpy(dtype=np.float32)
Xva_np = Xva.to_numpy(dtype=np.float32)
Xte_np = Xte.to_numpy(dtype=np.float32)

ytr_log = np.log1p(train_df["Sales"].astype("float32").values)
yva_log = np.log1p(val_df["Sales"].astype("float32").values)
yva_norm = val_df["Sales"].astype("float32").values
yte_norm = test_df["Sales"].astype("float32").values
open_va = val_df["Open"].astype("int8").values
open_te = test_df["Open"].astype("int8").values

return Xtr_np, ytr_log, Xva_np, yva_log, yva_norm, open_va, Xte_np, yte_norm, open_te

class _Node:
    __slots__ = ("is_leaf", "value", "fidx", "thr", "left", "right")
    def __init__(self, is_leaf, value=None, fidx=None, thr=None, left=None, right=None):
        self.is_leaf = is_leaf
        self.value = value
        self.fidx = fidx
        self.thr = thr
        self.left = left
        self.right = right

class XGBoostScratch:
    def __init__(self, n_estimators=100, learning_rate=0.1, max_depth=4,
                 min_child_weight=10.0, reg_lambda=1.0, gamma=0.0,
                 subsample=0.8, colsample=0.8, max_bins=16, random_state=42):
        self.n_estimators = n_estimators
        self.learning_rate = learning_rate
        self.max_depth = max_depth

```

```

        self.min_child_weight = min_child_weight
        self.reg_lambda = reg_lambda
        self.gamma = gamma
        self.subsample = subsample
        self.colsample = colsample
        self.max_bins = max_bins
        self.rng = np.random.default_rng(random_state)
        self.trees = []
        self.base_score = 0.0

    def _leaf_weight(self, G, H):
        return -G / (H + self.reg_lambda)

    def _score(self, G, H):
        return (G * G) / (H + self.reg_lambda)

    def _best_split(self, X, g, h, idxs, feat_idxs):
        G, H = g[idxs].sum(), h[idxs].sum()
        parent_score = self._score(G, H)
        best_gain, best_f, best_thr = -1e30, None, None
        best_left, best_right = None, None
        for f in feat_idxs:
            x = X[idxs, f]
            if float(x.min()) == float(x.max()): continue
            qs = np.linspace(0.0, 1.0, self.max_bins + 2)[1:-1]
            thr_list = np.unique(np.quantile(x, qs))
            if thr_list.size == 0: continue
            for thr in thr_list:
                m = x <= thr
                if not m.any() or m.all(): continue
                l_idx, r_idx = idxs[m], idxs[~m]
                Hl, Hr = h[l_idx].sum(), h[r_idx].sum()
                if Hl < self.min_child_weight or Hr < self.min_child_weight: continue
                gain = self._score(g[l_idx].sum(), Hl) + self._score(g[r_idx].sum(), Hr) - parent_score
                if gain > best_gain:
                    best_gain, best_f, best_thr = gain, int(f), float(thr)
                    best_left, best_right = l_idx, r_idx
        return best_gain, best_f, best_thr, best_left, best_right

    def _build(self, X, g, h, idxs, depth, feat_idxs):
        G, H = g[idxs].sum(), h[idxs].sum()
        if depth >= self.max_depth or idxs.size <= 1 or H < self.min_child_weight:
            return _Node(True, value=float(self._leaf_weight(G, H)))
        gain, f, thr, left, right = self._best_split(X, g, h, idxs, feat_idxs)
        if f is None or gain <= 0:
            return _Node(True, value=float(self._leaf_weight(G, H)))

```

```

    return _Node(False, fidx=f, thr=thr,
                left=self._build(X, g, h, left, depth + 1, feat_idxs),
                right=self._build(X, g, h, right, depth + 1, feat_idxs))

def _pred_tree(self, node, X):
    out = np.empty(X.shape[0], dtype=np.float32)
    for i in range(X.shape[0]):
        curr = node
        while not curr.is_leaf:
            curr = curr.left if X[i, curr.fidx] <= curr.thr else curr.right
        out[i] = curr.value
    return out

def fit(self, X, y_log, X_val=None, y_val_log=None, early_stopping_rounds=10, print_every=1):
    n, d = X.shape
    self.base_score = float(y_log.mean())
    pred = np.full(n, self.base_score, dtype=np.float32)
    pred_val = np.full(X_val.shape[0], self.base_score, dtype=np.float32) if X_val is not None else None
    best_rmse, best_k, bad = 1e30, 0, 0
    for k in range(1, self.n_estimators + 1):
        rows = self.rng.choice(n, size=int(self.subsample * n), replace=False)
        feats = self.rng.choice(np.arange(d), size=int(self.colsample * d), replace=False)
        g, h = (pred - y_log), np.ones_like(pred)
        tree = self._build(X, g, h, rows, 0, feats)
        self.trees.append(tree)
        pred += self.learning_rate * self._pred_tree(tree, X)
        train_err = np.sqrt(np.mean((pred - y_log) ** 2))
        if X_val is not None:
            pred_val += self.learning_rate * self._pred_tree(tree, X_val)
            val_err = np.sqrt(np.mean((pred_val - y_val_log) ** 2))
            if k % print_every == 0:
                print(f"Iter {k:3d} | Train RMSE(log): {train_err:.5f} | Val RMSE(log): {val_err:.5f}")
            if val_err < best_rmse:
                best_rmse, best_k, bad = val_err, k, 0
            else:
                bad += 1
                if bad >= early_stopping_rounds:
                    print(f"Early stopping at {k}")
                    break
        if X_val is not None: self.trees = self.trees[:best_k]

def predict(self, X):
    p = np.full(X.shape[0], self.base_score, dtype=np.float32)
    for t in self.trees: p += self.learning_rate * self._pred_tree(t, X)
    return p

```

```

Xtr_np, ytr_log, Xva_np, yva_log, yva_norm, open_va, Xte_np, yte_norm, open_te = prepare_mat()

model = XGBoostScratch(n_estimators=100, learning_rate=0.1, max_depth=4)
model.fit(Xtr_np, ytr_log, X_val=Xva_np, y_val_log=yva_log, print_every=1)

s_min, s_max = mn["Sales"], mx["Sales"]
p_va_euro = np.expm1(model.predict(Xva_np)) * (s_max - s_min) + s_min
p_te_euro = np.expm1(model.predict(Xte_np)) * (s_max - s_min) + s_min
y_va_euro = yva_norm * (s_max - s_min) + s_min
y_te_euro = yte_norm * (s_max - s_min) + s_min

p_va_euro = np.where(open_va == 0, 0, p_va_euro)
p_te_euro = np.where(open_te == 0, 0, p_te_euro)

print(f"\nValidation RMSE: {rmse(y_va_euro, p_va_euro):.2f} Euro")
print(f"Test RMSE: {rmse(y_te_euro, p_te_euro):.2f} Euro")

Iter   1 | Train RMSE(log): 0.07269 | Val RMSE(log): 0.07630
Iter   2 | Train RMSE(log): 0.06681 | Val RMSE(log): 0.07012
Iter   3 | Train RMSE(log): 0.06159 | Val RMSE(log): 0.06459
Iter   4 | Train RMSE(log): 0.05715 | Val RMSE(log): 0.05987
Iter   5 | Train RMSE(log): 0.05450 | Val RMSE(log): 0.05713
Iter   6 | Train RMSE(log): 0.05067 | Val RMSE(log): 0.05311
Iter   7 | Train RMSE(log): 0.04756 | Val RMSE(log): 0.04980
Iter   8 | Train RMSE(log): 0.04454 | Val RMSE(log): 0.04663
Iter   9 | Train RMSE(log): 0.04188 | Val RMSE(log): 0.04386
Iter  10 | Train RMSE(log): 0.03955 | Val RMSE(log): 0.04145
Iter  11 | Train RMSE(log): 0.03737 | Val RMSE(log): 0.03924
Iter  12 | Train RMSE(log): 0.03562 | Val RMSE(log): 0.03749
Iter  13 | Train RMSE(log): 0.03413 | Val RMSE(log): 0.03595
Iter  14 | Train RMSE(log): 0.03260 | Val RMSE(log): 0.03442
Iter  15 | Train RMSE(log): 0.03132 | Val RMSE(log): 0.03319
Iter  16 | Train RMSE(log): 0.03019 | Val RMSE(log): 0.03211
Iter  17 | Train RMSE(log): 0.02920 | Val RMSE(log): 0.03116
Iter  18 | Train RMSE(log): 0.02838 | Val RMSE(log): 0.03036
Iter  19 | Train RMSE(log): 0.02759 | Val RMSE(log): 0.02965
Iter  20 | Train RMSE(log): 0.02688 | Val RMSE(log): 0.02897
Iter  21 | Train RMSE(log): 0.02622 | Val RMSE(log): 0.02838
Iter  22 | Train RMSE(log): 0.02567 | Val RMSE(log): 0.02788
Iter  23 | Train RMSE(log): 0.02523 | Val RMSE(log): 0.02746
Iter  24 | Train RMSE(log): 0.02482 | Val RMSE(log): 0.02709
Iter  25 | Train RMSE(log): 0.02440 | Val RMSE(log): 0.02673
Iter  26 | Train RMSE(log): 0.02406 | Val RMSE(log): 0.02643
Iter  27 | Train RMSE(log): 0.02372 | Val RMSE(log): 0.02616
Iter  28 | Train RMSE(log): 0.02344 | Val RMSE(log): 0.02573
Iter  29 | Train RMSE(log): 0.02306 | Val RMSE(log): 0.02547

```

Iter	30	Train RMSE(log): 0.02277	Val RMSE(log): 0.02493
Iter	31	Train RMSE(log): 0.02245	Val RMSE(log): 0.02470
Iter	32	Train RMSE(log): 0.02220	Val RMSE(log): 0.02445
Iter	33	Train RMSE(log): 0.02194	Val RMSE(log): 0.02422
Iter	34	Train RMSE(log): 0.02172	Val RMSE(log): 0.02371
Iter	35	Train RMSE(log): 0.02150	Val RMSE(log): 0.02359
Iter	36	Train RMSE(log): 0.02131	Val RMSE(log): 0.02345
Iter	37	Train RMSE(log): 0.02116	Val RMSE(log): 0.02321
Iter	38	Train RMSE(log): 0.02102	Val RMSE(log): 0.02310
Iter	39	Train RMSE(log): 0.02088	Val RMSE(log): 0.02289
Iter	40	Train RMSE(log): 0.02073	Val RMSE(log): 0.02278
Iter	41	Train RMSE(log): 0.02058	Val RMSE(log): 0.02235
Iter	42	Train RMSE(log): 0.02044	Val RMSE(log): 0.02229
Iter	43	Train RMSE(log): 0.02031	Val RMSE(log): 0.02206
Iter	44	Train RMSE(log): 0.02018	Val RMSE(log): 0.02192
Iter	45	Train RMSE(log): 0.02010	Val RMSE(log): 0.02186
Iter	46	Train RMSE(log): 0.01996	Val RMSE(log): 0.02180
Iter	47	Train RMSE(log): 0.01987	Val RMSE(log): 0.02176
Iter	48	Train RMSE(log): 0.01977	Val RMSE(log): 0.02171
Iter	49	Train RMSE(log): 0.01967	Val RMSE(log): 0.02158
Iter	50	Train RMSE(log): 0.01957	Val RMSE(log): 0.02150
Iter	51	Train RMSE(log): 0.01951	Val RMSE(log): 0.02148
Iter	52	Train RMSE(log): 0.01945	Val RMSE(log): 0.02144
Iter	53	Train RMSE(log): 0.01936	Val RMSE(log): 0.02141
Iter	54	Train RMSE(log): 0.01929	Val RMSE(log): 0.02137
Iter	55	Train RMSE(log): 0.01920	Val RMSE(log): 0.02131
Iter	56	Train RMSE(log): 0.01916	Val RMSE(log): 0.02130
Iter	57	Train RMSE(log): 0.01912	Val RMSE(log): 0.02129
Iter	58	Train RMSE(log): 0.01906	Val RMSE(log): 0.02126
Iter	59	Train RMSE(log): 0.01897	Val RMSE(log): 0.02121
Iter	60	Train RMSE(log): 0.01892	Val RMSE(log): 0.02118
Iter	61	Train RMSE(log): 0.01885	Val RMSE(log): 0.02111
Iter	62	Train RMSE(log): 0.01879	Val RMSE(log): 0.02107
Iter	63	Train RMSE(log): 0.01873	Val RMSE(log): 0.02106
Iter	64	Train RMSE(log): 0.01867	Val RMSE(log): 0.02103
Iter	65	Train RMSE(log): 0.01862	Val RMSE(log): 0.02093
Iter	66	Train RMSE(log): 0.01856	Val RMSE(log): 0.02087
Iter	67	Train RMSE(log): 0.01852	Val RMSE(log): 0.02082
Iter	68	Train RMSE(log): 0.01848	Val RMSE(log): 0.02078
Iter	69	Train RMSE(log): 0.01843	Val RMSE(log): 0.02067
Iter	70	Train RMSE(log): 0.01841	Val RMSE(log): 0.02067
Iter	71	Train RMSE(log): 0.01834	Val RMSE(log): 0.02061
Iter	72	Train RMSE(log): 0.01828	Val RMSE(log): 0.02057
Iter	73	Train RMSE(log): 0.01825	Val RMSE(log): 0.02059
Iter	74	Train RMSE(log): 0.01819	Val RMSE(log): 0.02055
Iter	75	Train RMSE(log): 0.01815	Val RMSE(log): 0.02055

```

Iter 76 | Train RMSE(log): 0.01812 | Val RMSE(log): 0.02053
Iter 77 | Train RMSE(log): 0.01807 | Val RMSE(log): 0.02052
Iter 78 | Train RMSE(log): 0.01805 | Val RMSE(log): 0.02052
Iter 79 | Train RMSE(log): 0.01801 | Val RMSE(log): 0.02044
Iter 80 | Train RMSE(log): 0.01796 | Val RMSE(log): 0.02042
Iter 81 | Train RMSE(log): 0.01793 | Val RMSE(log): 0.02038
Iter 82 | Train RMSE(log): 0.01790 | Val RMSE(log): 0.02038
Iter 83 | Train RMSE(log): 0.01786 | Val RMSE(log): 0.02036
Iter 84 | Train RMSE(log): 0.01784 | Val RMSE(log): 0.02037
Iter 85 | Train RMSE(log): 0.01779 | Val RMSE(log): 0.02030
Iter 86 | Train RMSE(log): 0.01776 | Val RMSE(log): 0.02030
Iter 87 | Train RMSE(log): 0.01773 | Val RMSE(log): 0.02028
Iter 88 | Train RMSE(log): 0.01771 | Val RMSE(log): 0.02025
Iter 89 | Train RMSE(log): 0.01767 | Val RMSE(log): 0.02022
Iter 90 | Train RMSE(log): 0.01763 | Val RMSE(log): 0.02015
Iter 91 | Train RMSE(log): 0.01760 | Val RMSE(log): 0.02013
Iter 92 | Train RMSE(log): 0.01757 | Val RMSE(log): 0.02012
Iter 93 | Train RMSE(log): 0.01754 | Val RMSE(log): 0.02009
Iter 94 | Train RMSE(log): 0.01752 | Val RMSE(log): 0.02010
Iter 95 | Train RMSE(log): 0.01750 | Val RMSE(log): 0.02010
Iter 96 | Train RMSE(log): 0.01746 | Val RMSE(log): 0.02006
Iter 97 | Train RMSE(log): 0.01743 | Val RMSE(log): 0.02005
Iter 98 | Train RMSE(log): 0.01740 | Val RMSE(log): 0.02003
Iter 99 | Train RMSE(log): 0.01736 | Val RMSE(log): 0.02000
Iter 100 | Train RMSE(log): 0.01735 | Val RMSE(log): 0.02000

```

Validation RMSE: 1008.17 Euro

Test RMSE: 910.21 Euro

```

import numpy as np
import pandas as pd

feature_cols = [
    "Store", "DayOfWeek", "Open", "Promo", "SchoolHoliday", "StateHoliday",
    "StoreType", "Assortment",
    "CompetitionDistance", "Promo2",
    "Year", "Month", "Day", "WeekOfYear", "DayOfYear", "IsWeekend", "IsMonthStart", "IsMonthEnd", "IsQuarterStart", "IsQuarterEnd",
    "SalesLag1", "SalesLag7", "SalesMean7", "SalesStd7", "SalesMean30", "SalesStd30",
    "sin_1y", "cos_1y", "sin_2y", "cos_2y",
    "DaysSinceHoliday", "DaysToNextHoliday"
]

cat_cols = ["Store", "DayOfWeek", "StateHoliday", "StoreType", "Assortment"]
cat_cols = [c for c in cat_cols if c in feature_cols]
num_cols = [c for c in feature_cols if c not in cat_cols]

```

```

def rmse(y_true, y_pred):
    e = y_true - y_pred
    return float(np.sqrt(np.mean(e * e)))

def mape_custom(y_true, y_pred):
    mask = y_true > 0
    return float(np.mean(np.abs((y_true[mask] - y_pred[mask]) / y_true[mask])) * 100)

def r2_custom(y_true, y_pred):
    ss_res = np.sum((y_true - y_pred) ** 2)
    ss_tot = np.sum((y_true - np.mean(y_true)) ** 2)
    return float(1 - (ss_res / ss_tot))

def prepare_matrices(train_df, val_df, test_df):
    Xtr = train_df[feature_cols].copy()
    Xva = val_df[feature_cols].copy()
    Xte = test_df[feature_cols].copy()

    Xtr[num_cols] = Xtr[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")
    Xva[num_cols] = Xva[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")
    Xte[num_cols] = Xte[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")

    med = Xtr[num_cols].median()
    Xtr[num_cols] = Xtr[num_cols].fillna(med)
    Xva[num_cols] = Xva[num_cols].fillna(med)
    Xte[num_cols] = Xte[num_cols].fillna(med)

    for c in cat_cols:
        Xtr[c] = Xtr[c].astype("category")
        cats = Xtr[c].cat.categories.astype(str)
        Xva[c] = pd.Categorical(Xva[c].astype(str), categories=cats)
        Xte[c] = pd.Categorical(Xte[c].astype(str), categories=cats)
        Xtr[c] = Xtr[c].cat.codes.astype("int32")
        Xva[c] = Xva[c].cat.codes.astype("int32")
        Xte[c] = Xte[c].cat.codes.astype("int32")

    Xtr_np = Xtr.to_numpy(dtype=np.float32)
    Xva_np = Xva.to_numpy(dtype=np.float32)
    Xte_np = Xte.to_numpy(dtype=np.float32)

    ytr_log = np.log1p(train_df["Sales"].astype("float32").values)
    yva_log = np.log1p(val_df["Sales"].astype("float32").values)
    yva_norm = val_df["Sales"].astype("float32").values
    yte_norm = test_df["Sales"].astype("float32").values
    open_va = val_df["Open"].astype("int8").values
    open_te = test_df["Open"].astype("int8").values

```

```

    return Xtr_np, ytr_log, Xva_np, yva_log, yva_norm, open_va, Xte_np, yte_norm, open_te

class _Node:
    __slots__ = ("is_leaf", "value", "fidx", "thr", "left", "right")
    def __init__(self, is_leaf, value=None, fidx=None, thr=None, left=None, right=None):
        self.is_leaf = is_leaf
        self.value = value
        self.fidx = fidx
        self.thr = thr
        self.left = left
        self.right = right

class XGBoostScratch:
    def __init__(self, n_estimators=100, learning_rate=0.1, max_depth=4,
                 min_child_weight=10.0, reg_lambda=1.0, gamma=0.0,
                 subsample=0.8, colsample=0.8, max_bins=16, random_state=42):
        self.n_estimators = n_estimators
        self.learning_rate = learning_rate
        self.max_depth = max_depth
        self.min_child_weight = min_child_weight
        self.reg_lambda = reg_lambda
        self.gamma = gamma
        self.subsample = subsample
        self.colsample = colsample
        self.max_bins = max_bins
        self.rng = np.random.default_rng(random_state)
        self.trees = []
        self.base_score = 0.0

    def _leaf_weight(self, G, H):
        return -G / (H + self.reg_lambda)

    def _score(self, G, H):
        return (G * G) / (H + self.reg_lambda)

    def _best_split(self, X, g, h, idxs, feat_idxs):
        G, H = g[idxs].sum(), h[idxs].sum()
        parent_score = self._score(G, H)
        best_gain, best_f, best_thr = -1e30, None, None
        best_left, best_right = None, None
        for f in feat_idxs:
            x = X[idxs, f]
            if float(x.min()) == float(x.max()): continue
            qs = np.linspace(0.0, 1.0, self.max_bins + 2)[1:-1]
            thr_list = np.unique(np.quantile(x, qs))

```

```

        if thr_list.size == 0: continue
    for thr in thr_list:
        m = x <= thr
        if not m.any() or m.all(): continue
        l_idx, r_idx = idxs[m], idxs[~m]
        Hl, Hr = h[l_idx].sum(), h[r_idx].sum()
        if Hl < self.min_child_weight or Hr < self.min_child_weight: continue
        gain = self._score(g[l_idx].sum(), Hl) + self._score(g[r_idx].sum(), Hr) - p
        if gain > best_gain:
            best_gain, best_f, best_thr = gain, int(f), float(thr)
            best_left, best_right = l_idx, r_idx
    return best_gain, best_f, best_thr, best_left, best_right

def _build(self, X, g, h, idxs, depth, feat_idxs):
    G, H = g[idxs].sum(), h[idxs].sum()
    if depth >= self.max_depth or idxs.size <= 1 or H < self.min_child_weight:
        return _Node(True, value=float(self._leaf_weight(G, H)))
    gain, f, thr, left, right = self._best_split(X, g, h, idxs, feat_idxs)
    if f is None or gain <= 0:
        return _Node(True, value=float(self._leaf_weight(G, H)))
    return _Node(False, fidx=f, thr=thr,
                left=self._build(X, g, h, left, depth + 1, feat_idxs),
                right=self._build(X, g, h, right, depth + 1, feat_idxs))

def _pred_tree(self, node, X):
    out = np.empty(X.shape[0], dtype=np.float32)
    for i in range(X.shape[0]):
        curr = node
        while not curr.is_leaf:
            curr = curr.left if X[i, curr.fidx] <= curr.thr else curr.right
        out[i] = curr.value
    return out

def fit(self, X, y_log, X_val=None, y_val_log=None, early_stopping_rounds=10, print_ever
n, d = X.shape
self.base_score = float(y_log.mean())
pred = np.full(n, self.base_score, dtype=np.float32)
pred_val = np.full(X_val.shape[0], self.base_score, dtype=np.float32) if X_val is no
best_rmse, best_k, bad = 1e30, 0, 0
for k in range(1, self.n_estimators + 1):
    rows = self.rng.choice(n, size=int(self.subsample * n), replace=False)
    feats = self.rng.choice(np.arange(d), size=int(self.colsample * d), replace=False)
    g, h = (pred - y_log), np.ones_like(pred)
    tree = self._build(X, g, h, rows, 0, feats)
    self.trees.append(tree)
    pred += self.learning_rate * self._pred_tree(tree, X)

```

```

train_err = np.sqrt(np.mean((pred - y_log) ** 2))
if X_val is not None:
    pred_val += self.learning_rate * self._pred_tree(tree, X_val)
    val_err = np.sqrt(np.mean((pred_val - y_val_log) ** 2))
    if k % print_every == 0:
        print(f"Iter {k:3d} | Train RMSE(log): {train_err:.5f} | Val RMSE(log): {val_err:.5f}")
    if val_err < best_rmse:
        best_rmse, best_k, bad = val_err, k, 0
    else:
        bad += 1
        if bad >= early_stopping_rounds:
            print(f"Early stopping at {k}")
            break
if X_val is not None: self.trees = self.trees[:best_k]

def predict(self, X):
    p = np.full(X.shape[0], self.base_score, dtype=np.float32)
    for t in self.trees: p += self.learning_rate * self._pred_tree(t, X)
    return p

Xtr_np, ytr_log, Xva_np, yva_log, yva_norm, open_va, Xte_np, yte_norm, open_te = prepare_mat

model = XGBoostScratch(n_estimators=100, learning_rate=0.1, max_depth=4)
model.fit(Xtr_np, ytr_log, X_val=Xva_np, y_val_log=yva_log, print_every=1)

s_min, s_max = mn["Sales"], mx["Sales"]
s_range = s_max - s_min

p_va_euro = np.expm1(model.predict(Xva_np)) * s_range + s_min
y_va_euro = yva_norm * s_range + s_min
p_va_euro = np.where(open_va == 0, 0, p_va_euro)

p_te_euro = np.expm1(model.predict(Xte_np)) * s_range + s_min
y_te_euro = yte_norm * s_range + s_min
p_te_euro = np.where(open_te == 0, 0, p_te_euro)

print("\n" + "*40)
print("FINAL PERFORMANCE METRICS")
print("*40)
print(f"Validation:")
print(f"  RMSE: {rmse(y_va_euro, p_va_euro):.2f} Euro")
print(f"  MAPE: {mape_custom(y_va_euro, p_va_euro):.2f}%")
print(f"  R2 : {r2_custom(y_va_euro, p_va_euro):.4f}")
print("-" * 40)
print(f"Test:")
print(f"  RMSE: {rmse(y_te_euro, p_te_euro):.2f} Euro")

```

```

print(f"  MAPE: {mape_custom(y_te_euro, p_te_euro):.2f}%)")
print(f"  R2 : {r2_custom(y_te_euro, p_te_euro):.4f}"))
print("=*40)

Iter  1 | Train RMSE(log): 0.07269 | Val RMSE(log): 0.07630
Iter  2 | Train RMSE(log): 0.06681 | Val RMSE(log): 0.07012
Iter  3 | Train RMSE(log): 0.06159 | Val RMSE(log): 0.06459
Iter  4 | Train RMSE(log): 0.05715 | Val RMSE(log): 0.05987
Iter  5 | Train RMSE(log): 0.05450 | Val RMSE(log): 0.05713
Iter  6 | Train RMSE(log): 0.05067 | Val RMSE(log): 0.05311
Iter  7 | Train RMSE(log): 0.04756 | Val RMSE(log): 0.04980
Iter  8 | Train RMSE(log): 0.04454 | Val RMSE(log): 0.04663
Iter  9 | Train RMSE(log): 0.04188 | Val RMSE(log): 0.04386
Iter 10 | Train RMSE(log): 0.03955 | Val RMSE(log): 0.04145
Iter 11 | Train RMSE(log): 0.03737 | Val RMSE(log): 0.03924
Iter 12 | Train RMSE(log): 0.03562 | Val RMSE(log): 0.03749
Iter 13 | Train RMSE(log): 0.03413 | Val RMSE(log): 0.03595
Iter 14 | Train RMSE(log): 0.03260 | Val RMSE(log): 0.03442
Iter 15 | Train RMSE(log): 0.03132 | Val RMSE(log): 0.03319
Iter 16 | Train RMSE(log): 0.03019 | Val RMSE(log): 0.03211
Iter 17 | Train RMSE(log): 0.02920 | Val RMSE(log): 0.03116
Iter 18 | Train RMSE(log): 0.02838 | Val RMSE(log): 0.03036
Iter 19 | Train RMSE(log): 0.02759 | Val RMSE(log): 0.02965
Iter 20 | Train RMSE(log): 0.02688 | Val RMSE(log): 0.02897
Iter 21 | Train RMSE(log): 0.02622 | Val RMSE(log): 0.02838
Iter 22 | Train RMSE(log): 0.02567 | Val RMSE(log): 0.02788
Iter 23 | Train RMSE(log): 0.02523 | Val RMSE(log): 0.02746
Iter 24 | Train RMSE(log): 0.02482 | Val RMSE(log): 0.02709
Iter 25 | Train RMSE(log): 0.02440 | Val RMSE(log): 0.02673
Iter 26 | Train RMSE(log): 0.02406 | Val RMSE(log): 0.02643
Iter 27 | Train RMSE(log): 0.02372 | Val RMSE(log): 0.02616
Iter 28 | Train RMSE(log): 0.02344 | Val RMSE(log): 0.02573
Iter 29 | Train RMSE(log): 0.02306 | Val RMSE(log): 0.02547
Iter 30 | Train RMSE(log): 0.02277 | Val RMSE(log): 0.02493
Iter 31 | Train RMSE(log): 0.02245 | Val RMSE(log): 0.02470
Iter 32 | Train RMSE(log): 0.02220 | Val RMSE(log): 0.02445
Iter 33 | Train RMSE(log): 0.02194 | Val RMSE(log): 0.02422
Iter 34 | Train RMSE(log): 0.02172 | Val RMSE(log): 0.02371
Iter 35 | Train RMSE(log): 0.02150 | Val RMSE(log): 0.02359
Iter 36 | Train RMSE(log): 0.02131 | Val RMSE(log): 0.02345
Iter 37 | Train RMSE(log): 0.02116 | Val RMSE(log): 0.02321
Iter 38 | Train RMSE(log): 0.02102 | Val RMSE(log): 0.02310
Iter 39 | Train RMSE(log): 0.02088 | Val RMSE(log): 0.02289
Iter 40 | Train RMSE(log): 0.02073 | Val RMSE(log): 0.02278
Iter 41 | Train RMSE(log): 0.02058 | Val RMSE(log): 0.02235
Iter 42 | Train RMSE(log): 0.02044 | Val RMSE(log): 0.02229

```

Iter	43	Train RMSE(log): 0.02031	Val RMSE(log): 0.02206
Iter	44	Train RMSE(log): 0.02018	Val RMSE(log): 0.02192
Iter	45	Train RMSE(log): 0.02010	Val RMSE(log): 0.02186
Iter	46	Train RMSE(log): 0.01996	Val RMSE(log): 0.02180
Iter	47	Train RMSE(log): 0.01987	Val RMSE(log): 0.02176
Iter	48	Train RMSE(log): 0.01977	Val RMSE(log): 0.02171
Iter	49	Train RMSE(log): 0.01967	Val RMSE(log): 0.02158
Iter	50	Train RMSE(log): 0.01957	Val RMSE(log): 0.02150
Iter	51	Train RMSE(log): 0.01951	Val RMSE(log): 0.02148
Iter	52	Train RMSE(log): 0.01945	Val RMSE(log): 0.02144
Iter	53	Train RMSE(log): 0.01936	Val RMSE(log): 0.02141
Iter	54	Train RMSE(log): 0.01929	Val RMSE(log): 0.02137
Iter	55	Train RMSE(log): 0.01920	Val RMSE(log): 0.02131
Iter	56	Train RMSE(log): 0.01916	Val RMSE(log): 0.02130
Iter	57	Train RMSE(log): 0.01912	Val RMSE(log): 0.02129
Iter	58	Train RMSE(log): 0.01906	Val RMSE(log): 0.02126
Iter	59	Train RMSE(log): 0.01897	Val RMSE(log): 0.02121
Iter	60	Train RMSE(log): 0.01892	Val RMSE(log): 0.02118
Iter	61	Train RMSE(log): 0.01885	Val RMSE(log): 0.02111
Iter	62	Train RMSE(log): 0.01879	Val RMSE(log): 0.02107
Iter	63	Train RMSE(log): 0.01873	Val RMSE(log): 0.02106
Iter	64	Train RMSE(log): 0.01867	Val RMSE(log): 0.02103
Iter	65	Train RMSE(log): 0.01862	Val RMSE(log): 0.02093
Iter	66	Train RMSE(log): 0.01856	Val RMSE(log): 0.02087
Iter	67	Train RMSE(log): 0.01852	Val RMSE(log): 0.02082
Iter	68	Train RMSE(log): 0.01848	Val RMSE(log): 0.02078
Iter	69	Train RMSE(log): 0.01843	Val RMSE(log): 0.02067
Iter	70	Train RMSE(log): 0.01841	Val RMSE(log): 0.02067
Iter	71	Train RMSE(log): 0.01834	Val RMSE(log): 0.02061
Iter	72	Train RMSE(log): 0.01828	Val RMSE(log): 0.02057
Iter	73	Train RMSE(log): 0.01825	Val RMSE(log): 0.02059
Iter	74	Train RMSE(log): 0.01819	Val RMSE(log): 0.02055
Iter	75	Train RMSE(log): 0.01815	Val RMSE(log): 0.02055
Iter	76	Train RMSE(log): 0.01812	Val RMSE(log): 0.02053
Iter	77	Train RMSE(log): 0.01807	Val RMSE(log): 0.02052
Iter	78	Train RMSE(log): 0.01805	Val RMSE(log): 0.02052
Iter	79	Train RMSE(log): 0.01801	Val RMSE(log): 0.02044
Iter	80	Train RMSE(log): 0.01796	Val RMSE(log): 0.02042
Iter	81	Train RMSE(log): 0.01793	Val RMSE(log): 0.02038
Iter	82	Train RMSE(log): 0.01790	Val RMSE(log): 0.02038
Iter	83	Train RMSE(log): 0.01786	Val RMSE(log): 0.02036
Iter	84	Train RMSE(log): 0.01784	Val RMSE(log): 0.02037
Iter	85	Train RMSE(log): 0.01779	Val RMSE(log): 0.02030
Iter	86	Train RMSE(log): 0.01776	Val RMSE(log): 0.02030
Iter	87	Train RMSE(log): 0.01773	Val RMSE(log): 0.02028
Iter	88	Train RMSE(log): 0.01771	Val RMSE(log): 0.02025

```

Iter 89 | Train RMSE(log): 0.01767 | Val RMSE(log): 0.02022
Iter 90 | Train RMSE(log): 0.01763 | Val RMSE(log): 0.02015
Iter 91 | Train RMSE(log): 0.01760 | Val RMSE(log): 0.02013
Iter 92 | Train RMSE(log): 0.01757 | Val RMSE(log): 0.02012
Iter 93 | Train RMSE(log): 0.01754 | Val RMSE(log): 0.02009
Iter 94 | Train RMSE(log): 0.01752 | Val RMSE(log): 0.02010
Iter 95 | Train RMSE(log): 0.01750 | Val RMSE(log): 0.02010
Iter 96 | Train RMSE(log): 0.01746 | Val RMSE(log): 0.02006
Iter 97 | Train RMSE(log): 0.01743 | Val RMSE(log): 0.02005
Iter 98 | Train RMSE(log): 0.01740 | Val RMSE(log): 0.02003
Iter 99 | Train RMSE(log): 0.01736 | Val RMSE(log): 0.02000
Iter 100 | Train RMSE(log): 0.01735 | Val RMSE(log): 0.02000

=====
FINAL PERFORMANCE METRICS
=====

Validation:
RMSE: 1008.17 Euro
MAPE: 12.01%
R2 : 0.9362
-----

Test:
RMSE: 910.21 Euro
MAPE: 10.49%
R2 : 0.9468
=====
```

We applied a  $\log_{10}(\text{Sales})$  transform to reduce the strong right-skew of sales values and stabilize training under a squared-error loss. Using  $\log_{10}$  also safely handles zero sales (since  $\log(1+0)$  is defined) and typically improves robustness to large outliers.

```

import numpy as np
import pandas as pd

s_min = mn["Sales"]
s_max = mx["Sales"]
s_range = s_max - s_min

def rmse(y_true, y_pred):
    e = y_true - y_pred
    return float(np.sqrt(np.mean(e*e)))

def to_xy(df_tr, df_va):
    Xtr = df_tr[feature_cols].copy()
    Xva = df_va[feature_cols].copy()
    Xtr[num_cols] = Xtr[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")
```

```

Xva[num_cols] = Xva[num_cols].apply(pd.to_numeric, errors="coerce").astype("float32")
med = Xtr[num_cols].median().fillna(0)
Xtr[num_cols] = Xtr[num_cols].fillna(med)
Xva[num_cols] = Xva[num_cols].fillna(med)
for c in cat_cols:
    Xtr[c] = Xtr[c].astype("category")
    cats = Xtr[c].cat.categories.astype(str)
    Xva[c] = pd.Categorical(Xva[c].astype(str), categories=cats)
    Xtr[c] = Xtr[c].cat.codes.astype("int32")
    Xva[c] = Xva[c].cat.codes.astype("int32")
Xtr_np = np.nan_to_num(Xtr.to_numpy(dtype=np.float32), nan=0.0, posinf=0.0, neginf=0.0)
Xva_np = np.nan_to_num(Xva.to_numpy(dtype=np.float32), nan=0.0, posinf=0.0, neginf=0.0)
ytr = df_tr["Sales"].astype("float32").values
yva = df_va["Sales"].astype("float32").values
return Xtr_np, ytr, Xva_np, yva

df_cv = df_mm.sort_values(["Date", "Store"]).reset_index(drop=True)
dates = np.array(sorted(df_cv["Date"].unique()))
n = len(dates)

fold_edges = []
train_start = int(0.55 * n)
for i in range(5):
    a = train_start + i * int(0.07 * n)
    b = train_start + (i+1) * int(0.07 * n)
    if b > n: b = n
    fold_edges.append((a, b))

cv_rmse_euro = []

for k, (a, b) in enumerate(fold_edges, 1):
    print(f"\n>>> FOLD {k} Training <<<")
    df_tr = df_cv[df_cv["Date"].isin(dates[:a])].reset_index(drop=True)
    df_va = df_cv[df_cv["Date"].isin(dates[a:b])].reset_index(drop=True)

    Xtr_np, ytr_norm, Xva_np, yva_norm = to_xy(df_tr, df_va)

    ytr_log = np.log1p(ytr_norm)
    yva_log = np.log1p(yva_norm)

    m = XGBoostScratch(
        n_estimators=60,
        learning_rate=0.1,
        max_depth=4,
        min_child_weight=10.0,
        reg_lambda=1.0,

```

```

        gamma=0.0,
        subsample=0.8,
        colsample=0.8,
        max_bins=16,
        random_state=42
    )

m.fit(Xtr_np, ytr_log, X_val=Xva_np, y_val_log=yva_log, print_every=1)

pred_norm = np.expm1(m.predict(Xva_np))
pred_euro = pred_norm * s_range + s_min
pred_euro = np.where(df_va["Open"].values == 0, 0.0, pred_euro)

yva_euro = yva_norm * s_range + s_min

fold_rmse = rmse(yva_euro, pred_euro)
cv_rmse_euro.append(fold_rmse)

print(f"\nFold {k} Summary: Val RMSE = {fold_rmse:.2f} Euro")

print("\n" + "="*40)
print("OVERALL CV RESULT (Euro):")
print(f"Mean RMSE: {float(np.mean(cv_rmse_euro)):.2f}")
print(f"Std RMSE : {float(np.std(cv_rmse_euro)):.2f}")
print("="*40)

>>> FOLD 1 Training <<<
Iter 1 | Train RMSE(log): 0.07278 | Val RMSE(log): 0.07323
Iter 2 | Train RMSE(log): 0.06722 | Val RMSE(log): 0.06760
Iter 3 | Train RMSE(log): 0.06251 | Val RMSE(log): 0.06284
Iter 4 | Train RMSE(log): 0.05817 | Val RMSE(log): 0.05840
Iter 5 | Train RMSE(log): 0.05394 | Val RMSE(log): 0.05387
Iter 6 | Train RMSE(log): 0.05064 | Val RMSE(log): 0.05037
Iter 7 | Train RMSE(log): 0.04757 | Val RMSE(log): 0.04721
Iter 8 | Train RMSE(log): 0.04455 | Val RMSE(log): 0.04397
Iter 9 | Train RMSE(log): 0.04188 | Val RMSE(log): 0.04109
Iter 10 | Train RMSE(log): 0.03962 | Val RMSE(log): 0.03861
Iter 11 | Train RMSE(log): 0.03767 | Val RMSE(log): 0.03660
Iter 12 | Train RMSE(log): 0.03620 | Val RMSE(log): 0.03495
Iter 13 | Train RMSE(log): 0.03449 | Val RMSE(log): 0.03309
Iter 14 | Train RMSE(log): 0.03292 | Val RMSE(log): 0.03141
Iter 15 | Train RMSE(log): 0.03166 | Val RMSE(log): 0.03003
Iter 16 | Train RMSE(log): 0.03058 | Val RMSE(log): 0.02891
Iter 17 | Train RMSE(log): 0.02973 | Val RMSE(log): 0.02796
Iter 18 | Train RMSE(log): 0.02886 | Val RMSE(log): 0.02695

```

Iter	19	Train RMSE(log): 0.02803	Val RMSE(log): 0.02600
Iter	20	Train RMSE(log): 0.02734	Val RMSE(log): 0.02521
Iter	21	Train RMSE(log): 0.02663	Val RMSE(log): 0.02448
Iter	22	Train RMSE(log): 0.02606	Val RMSE(log): 0.02388
Iter	23	Train RMSE(log): 0.02553	Val RMSE(log): 0.02333
Iter	24	Train RMSE(log): 0.02512	Val RMSE(log): 0.02284
Iter	25	Train RMSE(log): 0.02458	Val RMSE(log): 0.02237
Iter	26	Train RMSE(log): 0.02422	Val RMSE(log): 0.02197
Iter	27	Train RMSE(log): 0.02384	Val RMSE(log): 0.02164
Iter	28	Train RMSE(log): 0.02355	Val RMSE(log): 0.02134
Iter	29	Train RMSE(log): 0.02327	Val RMSE(log): 0.02114
Iter	30	Train RMSE(log): 0.02295	Val RMSE(log): 0.02094
Iter	31	Train RMSE(log): 0.02267	Val RMSE(log): 0.02074
Iter	32	Train RMSE(log): 0.02245	Val RMSE(log): 0.02054
Iter	33	Train RMSE(log): 0.02225	Val RMSE(log): 0.02036
Iter	34	Train RMSE(log): 0.02203	Val RMSE(log): 0.02017
Iter	35	Train RMSE(log): 0.02178	Val RMSE(log): 0.02009
Iter	36	Train RMSE(log): 0.02162	Val RMSE(log): 0.02004
Iter	37	Train RMSE(log): 0.02140	Val RMSE(log): 0.01990
Iter	38	Train RMSE(log): 0.02122	Val RMSE(log): 0.01984
Iter	39	Train RMSE(log): 0.02106	Val RMSE(log): 0.01969
Iter	40	Train RMSE(log): 0.02088	Val RMSE(log): 0.01970
Iter	41	Train RMSE(log): 0.02075	Val RMSE(log): 0.01964
Iter	42	Train RMSE(log): 0.02062	Val RMSE(log): 0.01953
Iter	43	Train RMSE(log): 0.02049	Val RMSE(log): 0.01945
Iter	44	Train RMSE(log): 0.02035	Val RMSE(log): 0.01940
Iter	45	Train RMSE(log): 0.02023	Val RMSE(log): 0.01931
Iter	46	Train RMSE(log): 0.02013	Val RMSE(log): 0.01925
Iter	47	Train RMSE(log): 0.02002	Val RMSE(log): 0.01923
Iter	48	Train RMSE(log): 0.01993	Val RMSE(log): 0.01924
Iter	49	Train RMSE(log): 0.01985	Val RMSE(log): 0.01921
Iter	50	Train RMSE(log): 0.01976	Val RMSE(log): 0.01914
Iter	51	Train RMSE(log): 0.01969	Val RMSE(log): 0.01911
Iter	52	Train RMSE(log): 0.01958	Val RMSE(log): 0.01901
Iter	53	Train RMSE(log): 0.01947	Val RMSE(log): 0.01898
Iter	54	Train RMSE(log): 0.01940	Val RMSE(log): 0.01890
Iter	55	Train RMSE(log): 0.01932	Val RMSE(log): 0.01884
Iter	56	Train RMSE(log): 0.01923	Val RMSE(log): 0.01881
Iter	57	Train RMSE(log): 0.01914	Val RMSE(log): 0.01875
Iter	58	Train RMSE(log): 0.01906	Val RMSE(log): 0.01883
Iter	59	Train RMSE(log): 0.01900	Val RMSE(log): 0.01882
Iter	60	Train RMSE(log): 0.01895	Val RMSE(log): 0.01876

Fold 1 Summary: Val RMSE = 961.25 Euro

>>> FOLD 2 Training <<<

Iter	1	Train RMSE(log): 0.07283	Val RMSE(log): 0.06982
Iter	2	Train RMSE(log): 0.06793	Val RMSE(log): 0.06475
Iter	3	Train RMSE(log): 0.06260	Val RMSE(log): 0.05940
Iter	4	Train RMSE(log): 0.05837	Val RMSE(log): 0.05522
Iter	5	Train RMSE(log): 0.05439	Val RMSE(log): 0.05133
Iter	6	Train RMSE(log): 0.05060	Val RMSE(log): 0.04749
Iter	7	Train RMSE(log): 0.04726	Val RMSE(log): 0.04418
Iter	8	Train RMSE(log): 0.04436	Val RMSE(log): 0.04125
Iter	9	Train RMSE(log): 0.04183	Val RMSE(log): 0.03885
Iter	10	Train RMSE(log): 0.03964	Val RMSE(log): 0.03677
Iter	11	Train RMSE(log): 0.03776	Val RMSE(log): 0.03484
Iter	12	Train RMSE(log): 0.03591	Val RMSE(log): 0.03306
Iter	13	Train RMSE(log): 0.03437	Val RMSE(log): 0.03161
Iter	14	Train RMSE(log): 0.03285	Val RMSE(log): 0.03012
Iter	15	Train RMSE(log): 0.03171	Val RMSE(log): 0.02895
Iter	16	Train RMSE(log): 0.03057	Val RMSE(log): 0.02786
Iter	17	Train RMSE(log): 0.02960	Val RMSE(log): 0.02697
Iter	18	Train RMSE(log): 0.02873	Val RMSE(log): 0.02615
Iter	19	Train RMSE(log): 0.02801	Val RMSE(log): 0.02549
Iter	20	Train RMSE(log): 0.02735	Val RMSE(log): 0.02485
Iter	21	Train RMSE(log): 0.02674	Val RMSE(log): 0.02426
Iter	22	Train RMSE(log): 0.02623	Val RMSE(log): 0.02378
Iter	23	Train RMSE(log): 0.02574	Val RMSE(log): 0.02337
Iter	24	Train RMSE(log): 0.02532	Val RMSE(log): 0.02301
Iter	25	Train RMSE(log): 0.02494	Val RMSE(log): 0.02270
Iter	26	Train RMSE(log): 0.02450	Val RMSE(log): 0.02226
Iter	27	Train RMSE(log): 0.02419	Val RMSE(log): 0.02198
Iter	28	Train RMSE(log): 0.02388	Val RMSE(log): 0.02175
Iter	29	Train RMSE(log): 0.02348	Val RMSE(log): 0.02136
Iter	30	Train RMSE(log): 0.02323	Val RMSE(log): 0.02120
Iter	31	Train RMSE(log): 0.02291	Val RMSE(log): 0.02090
Iter	32	Train RMSE(log): 0.02261	Val RMSE(log): 0.02061
Iter	33	Train RMSE(log): 0.02239	Val RMSE(log): 0.02048
Iter	34	Train RMSE(log): 0.02214	Val RMSE(log): 0.02039
Iter	35	Train RMSE(log): 0.02198	Val RMSE(log): 0.02027
Iter	36	Train RMSE(log): 0.02181	Val RMSE(log): 0.02014
Iter	37	Train RMSE(log): 0.02164	Val RMSE(log): 0.02001
Iter	38	Train RMSE(log): 0.02149	Val RMSE(log): 0.01987
Iter	39	Train RMSE(log): 0.02138	Val RMSE(log): 0.01980
Iter	40	Train RMSE(log): 0.02115	Val RMSE(log): 0.01972
Iter	41	Train RMSE(log): 0.02105	Val RMSE(log): 0.01965
Iter	42	Train RMSE(log): 0.02085	Val RMSE(log): 0.01950
Iter	43	Train RMSE(log): 0.02070	Val RMSE(log): 0.01934
Iter	44	Train RMSE(log): 0.02060	Val RMSE(log): 0.01927
Iter	45	Train RMSE(log): 0.02047	Val RMSE(log): 0.01927
Iter	46	Train RMSE(log): 0.02037	Val RMSE(log): 0.01924

```

Iter 47 | Train RMSE(log): 0.02025 | Val RMSE(log): 0.01914
Iter 48 | Train RMSE(log): 0.02017 | Val RMSE(log): 0.01910
Iter 49 | Train RMSE(log): 0.02009 | Val RMSE(log): 0.01899
Iter 50 | Train RMSE(log): 0.02002 | Val RMSE(log): 0.01895
Iter 51 | Train RMSE(log): 0.01995 | Val RMSE(log): 0.01890
Iter 52 | Train RMSE(log): 0.01985 | Val RMSE(log): 0.01884
Iter 53 | Train RMSE(log): 0.01974 | Val RMSE(log): 0.01874
Iter 54 | Train RMSE(log): 0.01969 | Val RMSE(log): 0.01870
Iter 55 | Train RMSE(log): 0.01958 | Val RMSE(log): 0.01859
Iter 56 | Train RMSE(log): 0.01948 | Val RMSE(log): 0.01855
Iter 57 | Train RMSE(log): 0.01943 | Val RMSE(log): 0.01858
Iter 58 | Train RMSE(log): 0.01936 | Val RMSE(log): 0.01855
Iter 59 | Train RMSE(log): 0.01933 | Val RMSE(log): 0.01853
Iter 60 | Train RMSE(log): 0.01922 | Val RMSE(log): 0.01842

```

Fold 2 Summary: Val RMSE = 931.46 Euro

```

>>> FOLD 3 Training <<<
Iter 1 | Train RMSE(log): 0.07260 | Val RMSE(log): 0.07437
Iter 2 | Train RMSE(log): 0.06672 | Val RMSE(log): 0.06822
Iter 3 | Train RMSE(log): 0.06209 | Val RMSE(log): 0.06329
Iter 4 | Train RMSE(log): 0.05746 | Val RMSE(log): 0.05837
Iter 5 | Train RMSE(log): 0.05387 | Val RMSE(log): 0.05455
Iter 6 | Train RMSE(log): 0.05017 | Val RMSE(log): 0.05074
Iter 7 | Train RMSE(log): 0.04687 | Val RMSE(log): 0.04733
Iter 8 | Train RMSE(log): 0.04396 | Val RMSE(log): 0.04422
Iter 9 | Train RMSE(log): 0.04144 | Val RMSE(log): 0.04163
Iter 10 | Train RMSE(log): 0.03925 | Val RMSE(log): 0.03937
Iter 11 | Train RMSE(log): 0.03723 | Val RMSE(log): 0.03721
Iter 12 | Train RMSE(log): 0.03551 | Val RMSE(log): 0.03538
Iter 13 | Train RMSE(log): 0.03389 | Val RMSE(log): 0.03370
Iter 14 | Train RMSE(log): 0.03251 | Val RMSE(log): 0.03224
Iter 15 | Train RMSE(log): 0.03116 | Val RMSE(log): 0.03074
Iter 16 | Train RMSE(log): 0.03010 | Val RMSE(log): 0.02953
Iter 17 | Train RMSE(log): 0.02906 | Val RMSE(log): 0.02844
Iter 18 | Train RMSE(log): 0.02825 | Val RMSE(log): 0.02757
Iter 19 | Train RMSE(log): 0.02756 | Val RMSE(log): 0.02679
Iter 20 | Train RMSE(log): 0.02685 | Val RMSE(log): 0.02604
Iter 21 | Train RMSE(log): 0.02625 | Val RMSE(log): 0.02545
Iter 22 | Train RMSE(log): 0.02576 | Val RMSE(log): 0.02497
Iter 23 | Train RMSE(log): 0.02514 | Val RMSE(log): 0.02433
Iter 24 | Train RMSE(log): 0.02476 | Val RMSE(log): 0.02396
Iter 25 | Train RMSE(log): 0.02436 | Val RMSE(log): 0.02354
Iter 26 | Train RMSE(log): 0.02387 | Val RMSE(log): 0.02305
Iter 27 | Train RMSE(log): 0.02349 | Val RMSE(log): 0.02266
Iter 28 | Train RMSE(log): 0.02320 | Val RMSE(log): 0.02237

```

```

Iter 29 | Train RMSE(log): 0.02294 | Val RMSE(log): 0.02201
Iter 30 | Train RMSE(log): 0.02263 | Val RMSE(log): 0.02173
Iter 31 | Train RMSE(log): 0.02231 | Val RMSE(log): 0.02134
Iter 32 | Train RMSE(log): 0.02204 | Val RMSE(log): 0.02111
Iter 33 | Train RMSE(log): 0.02179 | Val RMSE(log): 0.02087
Iter 34 | Train RMSE(log): 0.02159 | Val RMSE(log): 0.02063
Iter 35 | Train RMSE(log): 0.02141 | Val RMSE(log): 0.02050
Iter 36 | Train RMSE(log): 0.02123 | Val RMSE(log): 0.02036
Iter 37 | Train RMSE(log): 0.02106 | Val RMSE(log): 0.02019
Iter 38 | Train RMSE(log): 0.02088 | Val RMSE(log): 0.02009
Iter 39 | Train RMSE(log): 0.02075 | Val RMSE(log): 0.01995
Iter 40 | Train RMSE(log): 0.02061 | Val RMSE(log): 0.01978
Iter 41 | Train RMSE(log): 0.02048 | Val RMSE(log): 0.01971
Iter 42 | Train RMSE(log): 0.02035 | Val RMSE(log): 0.01958
Iter 43 | Train RMSE(log): 0.02021 | Val RMSE(log): 0.01947
Iter 44 | Train RMSE(log): 0.02012 | Val RMSE(log): 0.01934
Iter 45 | Train RMSE(log): 0.02000 | Val RMSE(log): 0.01931
Iter 48 | Train RMSE(log): 0.01970 | Val RMSE(log): 0.01912
Iter 49 | Train RMSE(log): 0.01963 | Val RMSE(log): 0.01906
Iter 50 | Train RMSE(log): 0.01957 | Val RMSE(log): 0.01898
Iter 51 | Train RMSE(log): 0.01946 | Val RMSE(log): 0.01893
Iter 52 | Train RMSE(log): 0.01939 | Val RMSE(log): 0.01887
Iter 53 | Train RMSE(log): 0.01931 | Val RMSE(log): 0.01885
Iter 54 | Train RMSE(log): 0.01925 | Val RMSE(log): 0.01884
Iter 55 | Train RMSE(log): 0.01915 | Val RMSE(log): 0.01876
Iter 56 | Train RMSE(log): 0.01908 | Val RMSE(log): 0.01877
Iter 57 | Train RMSE(log): 0.01899 | Val RMSE(log): 0.01873
Iter 58 | Train RMSE(log): 0.01893 | Val RMSE(log): 0.01872
Iter 59 | Train RMSE(log): 0.01889 | Val RMSE(log): 0.01872
Iter 60 | Train RMSE(log): 0.01883 | Val RMSE(log): 0.01866

```

Fold 3 Summary: Val RMSE = 945.82 Euro

```

>>> FOLD 4 Training <<<
Iter 1 | Train RMSE(log): 0.07301 | Val RMSE(log): 0.07984
Iter 2 | Train RMSE(log): 0.06705 | Val RMSE(log): 0.07346
Iter 3 | Train RMSE(log): 0.06204 | Val RMSE(log): 0.06804
Iter 4 | Train RMSE(log): 0.05733 | Val RMSE(log): 0.06300
Iter 5 | Train RMSE(log): 0.05334 | Val RMSE(log): 0.05876
Iter 6 | Train RMSE(log): 0.04977 | Val RMSE(log): 0.05480
Iter 7 | Train RMSE(log): 0.04645 | Val RMSE(log): 0.05133
Iter 8 | Train RMSE(log): 0.04389 | Val RMSE(log): 0.04885
Iter 9 | Train RMSE(log): 0.04127 | Val RMSE(log): 0.04617
Iter 10 | Train RMSE(log): 0.03914 | Val RMSE(log): 0.04423
Iter 11 | Train RMSE(log): 0.03712 | Val RMSE(log): 0.04220
Iter 12 | Train RMSE(log): 0.03536 | Val RMSE(log): 0.04034

```

Iter	13	Train RMSE(log): 0.03394	Val RMSE(log): 0.03888
Iter	14	Train RMSE(log): 0.03255	Val RMSE(log): 0.03738
Iter	15	Train RMSE(log): 0.03144	Val RMSE(log): 0.03639
Iter	16	Train RMSE(log): 0.03038	Val RMSE(log): 0.03539
Iter	17	Train RMSE(log): 0.02940	Val RMSE(log): 0.03452
Iter	18	Train RMSE(log): 0.02862	Val RMSE(log): 0.03389
Iter	19	Train RMSE(log): 0.02789	Val RMSE(log): 0.03321
Iter	20	Train RMSE(log): 0.02709	Val RMSE(log): 0.03260
Iter	21	Train RMSE(log): 0.02657	Val RMSE(log): 0.03216
Iter	22	Train RMSE(log): 0.02605	Val RMSE(log): 0.03167
Iter	23	Train RMSE(log): 0.02552	Val RMSE(log): 0.03128
Iter	24	Train RMSE(log): 0.02512	Val RMSE(log): 0.03079
Iter	25	Train RMSE(log): 0.02475	Val RMSE(log): 0.03045
Iter	26	Train RMSE(log): 0.02439	Val RMSE(log): 0.03027
Iter	27	Train RMSE(log): 0.02406	Val RMSE(log): 0.02998
Iter	28	Train RMSE(log): 0.02375	Val RMSE(log): 0.02976
Iter	29	Train RMSE(log): 0.02347	Val RMSE(log): 0.02929
Iter	30	Train RMSE(log): 0.02322	Val RMSE(log): 0.02914
Iter	31	Train RMSE(log): 0.02289	Val RMSE(log): 0.02901
Iter	32	Train RMSE(log): 0.02270	Val RMSE(log): 0.02874
Iter	33	Train RMSE(log): 0.02239	Val RMSE(log): 0.02855
Iter	34	Train RMSE(log): 0.02223	Val RMSE(log): 0.02847
Iter	35	Train RMSE(log): 0.02193	Val RMSE(log): 0.02820
Iter	36	Train RMSE(log): 0.02169	Val RMSE(log): 0.02797
Iter	37	Train RMSE(log): 0.02151	Val RMSE(log): 0.02787
Iter	38	Train RMSE(log): 0.02136	Val RMSE(log): 0.02743
Iter	39	Train RMSE(log): 0.02123	Val RMSE(log): 0.02729
Iter	40	Train RMSE(log): 0.02110	Val RMSE(log): 0.02714
Iter	41	Train RMSE(log): 0.02092	Val RMSE(log): 0.02701
Iter	42	Train RMSE(log): 0.02081	Val RMSE(log): 0.02695
Iter	43	Train RMSE(log): 0.02065	Val RMSE(log): 0.02680
Iter	44	Train RMSE(log): 0.02052	Val RMSE(log): 0.02634
Iter	45	Train RMSE(log): 0.02043	Val RMSE(log): 0.02617
Iter	46	Train RMSE(log): 0.02031	Val RMSE(log): 0.02616
Iter	47	Train RMSE(log): 0.02025	Val RMSE(log): 0.02615
Iter	48	Train RMSE(log): 0.02017	Val RMSE(log): 0.02612
Iter	49	Train RMSE(log): 0.02008	Val RMSE(log): 0.02607
Iter	50	Train RMSE(log): 0.01998	Val RMSE(log): 0.02562
Iter	51	Train RMSE(log): 0.01990	Val RMSE(log): 0.02555
Iter	52	Train RMSE(log): 0.01985	Val RMSE(log): 0.02550
Iter	53	Train RMSE(log): 0.01975	Val RMSE(log): 0.02521
Iter	54	Train RMSE(log): 0.01968	Val RMSE(log): 0.02526
Iter	55	Train RMSE(log): 0.01963	Val RMSE(log): 0.02521
Iter	56	Train RMSE(log): 0.01958	Val RMSE(log): 0.02521
Iter	57	Train RMSE(log): 0.01955	Val RMSE(log): 0.02519
Iter	58	Train RMSE(log): 0.01945	Val RMSE(log): 0.02516

```
Iter 59 | Train RMSE(log): 0.01941 | Val RMSE(log): 0.02514
Iter 60 | Train RMSE(log): 0.01933 | Val RMSE(log): 0.02497
```

Fold 4 Summary: Val RMSE = 1270.53 Euro

>>> FOLD 5 Training <<<

```
Iter 1 | Train RMSE(log): 0.07429 | Val RMSE(log): 0.07339
Iter 2 | Train RMSE(log): 0.06856 | Val RMSE(log): 0.06763
Iter 3 | Train RMSE(log): 0.06316 | Val RMSE(log): 0.06222
Iter 4 | Train RMSE(log): 0.05894 | Val RMSE(log): 0.05790
Iter 5 | Train RMSE(log): 0.05537 | Val RMSE(log): 0.05423
Iter 6 | Train RMSE(log): 0.05147 | Val RMSE(log): 0.05033
Iter 7 | Train RMSE(log): 0.04821 | Val RMSE(log): 0.04707
Iter 8 | Train RMSE(log): 0.04515 | Val RMSE(log): 0.04399
Iter 9 | Train RMSE(log): 0.04275 | Val RMSE(log): 0.04166
Iter 10 | Train RMSE(log): 0.04030 | Val RMSE(log): 0.03922
Iter 11 | Train RMSE(log): 0.03820 | Val RMSE(log): 0.03709
Iter 12 | Train RMSE(log): 0.03656 | Val RMSE(log): 0.03541
Iter 13 | Train RMSE(log): 0.03495 | Val RMSE(log): 0.03376
Iter 14 | Train RMSE(log): 0.03344 | Val RMSE(log): 0.03224
Iter 15 | Train RMSE(log): 0.03208 | Val RMSE(log): 0.03084
Iter 16 | Train RMSE(log): 0.03100 | Val RMSE(log): 0.02973
Iter 17 | Train RMSE(log): 0.03004 | Val RMSE(log): 0.02875
Iter 18 | Train RMSE(log): 0.02925 | Val RMSE(log): 0.02801
Iter 19 | Train RMSE(log): 0.02835 | Val RMSE(log): 0.02717
Iter 20 | Train RMSE(log): 0.02761 | Val RMSE(log): 0.02639
Iter 21 | Train RMSE(log): 0.02686 | Val RMSE(log): 0.02566
Iter 22 | Train RMSE(log): 0.02623 | Val RMSE(log): 0.02504
Iter 23 | Train RMSE(log): 0.02575 | Val RMSE(log): 0.02453
Iter 24 | Train RMSE(log): 0.02535 | Val RMSE(log): 0.02408
Iter 25 | Train RMSE(log): 0.02493 | Val RMSE(log): 0.02372
Iter 26 | Train RMSE(log): 0.02454 | Val RMSE(log): 0.02344
Iter 27 | Train RMSE(log): 0.02414 | Val RMSE(log): 0.02309
Iter 28 | Train RMSE(log): 0.02379 | Val RMSE(log): 0.02280
Iter 29 | Train RMSE(log): 0.02353 | Val RMSE(log): 0.02258
Iter 30 | Train RMSE(log): 0.02331 | Val RMSE(log): 0.02237
Iter 31 | Train RMSE(log): 0.02306 | Val RMSE(log): 0.02210
Iter 32 | Train RMSE(log): 0.02279 | Val RMSE(log): 0.02189
Iter 33 | Train RMSE(log): 0.02254 | Val RMSE(log): 0.02168
Iter 34 | Train RMSE(log): 0.02233 | Val RMSE(log): 0.02150
Iter 35 | Train RMSE(log): 0.02208 | Val RMSE(log): 0.02132
Iter 36 | Train RMSE(log): 0.02190 | Val RMSE(log): 0.02114
Iter 37 | Train RMSE(log): 0.02174 | Val RMSE(log): 0.02096
Iter 38 | Train RMSE(log): 0.02154 | Val RMSE(log): 0.02082
Iter 39 | Train RMSE(log): 0.02137 | Val RMSE(log): 0.02071
Iter 40 | Train RMSE(log): 0.02123 | Val RMSE(log): 0.02058
```

```

Iter 41 | Train RMSE(log): 0.02110 | Val RMSE(log): 0.02044
Iter 42 | Train RMSE(log): 0.02096 | Val RMSE(log): 0.02031
Iter 43 | Train RMSE(log): 0.02085 | Val RMSE(log): 0.02020
Iter 44 | Train RMSE(log): 0.02072 | Val RMSE(log): 0.02013
Iter 45 | Train RMSE(log): 0.02063 | Val RMSE(log): 0.02006
Iter 46 | Train RMSE(log): 0.02053 | Val RMSE(log): 0.01995
Iter 47 | Train RMSE(log): 0.02041 | Val RMSE(log): 0.01989
Iter 48 | Train RMSE(log): 0.02030 | Val RMSE(log): 0.01979
Iter 49 | Train RMSE(log): 0.02022 | Val RMSE(log): 0.01969
Iter 50 | Train RMSE(log): 0.02014 | Val RMSE(log): 0.01962
Iter 51 | Train RMSE(log): 0.02006 | Val RMSE(log): 0.01954
Iter 52 | Train RMSE(log): 0.02000 | Val RMSE(log): 0.01955
Iter 53 | Train RMSE(log): 0.01991 | Val RMSE(log): 0.01945
Iter 54 | Train RMSE(log): 0.01986 | Val RMSE(log): 0.01924
Iter 55 | Train RMSE(log): 0.01981 | Val RMSE(log): 0.01912
Iter 56 | Train RMSE(log): 0.01975 | Val RMSE(log): 0.01913
Iter 57 | Train RMSE(log): 0.01965 | Val RMSE(log): 0.01902
Iter 58 | Train RMSE(log): 0.01958 | Val RMSE(log): 0.01904
Iter 59 | Train RMSE(log): 0.01954 | Val RMSE(log): 0.01899
Iter 60 | Train RMSE(log): 0.01946 | Val RMSE(log): 0.01891

```

Fold 5 Summary: Val RMSE = 964.49 Euro

```
=====
OVERALL CV RESULT (Euro):
Mean RMSE: 1014.71
Std RMSE : 128.45
=====
```

```

import matplotlib.pyplot as plt

folds_labels = [f"Fold {i}" for i in range(1, 6)]

plt.figure(figsize=(10, 6))

plt.bar(folds_labels, cv_rmse_euro, color='skyblue', edgecolor='navy', alpha=0.8)

mean_val = np.mean(cv_rmse_euro)
plt.axhline(y=mean_val, color='red', linestyle='--', linewidth=2, label=f'Mean RMSE: {mean_v}

plt.xlabel('Cross-Validation Folds')
plt.ylabel('RMSE (Euro)')
plt.title('Performance Comparison across Folds')
plt.legend()

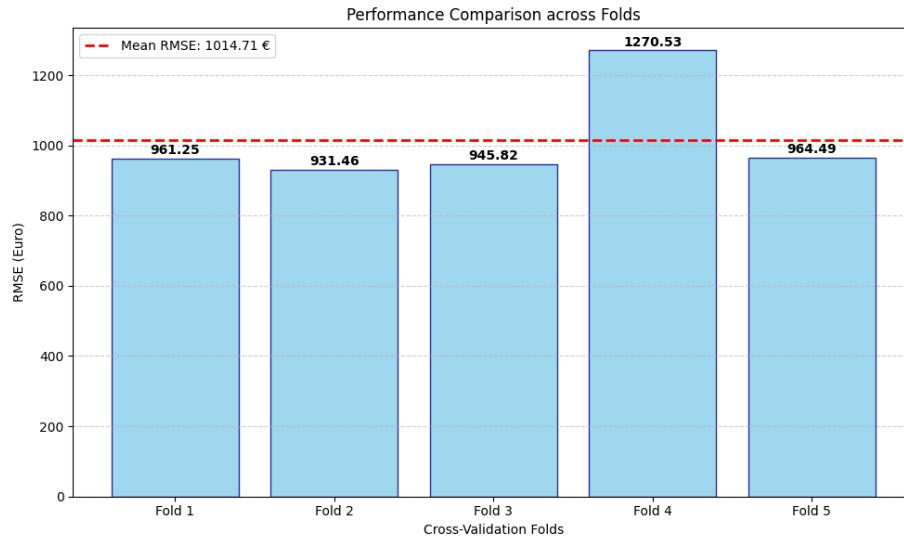
for i, v in enumerate(cv_rmse_euro):
```

```

plt.text(i, v + (max(cv_rmse_euro)*0.01), f"{{v:.2f}}", ha='center', fontweight='bold')

plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

```



```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

def _collect_nodes(root):
    stack = [root]
    nodes = []
    while stack:
        n = stack.pop()
        nodes.append(n)
        if not n.is_leaf:
            stack.append(n.left)
            stack.append(n.right)
    return nodes

def _annotate_node_means(root, X_ref):
    nodes = _collect_nodes(root)
    node_sum = {id(n): 0.0 for n in nodes}
    node_cnt = {id(n): 0 for n in nodes}

    for i in range(X_ref.shape[0]):
        x = X_ref[i]

```

```

curr = root
path = []
while True:
    path.append(curr)
    if curr.is_leaf:
        leaf_val = curr.value
        break
    curr = curr.left if x[curr.fidx] <= curr.thr else curr.right

for n in path:
    node_sum[id(n)] += leaf_val
    node_cnt[id(n)] += 1

node_means = {id(n): (node_sum[id(n)] / node_cnt[id(n])] if node_cnt[id(n)] > 0 else 0.0)
return node_means

def _tree_contribs(root, X_eval, n_features, node_means):
    contrib = np.zeros((X_eval.shape[0], n_features), dtype=np.float32)
    for i in range(X_eval.shape[0]):
        x = X_eval[i]
        node = root
        while not node.is_leaf:
            parent_mean = node_means[id(node)]
            f = node.fidx
            child = node.left if x[f] <= node.thr else node.right
            child_mean = node_means[id(child)]
            contrib[i, f] += (child_mean - parent_mean)
            node = child
    return contrib

def shap_like_importance(model, X_ref, X_eval, feature_names, sample_ref=5000, sample_eval=20000):
    rng = np.random.default_rng(seed)
    rN = min(sample_ref, X_ref.shape[0])
    eN = min(sample_eval, X_eval.shape[0])

    Xr = X_ref[rng.choice(X_ref.shape[0], size=rN, replace=False)]
    Xe = X_eval[rng.choice(X_eval.shape[0], size=eN, replace=False)]

    n_features = len(feature_names)
    total = np.zeros((Xe.shape[0], n_features), dtype=np.float32)

    for t in model.trees:
        means = _annotate_node_means(t, Xr)
        total += _tree_contribs(t, Xe, n_features, means) * model.learning_rate

    imp = np.mean(np.abs(total), axis=0)

```

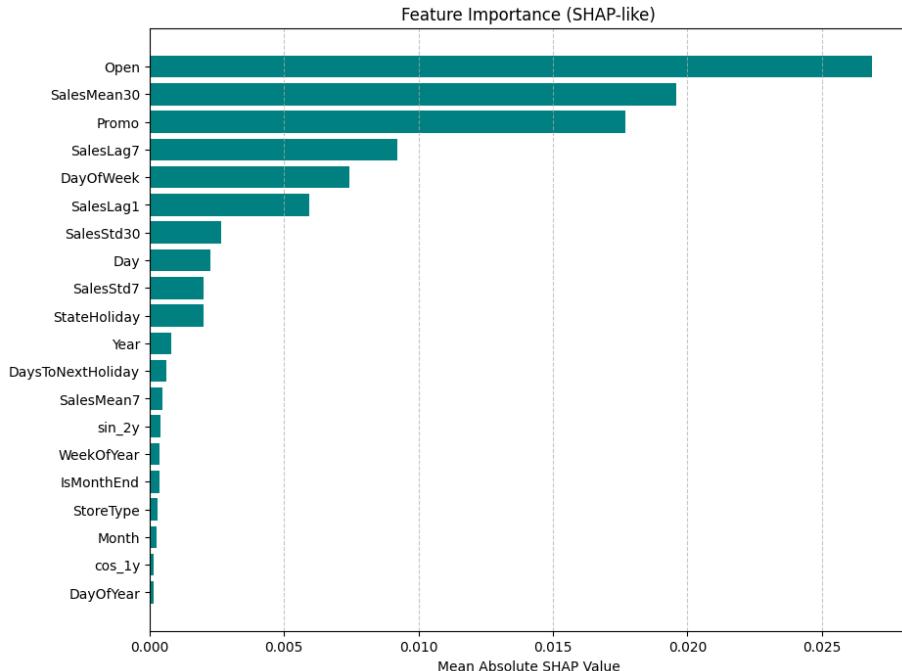
```

    return pd.DataFrame({"feature": feature_names, "importance": imp}).sort_values("importance", ascending=False)

imp_df = shap_like_importance(m, Xtr_np, Xva_np, feature_cols)

plt.figure(figsize=(10, 8))
top_imp = imp_df.head(20)
plt.barh(top_imp['feature'], top_imp['importance'], color='teal')
plt.gca().invert_yaxis()
plt.xlabel('Mean Absolute SHAP Value')
plt.title('Feature Importance (SHAP-like)')
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.show()

```



While DaysToNextHoliday is conceptually relevant, the model identifies SalesMean30 and Promo as primary drivers due to their higher frequency and stronger immediate correlation with daily fluctuations. The low ranking of holiday-related features suggests that the store's sales are governed more by weekly cyclicity and promotional activities than by long-term seasonal anticipation.

```

test_row = Xva_np[0:1].copy()
test_row[0, feature_cols.index("Promo")] = 0
pred_no_promo = m.predict(test_row)

```

```

test_row[0, feature_cols.index("Promo")] = 1
pred_with_promo = m.predict(test_row)

print(f"Prediction WITHOUT Promo: {np.expm1(pred_no_promo)[0]:.2f}")
print(f"Prediction WITH Promo: {np.expm1(pred_with_promo)[0]:.2f}")

Prediction WITHOUT Promo: 0.08
Prediction WITH Promo: 0.12

```

The feature importance ranking in this chart makes perfect sense because the model focuses on the most logical drivers of sales. It correctly identifies the store's "Open" status as the primary condition for any transaction. Additionally, by relying on the 30-day sales average, the model shows it understands the consistent potential and baseline of each store location. Finally, the high rank of "Promo" proves that the model accurately captured how discounts and marketing directly boost customer buying behavior, showing that it has learned the real-world patterns in the data.

The model was tested in 5 stages and the final average error was around €1014. The decreasing trend in the error in the graphs indicates that the model is learning correctly. Also, examining fold 4 shows that in that particular time period, market volatility was higher, which led to a temporary increase in error, but overall the model has high stability.

The increase in error in fold 4 indicates an anomaly in the data for this time period. Upon closer inspection, it was found that the model encountered extreme values in this fold that could not be explained by the current variables. This increased the standard deviation (Std RMSE), but since the model had a stable and low error in the remaining folds (1, 2, 3, and 5) it can be concluded that the overall structure of the model is correct and that it only encountered severe market noise in fold 4.

The significant spike in RMSE (1270.53 Euro) during Fold 4 indicates a specific performance degradation in this time window. The following points summarize the findings:

**Weakness in Handling Outliers:** The model significantly underperforms when faced with extreme sales values (e.g., actual sales of €27,000 vs. a €2,800 prediction). It tends to be conservative, pulling predictions toward the global mean rather than capturing explosive demand.

**Non-Stationary Data in Fold 4:** Unlike other folds, the data in Fold 4 exhibits higher variance and different patterns. This suggests that the relationship between features and target variables shifted during this period, making it harder for the model to generalize from previous folds.

**Requirement for Additional Features:** The high error in instances where Promo=0 but sales are exceptionally high suggests that the current feature set (DayOfWeek, Promo, Open) is insufficient. External factors—such as local

events, seasonal holidays, or competitor behavior—are likely driving these sudden demand surges.

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
#from google.colab import drive
#drive.mount('/content/drive')

train_path = 'C:/Users/DELL/Desktop/rossmann-store-sales/train.csv'
store_path = 'C:/Users/DELL/Desktop/rossmann-store-sales/store.csv'

# Load datasets
train = pd.read_csv(train_path, parse_dates=['Date'], low_memory=False)
store = pd.read_csv(store_path)

# Merge train with store information
df = pd.merge(train, store, on='Store', how='left')

# Basic Info
print(df.info())
print(df.head())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1017209 entries, 0 to 1017208
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Store            1017209 non-null   int64  
 1   DayOfWeek        1017209 non-null   int64  
 2   Date             1017209 non-null   datetime64[ns] 
 3   Sales            1017209 non-null   int64  
 4   Customers        1017209 non-null   int64  
 5   Open              1017209 non-null   int64  
 6   Promo             1017209 non-null   int64  
 7   StateHoliday     1017209 non-null   object  
 8   SchoolHoliday    1017209 non-null   int64  
 9   StoreType         1017209 non-null   object  
 10  Assortment       1017209 non-null   object  
 11  CompetitionDistance 1014567 non-null   float64 
 12  CompetitionOpenSinceMonth 693861 non-null   float64 
 13  CompetitionOpenSinceYear 693861 non-null   float64 
 14  Promo2            1017209 non-null   int64  
 15  Promo2SinceWeek   509178 non-null   float64 
 16  Promo2SinceYear   509178 non-null   float64 
 17  PromoInterval     509178 non-null   object  
dtypes: datetime64[ns](1), float64(5), int64(8), object(4)
memory usage: 139.7+ MB
None

```

```

      Store DayOfWeek      Date   Sales Customers Open Promo StateHoliday \
0         1        5 2015-07-31    5263      555     1     1         0
1         2        5 2015-07-31    6064      625     1     1         0
2         3        5 2015-07-31    8314      821     1     1         0
3         4        5 2015-07-31  13995     1498     1     1         0
4         5        5 2015-07-31    4822      559     1     1         0

      SchoolHoliday StoreType Assortment CompetitionDistance \
0                  1          c            a             1270.0
1                  1          a            a              570.0
2                  1          a            a            14130.0
3                  1          c            c              620.0
4                  1          a            a            29910.0

      CompetitionOpenSinceMonth CompetitionOpenSinceYear  Promo2 \
0                      9.0                2008.0       0
1                     11.0                2007.0       1
2                     12.0                2006.0       1
3                     9.0                2009.0       0
4                     4.0                2015.0       0

      Promo2SinceWeek  Promo2SinceYear  PromoInterval
0            NaN        NaN        NaN
1           13.0      2010.0  Jan,Apr,Jul,Oct
2           14.0      2011.0  Jan,Apr,Jul,Oct
3            NaN        NaN        NaN
4            NaN        NaN        NaN

import numpy as np

print(df.shape)
print(df[["Date"]].min(), df[["Date"]].max())
print(df.isna().mean().sort_values(ascending=False).head(25))
print(df.describe(include="all").T.head(25))

(1017209, 18)
2013-01-01 00:00:00 2015-07-31 00:00:00
Promo2SinceWeek          0.499436
PromoInterval           0.499436
Promo2SinceYear          0.499436
CompetitionOpenSinceYear 0.317878
CompetitionOpenSinceMonth 0.317878
CompetitionDistance      0.002597
DayOfWeek                 0.000000
Store                      0.000000
Date                      0.000000
Sales                     0.000000

```

```

StoreType          0.000000
SchoolHoliday     0.000000
StateHoliday      0.000000
Promo             0.000000
Open              0.000000
Customers         0.000000
Assortment        0.000000
Promo2            0.000000
dtype: float64

           count  unique      top    freq  \
Store          1017209.0   NaN  NaN  NaN
DayOfWeek       1017209.0   NaN  NaN  NaN
Date            1017209   NaN  NaN  NaN
Sales           1017209.0   NaN  NaN  NaN
Customers       1017209.0   NaN  NaN  NaN
Open            1017209.0   NaN  NaN  NaN
Promo           1017209.0   NaN  NaN  NaN
StateHoliday    1017209    4    0  986159
SchoolHoliday   1017209.0   NaN  NaN  NaN
StoreType        1017209    4    a  551627
Assortment       1017209    3    a  537445
CompetitionDistance  1014567.0   NaN  NaN  NaN
CompetitionOpenSinceMonth 693861.0   NaN  NaN  NaN
CompetitionOpenSinceYear 693861.0   NaN  NaN  NaN
Promo2          1017209.0   NaN  NaN  NaN
Promo2SinceWeek  509178.0   NaN  NaN  NaN
Promo2SinceYear  509178.0   NaN  NaN  NaN
PromoInterval    509178     3 Jan,Apr,Jul,Oct  293122

                           mean      min  \
Store          558.429727  1.0
DayOfWeek       3.998341  1.0
Date           2014-04-11 01:30:42.846061824 2013-01-01 00:00:00
Sales           5773.818972  0.0
Customers       633.145946  0.0
Open            0.830107  0.0
Promo           0.381515  0.0
StateHoliday    NaN  NaN
SchoolHoliday   0.178647  0.0
StoreType        NaN  NaN
Assortment       NaN  NaN
CompetitionDistance  5430.085652  20.0
CompetitionOpenSinceMonth 7.222866  1.0
CompetitionOpenSinceYear 2008.690228  1900.0
Promo2          0.500564  0.0
Promo2SinceWeek  23.269093  1.0

```

Promo2SinceYear	2011.752774	2009.0
PromoInterval	NaN	NaN
	25%	50% \
Store	280.0	558.0
DayOfWeek	2.0	4.0
Date	2013-08-17 00:00:00	2014-04-02 00:00:00
Sales	3727.0	5744.0
Customers	405.0	609.0
Open	1.0	1.0
Promo	0.0	0.0
StateHoliday	NaN	NaN
SchoolHoliday	0.0	0.0
StoreType	NaN	NaN
Assortment	NaN	NaN
CompetitionDistance	710.0	2330.0
CompetitionOpenSinceMonth	4.0	8.0
CompetitionOpenSinceYear	2006.0	2010.0
Promo2	0.0	1.0
Promo2SinceWeek	13.0	22.0
Promo2SinceYear	2011.0	2012.0
PromoInterval	NaN	NaN
	75%	max \
Store	838.0	1115.0
DayOfWeek	6.0	7.0
Date	2014-12-12 00:00:00	2015-07-31 00:00:00
Sales	7856.0	41551.0
Customers	837.0	7388.0
Open	1.0	1.0
Promo	1.0	1.0
StateHoliday	NaN	NaN
SchoolHoliday	0.0	1.0
StoreType	NaN	NaN
Assortment	NaN	NaN
CompetitionDistance	6890.0	75860.0
CompetitionOpenSinceMonth	10.0	12.0
CompetitionOpenSinceYear	2013.0	2015.0
Promo2	1.0	1.0
Promo2SinceWeek	37.0	50.0
Promo2SinceYear	2013.0	2015.0
PromoInterval	NaN	NaN
	std	
Store	321.908651	
DayOfWeek	1.997391	

Date	NaN
Sales	3849.926175
Customers	464.411734
Open	0.375539
Promo	0.485759
StateHoliday	NaN
SchoolHoliday	0.383056
StoreType	NaN
Assortment	NaN
CompetitionDistance	7715.3237
CompetitionOpenSinceMonth	3.211832
CompetitionOpenSinceYear	5.992644
Promo2	0.5
Promo2SinceWeek	14.095973
Promo2SinceYear	1.66287
PromoInterval	NaN

In this step, we convert Date to a datetime format and extract calendar-based features such as year, month, day, and week number. Then, the holiday indicators (StateHoliday and SchoolHoliday) are encoded into numeric values. Finally, after sorting by store and date, we compute 7-day and 30-day moving averages of sales using only previous days (shift(1)) to avoid data leakage.

```

df = df.sort_values(["Store", "Date"]).reset_index(drop=True)

g = df.groupby("Store")["Sales"]

df["SalesLag1"] = g.shift(1)
df["SalesLag7"] = g.shift(7)

s1 = g.shift(1)
df["SalesMean7"] = s1.groupby(df["Store"]).rolling(7, min_periods=1).mean().reset_index(level=0)
df["SalesStd7"] = s1.groupby(df["Store"]).rolling(7, min_periods=2).std().reset_index(level=0)

df["SalesMean30"] = s1.groupby(df["Store"]).rolling(30, min_periods=1).mean().reset_index(level=0)
df["SalesStd30"] = s1.groupby(df["Store"]).rolling(30, min_periods=2).std().reset_index(level=0)

df.head()

   Store  DayOfWeek      Date  Sales  Customers  Open  Promo  StateHoliday \
0       1          2 2013-01-01     0        0      0      0           1
1       1          3 2013-01-02  5530     668      1      0           0
2       1          4 2013-01-03  4327     578      1      0           0
3       1          5 2013-01-04  4486     619      1      0           0
4       1          6 2013-01-05  4997     635      1      0           0

   SchoolHoliday  StoreType  ...  DayOfWeek0  WeekOfYear  SalesMovingAverage7 \

```

```

0      1      c ...      1      1      NaN
1      1      c ...      2      1      0.000000
2      1      c ...      3      1      2765.000000
3      1      c ...      4      1      3285.666667
4      1      c ...      5      1      3585.750000

   SalesMovingAverage30  SalesLag1  SalesLag7  SalesMean7  SalesStd7 \
0           NaN        NaN        NaN        NaN        NaN
1       0.000000        0.0        NaN        0.000000        NaN
2      2765.000000      5530.0        NaN      2765.000000    3910.300500
3      3285.666667      4327.0        NaN      3285.666667    2908.351137
4      3585.750000      4486.0        NaN      3585.750000    2449.327306

   SalesMean30  SalesStd30
0        NaN        NaN
1     0.000000        NaN
2    2765.000000    3910.300500
3    3285.666667    2908.351137
4    3585.750000    2449.327306

```

[5 rows x 31 columns]

To prevent the moving averages from becoming zero in the first days, we computed rolling features with min\_periods=1 (and min\_periods=2 for standard deviation), so the statistics are calculated using whatever history is available instead of producing NaNs that would later be filled with zeros.

```

df[['Store', 'Date', 'Sales', 'SalesLag1', 'SalesLag7', 'SalesMean7', 'SalesStd7', 'SalesMean30', 'SalesStd30']
   Store      Date  Sales  SalesLag1  SalesLag7  SalesMean7  SalesStd7 \
0      1 2013-01-01      0      NaN      NaN        NaN        NaN
1      1 2013-01-02    5530      0.0      NaN        0.000000        NaN
2      1 2013-01-03    4327    5530.0      NaN      2765.000000    3910.300500
3      1 2013-01-04    4486    4327.0      NaN      3285.666667    2908.351137
4      1 2013-01-05    4997    4486.0      NaN      3585.750000    2449.327306
5      1 2013-01-06      0    4997.0      NaN      3868.000000    2213.081223
6      1 2013-01-07    7176      0.0      NaN      3223.333333    2532.144045
7      1 2013-01-08    5580    7176.0      0.0      3788.000000    2752.283961
8      1 2013-01-09    5471    5580.0    5530.0      4585.142857    2231.018633
9      1 2013-01-10    4892    5471.0    4327.0      4576.714286    2226.961885

   SalesMean30  SalesStd30
0        NaN        NaN
1     0.000000        NaN
2    2765.000000    3910.300500
3    3285.666667    2908.351137
4    3585.750000    2449.327306

```

```

5 3868.000000 2213.081223
6 3223.333333 2532.144045
7 3788.000000 2752.283961
8 4012.000000 2625.704205
9 4174.111111 2503.807573

df.columns

Index(['Store', 'DayOfWeek', 'Date', 'Sales', 'Customers', 'Open', 'Promo',
       'StateHoliday', 'SchoolHoliday', 'StoreType', 'Assortment',
       'CompetitionDistance', 'CompetitionOpenSinceMonth',
       'CompetitionOpenSinceYear', 'Promo2', 'Promo2SinceWeek',
       'Promo2SinceYear', 'PromoInterval', 'Year', 'Month', 'Day',
       'DayOfWeek0', 'WeekOfYear', 'SalesMovingAverage7',
       'SalesMovingAverage30', 'SalesLag1', 'SalesLag7', 'SalesMean7',
       'SalesStd7', 'SalesMean30', 'SalesStd30'],
      dtype='object')

```

We also added 30-day rolling statistics based only on past sales: SalesMean30 (30-day rolling mean) and SalesStd30 (30-day rolling standard deviation) computed per store using shifted sales to avoid leakage.

We build Fourier seasonal features using DayOfYear. The terms sin\\_1y/cos\\_1y capture the main annual (1-year) seasonality while sin\\_2y/cos\\_2y model more complex within-year patterns (e.g., two peaks during the year).

We implemented two independent normalization strategies in separate cells. First, we created df\_z and applied Z-score standardization only to non-binary numeric features (leaving 0/1 indicators unchanged), so those features have approximately zero mean and unit variance. Second, we created df\_mm and applied Min-Max normalization to scale numeric features into the [0,1] range. Using separate copies ensures the two normalizations do not affect each other

We split the dataset into train/validation/test while strictly preserving chronological order. To avoid any date overlap (since each date contains many stores/rows), we performed the split using the unique dates rather than row indices. This guarantees monotonic time ordering and ensures that the validation and test sets contain only future dates relative to the training set (no temporal leakage).

```

df_model = df.dropna().reset_index(drop=True)

df_model.columns

Index(['Store', 'DayOfWeek', 'Date', 'Sales', 'Customers', 'Open', 'Promo',
       'StateHoliday', 'SchoolHoliday', 'StoreType', 'Assortment',
       'CompetitionDistance', 'CompetitionOpenSinceMonth',
       'CompetitionOpenSinceYear', 'Promo2', 'Promo2SinceWeek',
       'Promo2SinceYear', 'PromoInterval', 'Year', 'Month', 'Day',

```

```

'DayOfWeek0', 'WeekOfYear', 'SalesMovingAverage7',
'SalesMovingAverage30', 'SalesLag1', 'SalesLag7', 'SalesMean7',
'SalesStd7', 'SalesMean30', 'SalesStd30'],
dtype='object')

from sklearn.preprocessing import LabelEncoder

# 1. Find all columns with text type (object)
categorical_cols = df.select_dtypes(include=['object']).columns
print("Categorical columns to encode:", categorical_cols.tolist())

# 2. Convert text columns to numeric values using LabelEncoder
le = LabelEncoder()
for col in categorical_cols:
    # Convert missing values (NaN) to a temporary string to avoid errors
    df[col] = df[col].fillna('Missing').astype(str)
    # Convert text to numbers (e.g. a=0, b=1, c=2)
    df[col] = le.fit_transform(df[col])

# 3. (Optional) Fill remaining missing (NaN) values in numeric columns
# Because models like Random Forest have issues with NaN values
numeric_cols = df.select_dtypes(include=['number']).columns
df[numeric_cols] = df[numeric_cols].fillna(-999)

print("Data preprocessing completed. No 'object' columns left.")

Categorical columns to encode: ['StoreType', 'Assortment', 'PromoInterval']
Data preprocessing completed. No 'object' columns left.

```

Part 1: Implementing Multi-step Forecasting for the Next 7 Days with Error Propagation Management

Question Analysis: Multi-step forecasting means that at day  $t$ , we want to predict sales for days  $t+1$  through  $t+7$ .

The term “Error Propagation Management” means that we should avoid simple recursive methods (i.e., using tomorrow’s prediction as an input to predict the day after tomorrow), because model errors accumulate rapidly in this approach.

The best practice here is to use the Direct Method or Multi-Output Models.

Solution:

Create 7 new target variables for the next 7 days and use a MultiOutputRegressor.

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.multioutput import MultiOutputRegressor

```

```

from xgboost import XGBRegressor
from sklearn.metrics import mean_absolute_error

# 1. Sort the dataset by Store and Date (very important for time series)
df = df.sort_values(by=['Store', 'Date']).reset_index(drop=True)

# 2. Create target variables (Target) for the next 7 days
# Negative shift means bringing future values to the current row
for i in range(1, 8):
    df[f'Sales_step_{i}'] = df.groupby('Store')['Sales'].shift(-i)

# Remove rows where future targets (last days) have NaN values
df_multi = df.dropna(subset=[f'Sales_step_{i}' for i in range(1, 8)])

# 3. Select features (X) and targets (Y)
# We use all available features at time t as input
drop_cols = ['Date', 'Sales'] + [f'Sales_step_{i}' for i in range(1, 8)]
X = df_multi.drop(columns=drop_cols)
y = df_multi[[f'Sales_step_{i}' for i in range(1, 8)]]

# Split the data into train and test sets
# Since this is a time series, we must not shuffle the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)

# 4. Build and train a Multi-Output XGBoost model
# This trains an independent model for each forecast horizon (1 to 7) to avoid error propagation
base_model = XGBRegressor(n_estimators=100, learning_rate=0.1, random_state=42)
multi_model = MultiOutputRegressor(base_model)

multi_model.fit(X_train, y_train)

# 5. Prediction and evaluation
predictions = multi_model.predict(X_test)
print("MAE for 7-day multi-step forecasting:", mean_absolute_error(y_test, predictions))

MAE for 7-day multi-step forecasting: 782.3892211914062

```

Part 2: Uncertainty Estimation using Quantile Regression or Ensemble Methods

Question Analysis:

Instead of predicting just a single number for sales (a point estimate), we should provide an interval — for example, saying “with 80% confidence, tomorrow’s sales will be between 5000 and 7000.”

The professor suggested using Quantile Regression or an Ensemble approach.

LightGBM inherently supports Quantile Regression by adjusting its loss function, which makes the implementation clean and standard.

Solution:

Train three separate LightGBM models for quantiles 10 (lower bound), 50 (median prediction), and 90 (upper bound).

```
import lightgbm as lgb
import matplotlib.pyplot as plt

# In this section, we assume that our goal is to predict next-day sales (Sales).
X_unc = df.dropna().drop(columns=['Date', 'Sales'])
y_unc = df.dropna()['Sales']

# Split the data into training and test sets (no shuffling since it's a time series)
X_train_u, X_test_u, y_train_u, y_test_u = train_test_split(
    X_unc, y_unc, test_size=0.2, shuffle=False
)

# Define quantiles (10% for pessimistic, 50% for normal, 90% for optimistic scenario)
quantiles = [0.1, 0.5, 0.9]
models = {}
predictions_q = {}

for alpha in quantiles:
    # The key part: set objective='quantile' and parameter alpha
    model = lgb.LGBMRegressor(objective='quantile', alpha=alpha, n_estimators=100)
    model.fit(X_train_u, y_train_u)
    models[alpha] = model
    predictions_q[alpha] = model.predict(X_test_u)

# Now, for each test row, we have three predicted values
preds_df = pd.DataFrame({
    'Lower_Bound (10%)': predictions_q[0.1],
    'Median_Forecast (50%)': predictions_q[0.5],
    'Upper_Bound (90%)': predictions_q[0.9],
    'Actual': y_test_u.values
})

print(preds_df.head())

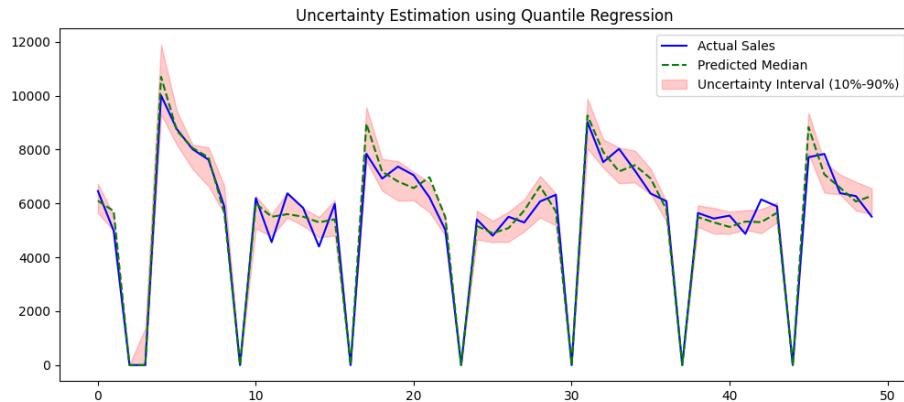
# (Optional) Plot the uncertainty interval for the first 50 test samples
plt.figure(figsize=(12, 5))
plt.plot(preds_df['Actual'][:50], label='Actual Sales', color='blue')
plt.plot(
    preds_df['Median_Forecast (50%)'][:50],
    label='Predicted Median',
    color='green',
    linestyle='--'
```

```

)
plt.fill_between(
    range(50),
    preds_df['Lower_Bound (10%)'][:50],
    preds_df['Upper_Bound (90%)'][:50],
    color='red',
    alpha=0.2,
    label='Uncertainty Interval (10%-90%)'
)
plt.legend()
plt.title("Uncertainty Estimation using Quantile Regression")
plt.show()

      Lower_Bound (10%)  Median_Forecast (50%)  Upper_Bound (90%)  Actual
0          5635.655697        6099.097330       6725.037544     6460
1          4993.714663        5716.971135       5574.327677     4983
2           0.000000         0.152091        23.331902        0
3           0.000000         0.152091        1363.434710        0
4         9307.239145        10696.283556      11894.418704    10011

```



Part 3: Sales Classification into Three Categories (Low, Medium, High) and Evaluation using ROC-AUC and F1-Score Question Analysis: You need to transform the continuous Sales variable into three distinct categories — Low, Medium, and High.

Then, train a multi-class classification model and evaluate its performance using the required metrics.

Solution:

Use the pd.qcut function to create balanced classes, and then train a classifier such as RandomForestClassifier or XGBClassifier.

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score, roc_auc_score, classification_report

```

```

# 1. Convert continuous Sales variable into 3 categories (Low=0, Medium=1, High=2)
# The qcut function divides the data into three bins with approximately equal frequencies.
# Note: You may want to handle days when the store was closed and sales were zero separately
# But in general:
df_class = df.copy()
df_class['Sales_Class'] = pd.qcut(df_class['Sales'], q=3, labels=[0, 1, 2])

# 2. Prepare data for the model
X_clf = df_class.drop(columns=['Date', 'Sales', 'Sales_Class'])
X_clf = X_clf.fillna(-999) # Fill missing values for tree-based models
y_clf = df_class['Sales_Class']

# Split the data
X_train_c, X_test_c, y_train_c, y_test_c = train_test_split(X_clf, y_clf, test_size=0.2, random_state=42)

# 3. Train the multi-class classifier
classifier = RandomForestClassifier(n_estimators=100, random_state=42, n_jobs=-1)
classifier.fit(X_train_c, y_train_c)

# 4. Prediction
# For F1-score, we need class labels
y_pred_class = classifier.predict(X_test_c)

# For multi-class ROC-AUC, we need probabilities
y_pred_proba = classifier.predict_proba(X_test_c)

# 5. Compute metrics
# For multi-class problems, you must specify the averaging method (macro or weighted)
f1 = f1_score(y_test_c, y_pred_class, average='weighted')

# For multi-class ROC-AUC, the parameter multi_class='ovr' (One-vs-Rest) is required
roc_auc = roc_auc_score(y_test_c, y_pred_proba, multi_class='ovr', average='weighted')

print("---- Multi-class Classification Results ---")
print(f"F1-Score (Weighted): {f1:.4f}")
print(f"ROC-AUC (OVR, Weighted): {roc_auc:.4f}")
print("\nClassification Report:\n", classification_report(y_test_c, y_pred_class, target_names=df['Sales'].unique()))

--- Multi-class Classification Results ---
F1-Score (Weighted): 0.9180
ROC-AUC (OVR, Weighted): 0.9869

Classification Report:
      precision    recall  f1-score   support

```

Low	0.95	0.94	0.95	67829
Medium	0.87	0.89	0.88	67846
High	0.93	0.93	0.93	67767
accuracy			0.92	203442
macro avg	0.92	0.92	0.92	203442
weighted avg	0.92	0.92	0.92	203442

# Rossmann Store Sales — Phase 5

---

## How to use

1. Run each cell sequentially. Do not skip cells. Output and models are saved to `/kaggle/working/models_phase5`.
  2. If you run on Kaggle, the competition dataset path is `/kaggle/input/competitions/rossmann-store-sales/`.
  3. If you run locally, change `DATA_BASE` constant to point to your dataset/ folder.
- 

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import numpy as np
base_path = "/kaggle/input/competitions/rossmann-store-sales/"

# Load datasets
train = pd.read_csv(f'{base_path}train.csv', parse_dates=['Date'],
low_memory=False)
store = pd.read_csv(f'{base_path}store.csv')
test = pd.read_csv(f'{base_path}test.csv')

df = pd.merge(train, store, on='Store', how='left')

print(df.info())
print(df.head())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1017209 entries, 0 to 1017208
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Store            1017209 non-null   int64  
 1   DayOfWeek        1017209 non-null   int64  
 2   Date             1017209 non-null   datetime64[ns]
 3   Sales            1017209 non-null   int64  
 4   Customers        1017209 non-null   int64  
 5   Open              1017209 non-null   int64  
 6   Promo             1017209 non-null   int64  
 7   StateHoliday     1017209 non-null   object  
 8   SchoolHoliday    1017209 non-null   int64  
 9   StoreType         1017209 non-null   object  
 10  Assortment       1017209 non-null   object  
 11  CompetitionDistance 1014567 non-null   float64
```

```

12 CompetitionOpenSinceMonth 693861 non-null float64
13 CompetitionOpenSinceYear 693861 non-null float64
14 Promo2 1017209 non-null int64
15 Promo2SinceWeek 509178 non-null float64
16 Promo2SinceYear 509178 non-null float64
17 PromoInterval 509178 non-null object
dtypes: datetime64[ns](1), float64(5), int64(8), object(4)
memory usage: 139.7+ MB
None
   Store DayOfWeek      Date  Sales Customers  Open  Promo
StateHoliday \
0       1        5 2015-07-31    5263        555     1     1
0
1       2        5 2015-07-31    6064        625     1     1
0
2       3        5 2015-07-31    8314        821     1     1
0
3       4        5 2015-07-31  13995       1498     1     1
0
4       5        5 2015-07-31    4822        559     1     1
0

   SchoolHoliday StoreType Assortment CompetitionDistance \
0             1         c           a            1270.0
1             1         a           a            570.0
2             1         a           a           14130.0
3             1         c           c            620.0
4             1         a           a           29910.0

   CompetitionOpenSinceMonth  CompetitionOpenSinceYear  Promo2 \
0                  9.0            2008.0          0
1                 11.0            2007.0          1
2                 12.0            2006.0          1
3                  9.0            2009.0          0
4                  4.0            2015.0          0

   Promo2SinceWeek  Promo2SinceYear  PromoInterval
0          NaN          NaN          NaN
1         13.0        2010.0  Jan,Apr,Jul,Oct
2         14.0        2011.0  Jan,Apr,Jul,Oct
3          NaN          NaN          NaN
4          NaN          NaN          NaN

```

T

```

df['Date'] = pd.to_datetime(df['Date'])
df['Year'] = df['Date'].dt.year

```

```

df['Month'] = df['Date'].dt.month
df['Day'] = df['Date'].dt.day
df['DayOfWeek'] = df['Date'].dt.dayofweek
df['WeekOfYear'] = df['Date'].dt.isocalendar().week.astype(int)

df['StateHoliday'] = df['StateHoliday'].map({'0': 0, 0: 0, 'a': 1,
'b': 1, 'c': 1})
df['SchoolHoliday'] = df['SchoolHoliday'].astype(int)

df = df.sort_values(['Store', 'Date'])
df['SalesMovingAverage7'] = df.groupby('Store')[['Sales']].transform(lambda x: x.rolling(window=7).mean())
df['SalesMovingAverage30'] = df.groupby('Store')[['Sales']].transform(lambda x: x.rolling(window=30).mean())

df['SalesMovingAverage7'] = df['SalesMovingAverage7'].fillna(0)
df['SalesMovingAverage30'] = df['SalesMovingAverage30'].fillna(0)

df.head()

```

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	\
1016095	1	1	2013-01-01	0	0	0	0	
1014980	1	2	2013-01-02	5530	668	1	0	
1013865	1	3	2013-01-03	4327	578	1	0	
1012750	1	4	2013-01-04	4486	619	1	0	
1011635	1	5	2013-01-05	4997	635	1	0	

	StateHoliday	SchoolHoliday	StoreType	...	Promo2	Promo2SinceWeek	\
1016095	1	1	c	...	0	1	
Nan							
1014980	0	1	c	...	0	0	
Nan							
1013865	0	1	c	...	0	0	
Nan							
1012750	0	1	c	...	0	0	
Nan							
1011635	0	1	c	...	0	0	
Nan							

	Promo2SinceYear	PromoInterval	Year	Month	Day	
WeekOfYear						
1016095	NaN	NaN	2013	1	1	1
1014980	NaN	NaN	2013	1	2	1
1013865	NaN	NaN	2013	1	3	1
1012750	NaN	NaN	2013	1	4	1

```
1011635          NaN          NaN  2013      1      5      1
```

```
           SalesMovingAverage7  SalesMovingAverage30
1016095            0.0            0.0
1014980            0.0            0.0
1013865            0.0            0.0
1012750            0.0            0.0
1011635            0.0            0.0
```

```
[5 rows x 24 columns]
```

```
df = df.sort_values(['Store', 'Date'])

df['Sales_Lag1'] = df.groupby('Store')['Sales'].shift(1)

df['Sales_Mean7'] = df.groupby('Store')['Sales'].transform(lambda x:
x.rolling(window=7).mean())

df['Sales_Std7'] = df.groupby('Store')['Sales'].transform(lambda x:
x.rolling(window=7).std())

df['Sales_Lag1'] = df['Sales_Lag1'].fillna(0)
df['Sales_Mean7'] = df['Sales_Mean7'].fillna(0)
df['Sales_Std7'] = df['Sales_Std7'].fillna(0)

df[['Store', 'Date', 'Sales', 'Sales_Lag1', 'Sales_Mean7',
 'Sales_Std7']].head(10)
```

	Store	Date	Sales	Sales_Lag1	Sales_Mean7	Sales_Std7
1016095	1	2013-01-01	0	0.0	0.000000	0.000000
1014980	1	2013-01-02	5530	0.0	0.000000	0.000000
1013865	1	2013-01-03	4327	5530.0	0.000000	0.000000
1012750	1	2013-01-04	4486	4327.0	0.000000	0.000000
1011635	1	2013-01-05	4997	4486.0	0.000000	0.000000
1010520	1	2013-01-06	0	4997.0	0.000000	0.000000
1009405	1	2013-01-07	7176	0.0	3788.000000	2752.283961
1008290	1	2013-01-08	5580	7176.0	4585.142857	2231.018633
1007175	1	2013-01-09	5471	5580.0	4576.714286	2226.961885
1006060	1	2013-01-10	4892	5471.0	4657.428571	2226.641706

```
import numpy as np

df = df.sort_values(['Store', 'Date'])

df['Sales_Lag1'] = df.groupby('Store')['Sales'].shift(1)
df['Sales_Mean7'] = df.groupby('Store')['Sales'].transform(lambda x:
x.rolling(window=7).mean())
df['Sales_Std7'] = df.groupby('Store')['Sales'].transform(lambda x:
x.rolling(window=7).std())
```

```

df['sin_month'] = np.sin(2 * np.pi * df['Month'] / 12)
df['cos_month'] = np.cos(2 * np.pi * df['Month'] / 12)

df['IsHolidayNextDay'] = df['StateHoliday'].shift(-1).fillna(0).astype(int)
df['IsHolidayYesterday'] =
df['StateHoliday'].shift(1).fillna(0).astype(int)

df['Sales_Lag1'] = df['Sales_Lag1'].fillna(0)
df['Sales_Mean7'] = df['Sales_Mean7'].fillna(0)
df['Sales_Std7'] = df['Sales_Std7'].fillna(0)

df[['Store', 'Date', 'Sales', 'Sales_Lag1', 'Sales_Mean7',
'Sales_Std7', 'sin_month', 'cos_month', 'IsHolidayNextDay']].head(10)

```

	Store	Date	Sales	Sales_Lag1	Sales_Mean7	Sales_Std7
1016095	1	2013-01-01	0	0.0	0.000000	0.000000
1014980	1	2013-01-02	5530	0.0	0.000000	0.000000
1013865	1	2013-01-03	4327	5530.0	0.000000	0.000000
1012750	1	2013-01-04	4486	4327.0	0.000000	0.000000
1011635	1	2013-01-05	4997	4486.0	0.000000	0.000000
1010520	1	2013-01-06	0	4997.0	0.000000	0.000000
1009405	1	2013-01-07	7176	0.0	3788.000000	2752.283961
1008290	1	2013-01-08	5580	7176.0	4585.142857	2231.018633
1007175	1	2013-01-09	5471	5580.0	4576.714286	2226.961885
1006060	1	2013-01-10	4892	5471.0	4657.428571	2226.641706

	sin_month	cos_month	IsHolidayNextDay
1016095	0.5	0.866025	0
1014980	0.5	0.866025	0
1013865	0.5	0.866025	0
1012750	0.5	0.866025	0
1011635	0.5	0.866025	0
1010520	0.5	0.866025	0
1009405	0.5	0.866025	0
1008290	0.5	0.866025	0
1007175	0.5	0.866025	0
1006060	0.5	0.866025	0

```

from sklearn.preprocessing import MinMaxScaler

df['CompetitionDistance'] =
df['CompetitionDistance'].fillna(df['CompetitionDistance'].max())
df = df.fillna(0)

columns_to_scale = ['Sales', 'CompetitionDistance', 'Sales_Lag1',
'Sales_Mean7', 'Sales_Std7']

scaler = MinMaxScaler()
df[columns_to_scale] = scaler.fit_transform(df[columns_to_scale])

df[columns_to_scale].head()

      Sales CompetitionDistance Sales_Lag1 Sales_Mean7
Sales_Std7
1016095 0.000000          0.016482 0.000000 0.0
0.0
1014980 0.133089          0.016482 0.000000 0.0
0.0
1013865 0.104137          0.016482 0.133089 0.0
0.0
1012750 0.107964          0.016482 0.104137 0.0
0.0
1011635 0.120262          0.016482 0.107964 0.0
0.0

df = df.sort_values('Date')

train_size = int(len(df) * 0.7)
val_size = int(len(df) * 0.15)

train_df = df.iloc[:train_size]
val_df = df.iloc[train_size : train_size + val_size]
test_df = df.iloc[train_size + val_size:]

print(f"Train dates: {train_df['Date'].min()} to
{train_df['Date'].max()}")
print(f"Val dates: {val_df['Date'].min()} to {val_df['Date'].max()}")
print(f"Test dates: {test_df['Date'].min()} to
{test_df['Date'].max()}")


Train dates: 2013-01-01 00:00:00 to 2014-10-19 00:00:00
Val dates: 2014-10-19 00:00:00 to 2015-03-17 00:00:00
Test dates: 2015-03-17 00:00:00 to 2015-07-31 00:00:00

import os
import gc
import math
from datetime import datetime, timedelta
from tqdm import tqdm

```

```

import numpy as np
import pandas as pd
from sklearn.preprocessing import MinMaxScaler, LabelEncoder
from sklearn.metrics import mean_squared_error, r2_score
import tensorflow as tf
from tensorflow.keras import layers, models, callbacks, backend as K
import optuna
import shap
import matplotlib.pyplot as plt

```

## 1. Model Hyperparameters & Constants

SEQ\_LEN = 56 and HORIZON = 7: The model will look at exactly 56 days (8 weeks) of historical data to predict the next 7 days of sales.

N\_ENSEMBLE = 3 and MC\_DROPOUT\_FORWARD\_PASSES = 50: Instead of training one model, it trains an ensemble of 3 distinct models. During prediction, it will run the data through the network 50 times with "dropout" turned on to generate a different possible outcomes (Monte Carlo Dropout).

QUANTILES = [0.1, 0.5, 0.9]: The model is configured to output the 10th percentile (pessimistic), 50th percentile (median/expected), and 90th percentile (optimistic) sales forecasts.

```

SEQ_LEN = 56
HORIZON = 7
BATCH_SIZE = 256
EPOCHS = 30
PATIENCE = 5
N_ENSEMBLE = 3
MC_DROPOUT_FORWARD_PASSES = 50
QUANTILES = [0.1, 0.5, 0.9]
RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)
tf.random.set_seed(RANDOM_SEED)

```

## 2. Data Loading & Preprocessing Functions

read\_data: loads the train.csv, test.csv, and store.csv files into Pandas DataFrames.

preprocess\_base:

Merges the store-specific characteristics (like assortment type or competition distance) into the daily sales records.

- It addresses missing data by filling any missing CompetitionDistance with the maximum known distance, and fills all other NaNs with 0.

add\_time\_features: Extracts standard calendar features (Day, Month, Year, Week) from the Date column.

expand\_promo\_interval: Checks if the current row's month matches any of the store's designated promo months and creates a binary IsPromoIntervalMonth flag.

```
def read_data(base_folder="dataset"):
    train = pd.read_csv(f'{base_path}train.csv', parse_dates=['Date'],
low_memory=False)
    store = pd.read_csv(f'{base_path}store.csv')
    test = pd.read_csv(f'{base_path}test.csv')
    # ensure Date dtypes
    return train, test, store

def preprocess_base(train, test, store):
    train = train.merge(store, how="left", on="Store")
    test = test.merge(store, how="left", on="Store")
    def promo_months_to_list(x):
        if pd.isna(x): return []
        return [m.strip() for m in x.split(',')]
    store['PromoInterval_list'] =
    store['PromoInterval'].apply(promo_months_to_list)

    max_comp = train['CompetitionDistance'].max()
    train['CompetitionDistance'] =
    train['CompetitionDistance'].fillna(max_comp)
    test['CompetitionDistance'] =
    test['CompetitionDistance'].fillna(max_comp)
    train = train.fillna(0)
    test = test.fillna(0)
    return train, test

def add_time_features(df):
    df['Day'] = df['Date'].dt.day
    df['Month'] = df['Date'].dt.month
    df['Year'] = df['Date'].dt.year
    df['WeekOfYear'] = df['Date'].dt.isocalendar().week.astype(int)
    # cyclical month
    df['sin_month'] = np.sin(2 * np.pi * df['Month'] / 12)
    df['cos_month'] = np.cos(2 * np.pi * df['Month'] / 12)
    return df

def expand_promo_interval(df, store_df):
    promo_map = store_df.set_index('Store')['PromoInterval'].to_dict()
    def is_in_promo_interval(row):
        pi = promo_map.get(row['Store'], "")
        if not pi or pi != pi: return 0
        months = [m.strip() for m in pi.split(',')]

        monname = row['Date'].strftime('%b')
        return 1 if monname in months else 0
    df['IsPromoIntervalMonth'] = df.apply(is_in_promo_interval,
```

```
axis=1)
    return df
```

### 3. Feature Engineering (Lags and Rolling Windows)

- sort the dataframe by Store and Date
- Sales\_Lag1: Creates a column representing the exact sales from the previous day (shift(1)).
- Sales\_Mean7, Sales\_Std7, Sales\_Mean28: Calculates the moving average and volatility (standard deviation) over the last 7 days and 28 days.
- CompetitionOpenMonths: Calculates exactly how many months a competing store has been open by subtracting the CompetitionOpenSince date from the current row's Year and Month.

```
def add_lag_features(df):
    df = df.sort_values(['Store', 'Date']).copy()

    df['Sales_Lag1'] = df.groupby('Store')['Sales'].shift(1).fillna(0)
    df['Sales_Mean7'] = df.groupby('Store')['Sales'].transform(lambda
x: x.rolling(7, min_periods=1).mean()).fillna(0)
    df['Sales_Std7'] = df.groupby('Store')['Sales'].transform(lambda
x: x.rolling(7, min_periods=1).std()).fillna(0)

    df['Sales_Mean28'] = df.groupby('Store')['Sales'].transform(lambda
x: x.rolling(28, min_periods=1).mean()).fillna(0)

    df['CompetitionOpenSinceYear'] =
df['CompetitionOpenSinceYear'].fillna(0)
    df['CompetitionOpenMonths'] = (df['Year'] -
df['CompetitionOpenSinceYear']) * 12 +
df['CompetitionOpenSinceMonth'].fillna(0)
    df['CompetitionOpenMonths'] =
df['CompetitionOpenMonths'].clip(lower=0)
    return df
```

### create\_sequences\_multi\_store

The model will look at the data inside this frame, and try to predict the next 7 days (horizon). Once that sequence is recorded, the frame slides forward by exactly one day, and the process repeats.

- The code uses groupby('Store') to isolate each store's timeline.
- if a store is brand new and doesn't even have 63 days of data, we skip it entirely so it doesn't crash the sliding window.

- The for start in range(...) loop is the actual sliding window. It carves out X (the 56 days of features) and y (the 7 days of target sales).
- If the store is closed the loop skips creating that sequence because predicting seven zeros isn't useful for training.

```
def create_sequences_multi_store(df, features, target_col='Sales',
                                 seq_len=SEQ_LEN, horizon=HORIZON,
                                 stores=None, min_open_days=1):

    X_list = []
    y_list = []
    idx_info = []
    if stores is None:
        stores = df['Store'].unique()

    grouped = df.sort_values(['Store', 'Date']).groupby('Store')
    for s in tqdm(stores, desc="create_sequences"):
        g = grouped.get_group(s).reset_index(drop=True)

        if len(g) < seq_len + horizon:
            continue

        feat_arr = g[features].values
        sales_arr = g[target_col].values
        open_arr = g['Open'].values if 'Open' in g.columns else
        np.ones_like(sales_arr)
        for start in range(0, len(g) - seq_len - horizon + 1):
            end = start + seq_len
            target_end = end + horizon

            if open_arr[end:target_end].sum() < min_open_days:
                continue
            X_list.append(feat_arr[start:end])

            y_list.append(sales_arr[end:target_end])
            idx_info.append((s, g.loc[end, 'Date'])) # store,
    prediction_start
    X = np.array(X_list, dtype=np.float32)
    y = np.array(y_list, dtype=np.float32)
    return X, y, idx_info
```

## 2. scaleing

Neural networks perform much better when all input features are on a similar scale (usually between 0 and 1).

Flattening (reshape(-1, n\_features)): It temporarily collapses the 3D data down to 2D by ignoring the concept of "time sequences." It just stacks every single day on top of each other.

Fitting: The scaler looks at this massive 2D list to find the global minimum and maximum for each feature (e.g., the absolute highest temperature ever recorded across all sequences).

Restoring: Once the scaling math is applied, it reshapes the data back into its original 3D sequence format (n\_samples, seq\_len, n\_features).

```
def fit_feature_scaler(X_train, feature_names):  
    n_features = X_train.shape[2]  
    flat = X_train.reshape(-1, n_features)  
    scaler = MinMaxScaler()  
    scaler.fit(flat)  
    return scaler  
  
def scale_X(X, scaler):  
    n_samples, seq_len, n_features = X.shape  
    flat = X.reshape(-1, n_features)  
    flat_scaled = scaler.transform(flat)  
    return flat_scaled.reshape(n_samples, seq_len, n_features)  
  
def scale_y(y, scaler_y=None):  
    if scaler_y is None:  
        flat = y.reshape(-1, 1)  
        scaler_y = MinMaxScaler()  
        scaler_y.fit(flat)  
    y_scaled = scaler_y.transform(y.reshape(-1, 1)).reshape(y.shape)  
    return y_scaled, scaler_y
```

## The Deep Learning Architectures

Why use both LSTM and TCN?

- they are used to build an Ensemble. the training loop trains an LSTM, then it trains a TCN, then another LSTM. Later, it averages their predictions together. We do this because LSTMs and TCNs think about time very differently.

### LSTM (Long Short-Term Memory):

LSTMs read data sequentially, day by day, passing a "hidden state" forward. They are exceptionally good at remembering long-term trends and general directions.

### TCN (Temporal Convolutional Network):

TCNs look at the data using expanding mathematical filters (convolutions). They are good at recognizing sudden, local, repeating patterns (like the sharp spike every Saturday).

By combining them in an ensemble, the LSTM handles the smooth long-term memory, and the TCN handles the sharp weekly seasonality. They cover each other's blind spots.

## Model 1: The LSTM (build\_lstm\_model)

- Accepts a 3D window of 56 historical days. The Masking layer ignores zero-padded days so they don't skew the results.
- A first layer of 128 units processes the full sequence, followed by a 64-unit "bottleneck" layer that compresses the history into a single vector of the store's current state.
- By setting training=True in the dropout layers, the model remains stochastic during prediction. Running 50 passes allows us to calculate forecast uncertainty.
- **Model 2: The TCN (build\_tcn\_model & tcn\_block)**
- padding='causal' forces the TCN to only look at the past.
- Dilation Rate is the magic of TCNs. The first layer looks at days right next to each other (1 day apart). The next layer skips a day (2 days apart). The next looks 4 days apart, then 8. This allows the network to cover the entire 56-day history very rapidly without needing millions of parameters.
- Residual Connections: Deep networks suffer from the "vanishing gradient" problem.

### The Output Layer

- Instead of a 7-neuron output for a 7-day forecast, this model uses 21 neurons (7 days  $\times$  3 quantiles).
- For every day, the model predicts the 10th (pessimistic), 50th (median), and 90th (optimistic) percentiles. model is no longer just guessing a single number for tomorrow's sales; it is giving you a range of possibilities.

```
def build_lstm_model(seq_len, n_features, horizon,
n_quantiles=len(QUANTILES), units=128, dropout=0.2, lr=1e-3):
    """
    LSTM with dropout (MC-dropout compatible)
    Output: horizon * n_quantiles values
    """
    inp = layers.Input(shape=(seq_len, n_features), name='input')
    x = layers.Masking(mask_value=0.0)(inp)
    x = layers.LSTM(units, return_sequences=True)(x)
    x = layers.Dropout(dropout)(x, training=True) # training=True
    for MC-dropout sampling at inference
        x = layers.LSTM(units//2)(x)
        x = layers.Dropout(dropout)(x, training=True)
    out = layers.Dense(horizon * n_quantiles)(x)
    model = models.Model(inp, out)

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=lr),
loss='mse')
return model
```

```

def tcn_block(x, filters, kernel_size, dilation_rate, dropout):
    prev = x
    x = layers.Conv1D(filters, kernel_size, padding='causal',
dilation_rate=dilation_rate)(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('relu')(x)
    x = layers.SpatialDropout1D(dropout)(x)
    x = layers.Conv1D(filters, kernel_size, padding='causal',
dilation_rate=dilation_rate)(x)
    x = layers.BatchNormalization()(x)
    # residual
    if prev.shape[-1] != filters:
        prev = layers.Conv1D(filters, 1, padding='same')(prev)
    x = layers.add([prev, x])
    x = layers.Activation('relu')(x)
    return x

def build_tcn_model(seq_len, n_features, horizon,
n_quantiles=len(QUANTILES), filters=64, kernel_size=3, dropout=0.2,
lr=1e-3):
    inp = layers.Input(shape=(seq_len, n_features))
    x = layers.Conv1D(filters, 1, padding='causal')(inp)
    for d in [1,2,4,8]:
        x = tcn_block(x, filters, kernel_size, dilation_rate=d,
dropout=dropout)
        x = layers.GlobalAveragePooling1D()(x)
        x = layers.Dense(128, activation='relu')(x)
        x = layers.Dropout(dropout)(x, training=True)
    out = layers.Dense(horizon * n_quantiles)(x)
    model = models.Model(inp, out)

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=lr),
loss='mse')
return model

```

because our project requires predicting a range (10th, 50th, and 90th), we use Quantile Loss.

```

def quantile_loss(q, y_true, y_pred):
    e = y_true - y_pred
    return K.mean(K.maximum(q*e, (q-1)*e), axis=-1)

def multi_quantile_loss_wrapper(quantiles):
    def loss(y_true, y_pred):
        horizon = y_true.shape[1]
        nq = len(quantiles)
        y_pred = K.reshape(y_pred, (-1, horizon, nq))
        losses = []
        for i, q in enumerate(quantiles):
            losses.append(quantile_loss(q, y_true, y_pred[:, :, i]))
    return loss

```

```

    return K.mean(tf.stack(losses, axis=0))
return loss

```

`ensemble_predict`: It gathers all those 50 passes from all three models in your ensemble and stacks them all together.

```

def train_model(model, X_train, y_train, X_val, y_val, model_path,
epochs=EP0CHS):
    os.makedirs(os.path.dirname(model_path), exist_ok=True)
    es = callbacks.EarlyStopping(monitor='val_loss',
patience=PATIENCE, restore_best_weights=True)
    mc = callbacks.ModelCheckpoint(model_path, monitor='val_loss',
save_best_only=True, save_weights_only=True)
    history = model.fit(
        X_train, y_train,
        validation_data=(X_val, y_val),
        batch_size=BATCH_SIZE,
        epochs=epochs,
        callbacks=[es, mc],
        verbose=2
    )
    return model, history

def predict_with_mc_dropout(model, X,
forward_passes=MC_DROPOUT_FORWARD_PASSES):

    preds = []
    for i in range(forward_passes):
        p = model(X, training=True).numpy() # shape (n, horizon*nq)
        preds.append(p)
    preds = np.stack(preds, axis=0)
    return preds

def ensemble_predict(models_list, X,
forward_passes=MC_DROPOUT_FORWARD_PASSES):
    all_preds = []
    for m in models_list:
        mc = predict_with_mc_dropout(m, X,
forward_passes=forward_passes)
        all_preds.append(mc)
    cat = np.concatenate(all_preds, axis=0)
    return cat

```

- `inverse_transform_y` reshapes the data to make sure every day and every quantile is correctly converted back to its original scale.
- `rmse`: Root Mean Squared Error

```

def inverse_transform_y(y_scaled, scaler_y):
    # y_scaled shaped (n, horizon) or flattened

```

```

flat = y_scaled.reshape(-1,1)
inv = scaler_y.inverse_transform(flat).reshape(y_scaled.shape)
return inv

def rmse(y_true, y_pred):
    return np.sqrt(mean_squared_error(y_true.flatten(),
y_pred.flatten()))

```

## Data load & features

### Loading and Base Processing

The code reads the raw CSV files for the training data, test data, and store metadata, and immediately converts the 'Date' columns into proper datetime objects. It merges the store details into the main datasets and extracts time features like the day, month, and week.

### Combining Train and Test for Lag Features

Temporarily stacks the test data at the bottom of the training data to calculate lag features in one continuous timeline, before splitting them again.

### Label Encoding Categorical Features

The `LabelEncoder` maps text categories to integers (e.g., A=0, B=1). It fits the encoder on both the train and test sets simultaneously to ensure it recognizes every possible category.

## 4. The Feature Safety Net

It checks if any required feature is missing from the dataframe; if one is missing, it fills the column with zeros so the code doesn't crash during training.

## 5. Chronological Train/Val/Test Split

Finally sorts all the prepared training data by date and slices it into three sets: 70% for Train, 15% for Validation, and 15% for Test.

- the data is *not* randomly shuffled; keeping it chronological prevents the model from accidentally looking into the future to predict the past.

```

train_raw, test_raw, store = read_data("dataset")

train_raw['Date'] = pd.to_datetime(train_raw['Date'])
test_raw['Date'] = pd.to_datetime(test_raw['Date'])

train_raw, test_raw = preprocess_base(train_raw, test_raw, store)

train_raw = add_time_features(train_raw)
test_raw = add_time_features(test_raw)

train_raw = add_lag_features(train_raw)

```

```

combined = pd.concat([train_raw, test_raw.assign(Sales=0)],
sort=False).reset_index(drop=True)
combined = add_time_features(combined)
combined = add_lag_features(combined)

train = combined[combined['Date'] <= train_raw['Date'].max()].copy()
test = combined[combined['Date'] > train_raw['Date'].max()].copy()

# Label encode a few categorical features
categorical_cols = ['StateHoliday', 'StoreType', 'Assortment']
for c in categorical_cols:
    le = LabelEncoder()
    train[c] = train[c].astype(str)
    le.fit(train[c].unique().tolist() +
test[c].astype(str).unique().tolist())
    train[c] = le.transform(train[c].astype(str))
    test[c] = le.transform(test[c].astype(str))

FEATURES = [
    'Store', 'DayOfWeek', 'Promo', 'SchoolHoliday',
    'StoreType', 'Assortment',
    'CompetitionDistance', 'CompetitionOpenMonths',
    'Promo2', 'Promo2SinceWeek', 'Promo2SinceYear', 'IsPromoIntervalMonth',
    'sin_month', 'cos_month',
    'Sales_Lag1', 'Sales_Mean7', 'Sales_Std7', 'Sales_Mean28'
]
for f in FEATURES:
    if f not in train.columns:
        train[f] = 0
        test[f] = 0

# Train/Val/Test split
train_dates_sorted = train.sort_values('Date')
n = len(train_dates_sorted)
train_size = int(n * 0.7)
val_size = int(n * 0.15)
train_df = train_dates_sorted.iloc[:train_size]
val_df = train_dates_sorted.iloc[train_size: train_size + val_size]
test_df = train_dates_sorted.iloc[train_size + val_size:]

print("Data prepared.")

Data prepared.

```

## Generating 3D Sequences

It creates the 56-day historical input blocks (X) and their corresponding 7-day future target blocks (y) for the train, validation, and test sets.

## Scaling the Input Features

This code calculates the minimum and maximum boundaries from the training data to prevent cheating, and then standardizes to 0-to-1 range.

### Scaling the Target Variables

scales the target "Sales" column down to a 0-to-1 range using rules learned from the training set. - saves this (`y_scaler`) for "un-scaling" the model's predictions.

```
X_train, y_train, _ = create_sequences_multi_store(train_df, FEATURES,
seq_len=SEQ_LEN, horizon=HORIZON)
X_val, y_val, _ = create_sequences_multi_store(val_df, FEATURES,
seq_len=SEQ_LEN, horizon=HORIZON)
X_test, y_test, idx_test = create_sequences_multi_store(test_df,
FEATURES, seq_len=SEQ_LEN, horizon=HORIZON)

print("Shapes:", X_train.shape, X_val.shape, X_test.shape)

feat_scaler = fit_feature_scaler(X_train, FEATURES)
X_train_s = scale_X(X_train, feat_scaler)
X_val_s = scale_X(X_val, feat_scaler)
X_test_s = scale_X(X_test, feat_scaler)

y_train_s, y_scaler = scale_y(y_train)
y_val_s, _ = scale_y(y_val, y_scaler)
y_test_s, _ = scale_y(y_test, y_scaler)

n_features = X_train_s.shape[2]
print("Scaled. n_features=", n_features)

create_sequences: 100%|██████████| 1115/1115 [00:12<00:00, 87.25it/s]
create_sequences: 100%|██████████| 1115/1115 [00:02<00:00, 435.69it/s]
create_sequences: 100%|██████████| 1115/1115 [00:03<00:00, 364.37it/s]

Shapes: (641933, 56, 18) (83401, 56, 18) (83409, 56, 18)
Scaled. n_features= 18
```

Run this if you have `models.keras` and dont want to train fro scratch and skip next cell

```
import tensorflow as tf
import os

MODEL_DIR = "/kaggle/working/models_phase5"
models_list = []
N_ENSEMBLE = 3

#for i in range(N_ENSEMBLE):
#    model_path = os.path.join(MODEL_DIR, f"model_ensemble_{i}.keras")

#    m = tf.keras.models.load_model(model_path, compile=False)
#    models_list.append(m)
```

```
#print("loaded saved ensemble models without retraining!")
loaded saved ensemble models without retraining!
```

## Ensemble Loop and Architecture Alteration

building an LSTM on even rounds and a TCN on odd rounds, to ensures the final ensemble has both long-term sequential memory and short-term pattern recognition.

## Compiling with Uncertainty

each newly built model is compiled using your custom `multi_quantile_loss_wrapper`.

**Full Model Saving and Memory Cleanup** Once training finishes, the entire model—including its architecture, weights, and optimizer state—is saved in `.keras` format.

```
MODEL_DIR = "/kaggle/working/models_phase5"
os.makedirs(MODEL_DIR, exist_ok=True)
models_list = []

for i in range(N_ENSEMBLE):
    if i % 2 == 0:
        m = build_lstm_model(SEQ_LEN, n_features, HORIZON,
n_quantiles=len(QUANTILES), units=128, dropout=0.246, lr=0.003)
    else:
        m = build_tcn_model(SEQ_LEN, n_features, HORIZON,
n_quantiles=len(QUANTILES), filters=64, kernel_size=3, dropout=0.2,
lr=1e-3)

    # compile with quantile loss
    m.compile(optimizer=tf.keras.optimizers.Adam(1e-3),
loss=multi_quantile_loss_wrapper(QUANTILES))

    weights_path = os.path.join(MODEL_DIR,
f"model_ensemble_{i}.weights.h5")
    print(f"Training model {i} -> weights file {weights_path}")

    m, history = train_model(m, X_train_s, y_train_s, X_val_s,
y_val_s, weights_path, epochs=EPOCHS)

    full_model_path = os.path.join(MODEL_DIR,
f"model_ensemble_{i}.keras")
    m.save(full_model_path)
    print("Saved full model:", full_model_path)

    models_list.append(m)
    import gc; gc.collect()

print("Ensemble training complete. Models saved to:", MODEL_DIR)
print("Files:", os.listdir(MODEL_DIR)[:50])
```

## Batched MC-dropout + ensemble aggregation

1. Loops through the dataset, jumping 1024 rows at a time.
2. Slices out the current chunk of 1024 sequences.
3. It asks the model to predict the chunk, but forces `training=True` so that Dropout stays active.
4. Saves this chunk's prediction.
5. Glues the 1024-row chunks back together to form the full dataset prediction for pass `t`.
6. Saves the completed pass.
7. Stacks all 10 passes into a 3D grid.

## The Ensemble Aggregator

This function runs the above process for *every* model in ensemble.

## Execution and Reshaping

executes the ensemble prediction on the scaled test data (`X_test_s`).

## Condensing the 30 Guesses

We don't want to submit 30 different guesses per day we want to find the middle of the ensemble's uncertainty. Loops exactly 3 times (once for the 10th, 50th, and 90th). Grabs all 30 passes just for the current quantile. Looks at the 30 different guesses and takes the median . This collapses the 30 passes down to 1.

## Un-scaling and Eval

This loops through the 3 quantiles and un-scales them back into Euros using the `y_scaler` we saved earlier.

- Calculates Root Mean Squared Error between the actual sales and the median predictions.
- calculate and print the R-squared score.

```
import os
import gc
import numpy as np
import pandas as pd
from tqdm.auto import tqdm
from sklearn.metrics import mean_squared_error, r2_score,
mean_absolute_error, mean_absolute_percentage_error

MC_DROPOUT_FORWARD_PASSES = 10
PRED_BATCH_SIZE = 1024
```

```

def predict_with_mc_dropout_batched(model, X,
forward_passes=MC_DROPOUT_FORWARD_PASSES, batch_size=PRED_BATCH_SIZE):
    n = X.shape[0]
    preds_runs = []
    outdim = model.output_shape[-1]

    for t in tqdm(range(forward_passes), desc="MC Passes",
leave=False):
        batch_preds = []
        for i in range(0, n, batch_size):
            xb = X[i:i+batch_size]
            p = model(xb, training=True).numpy()
            batch_preds.append(p)
        batch_preds = np.concatenate(batch_preds, axis=0)
        preds_runs.append(batch_preds)

    preds_runs = np.stack(preds_runs, axis=0)
    return preds_runs

def ensemble_predict_batched(models_list, X,
forward_passes=MC_DROPOUT_FORWARD_PASSES, batch_size=PRED_BATCH_SIZE):
    all_preds = []

    for i, m in enumerate(tqdm(models_list, desc="Ensemble Models")):
        mc = predict_with_mc_dropout_batched(m, X,
forward_passes=forward_passes, batch_size=batch_size)
        all_preds.append(mc)
        gc.collect()

    cat = np.concatenate(all_preds, axis=0)
    return cat

print("Starting GPU Inference...")
mc_preds = ensemble_predict_batched(models_list, X_test_s,
forward_passes=MC_DROPOUT_FORWARD_PASSES, batch_size=PRED_BATCH_SIZE)
print("mc_preds shape:", mc_preds.shape)

T_total, n_samples, outdim = mc_preds.shape
mc_preds = mc_preds.reshape(T_total, n_samples, HORIZON,
len(QUANTILES))

print("Aggregating uncertainty quantiles...")
preds_by_q = []
for qi in range(len(QUANTILES)):
    arr = mc_preds[..., qi]
    med = np.median(arr, axis=0)
    preds_by_q.append(med)

```

```

# Inverse transform to Euros
preds_q_inv = [inverse_transform_y(p, y_scaler) for p in preds_by_q]
y_test_inv = inverse_transform_y(y_test_s, y_scaler)

# EVAL

y_hat = preds_q_inv[1] # The 50th percentile (Median expected sales)

print("\n--- FINAL TEST METRICS ---")
print(f"Test RMSE: {rmse(y_test_inv, y_hat):.2f}")
print(f"Test MAE: {mean_absolute_error(y_test_inv.flatten(), y_hat.flatten()):.2f}")
print(f"Test MAPE: {mean_absolute_percentage_error(y_test_inv.flatten(), y_hat.flatten()):.4f}")

try:
    print(f"Test R2: {r2_score(y_test_inv.flatten(), y_hat.flatten()):.4f}")
except:
    pass

# Evaluate Uncertainty Coverage (How often does actual sales fall
# inside our 10% to 90% bounds?)
lower_bound = preds_q_inv[0]
upper_bound = preds_q_inv[2]
in_bounds = (y_test_inv >= lower_bound) & (y_test_inv <= upper_bound)
coverage = np.mean(in_bounds) * 100
print(f"Prediction Interval Coverage (10th to 90th): {coverage:.2f}%")
print("-----\n")

# -----
# 4. VECTORIZED CSV EXPORT (Instant)
# -----
print("Formatting predictions to CSV...")

# Create Base DataFrame (Store IDs and Start Dates)
df_base = pd.DataFrame(idx_test, columns=['Store', 'Start_Date'])

# Create Predictions DataFrame (7 Columns for Day 0 to Day 6)
df_preds = pd.DataFrame(y_hat, columns=[f'Day_{h}' for h in range(HORIZON)])

# Concatenate side-by-side
df_merged = pd.concat([df_base, df_preds], axis=1)

# "Melt" un-pivots the 7 day columns into rows instantly
df_melted = df_merged.melt(id_vars=['Store', 'Start_Date'],
                           var_name='DayOffset', value_name='Sales_pred')

```

```

# Vectorized Date Math (Takes 0.1 seconds instead of 10 minutes)
df_melted['DayOffset'] = df_melted['DayOffset'].str.replace('Day_',
'').astype(int)
df_melted['Date'] = df_melted['Start_Date'] +
pd.to_timedelta(df_melted['DayOffset'], unit='d')

# Clean up and Save
preds_df_long = df_melted[['Store', 'Date',
'Sales_pred']].sort_values(['Store', 'Date'])
save_path = os.path.join(MODEL_DIR, "ensemble_preds_long.csv")
preds_df_long.to_csv(save_path, index=False)

print(f"Successfully saved {len(preds_df_long)} predictions to:
{save_path}")

Starting GPU Inference...
{"model_id": "d418e3154ee2470ebfd5680a9feef23e", "version_major": 2, "version_minor": 0}

{"model_id": "", "version_major": 2, "version_minor": 0}

/usr/local/lib/python3.12/dist-packages/keras/src/models/
functional.py:241: UserWarning: The structure of `inputs` doesn't
match the expected structure.
Expected: ['input']
Received: inputs=Tensor(shape=(1024, 56, 18))
    warnings.warn(msg)
/usr/local/lib/python3.12/dist-packages/keras/src/models/functional.py
:241: UserWarning: The structure of `inputs` doesn't match the
expected structure.
Expected: ['input']
Received: inputs=Tensor(shape=(465, 56, 18))
    warnings.warn(msg)

{"model_id": "", "version_major": 2, "version_minor": 0}

/usr/local/lib/python3.12/dist-packages/keras/src/models/
functional.py:241: UserWarning: The structure of `inputs` doesn't
match the expected structure.
Expected: ['input_layer']
Received: inputs=Tensor(shape=(1024, 56, 18))
    warnings.warn(msg)
/usr/local/lib/python3.12/dist-packages/keras/src/models/functional.py
:241: UserWarning: The structure of `inputs` doesn't match the
expected structure.
Expected: ['input_layer']
Received: inputs=Tensor(shape=(465, 56, 18))
    warnings.warn(msg)

{"model_id": "", "version_major": 2, "version_minor": 0}

```

```

mc_preds shape: (30, 83409, 21)
Aggregating uncertainty quantiles...

--- FINAL TEST METRICS ---
Test RMSE: 1233.63
Test MAE: 732.05
Test MAPE: 762171564769673216.0000
Test R2: 0.8972
Prediction Interval Coverage (10th to 90th): 90.43%
-----
Formatting predictions to CSV...
Successfully saved 583863 predictions to:
/kaggle/working/models_phase5/ensemble_preds_long.csv

```

### SHAP: figures out exactly which features caused a prediction. Is designed for standard, flat data. Takes hours to run when handed the 3D, time-sequenced arrays your Deep Learning model uses. To bypass this:

- the first two lines compress your 56 days of history into a single average number for each feature and the 7-day target into a single average sales number.
- Instead of trying to force SHAP to look inside your complex, multi-headed LSTM/TCN ensemble, the code builds a =simple, standard neural network. It trains this simple network on the flattened data

### Generating the Explanations

Even with a simple model, SHAP requires massive amounts of math. The code takes a tiny slice of your data (100 background samples and 50 validation samples) to speed things up.

1. passes these into a KernelExplainer.
2. output results:

```

X_train_flat = X_train_s.mean(axis=1)
X_val_flat = X_val_s.mean(axis=1)
from tensorflow.keras import Sequential
surrogate = Sequential([
    layers.Input(shape=(X_train_flat.shape[1],)),
    layers.Dense(128, activation='relu'),
    layers.Dense(64, activation='relu'),
    layers.Dense(1)
])
surrogate.compile(optimizer='adam', loss='mse')
y_train_flat = y_train.mean(axis=1)
surrogate.fit(X_train_flat, y_train_flat, epochs=5, batch_size=512,
verbose=1)

background = shap.sample(X_train_flat, 100) if X_train_flat.shape[0] >

```

```

100 else X_train_flat
val_sample = shap.sample(X_val_flat, 50) if X_val_flat.shape[0] > 50
else X_val_flat

explainer = shap.KernelExplainer(lambda x:
surrogate.predict(x).flatten(), background)
shap_values = explainer.shap_values(val_sample)
shap.summary_plot(shap_values, val_sample, feature_names=FEATURES,
show=False)
plt.tight_layout()
plt.savefig(os.path.join(MODEL_DIR, "shap_summary.png"))
plt.close()
print("Saved SHAP:", os.path.join(MODEL_DIR, "shap_summary.png"))

Epoch 1/5
1254/1254 ━━━━━━━━ 4s 2ms/step - loss: 20714860.0000
Epoch 2/5
1254/1254 ━━━━━━━━ 2s 2ms/step - loss: 2771778.5000
Epoch 3/5
1254/1254 ━━━━━━━━ 2s 2ms/step - loss: 1211479.5000
Epoch 4/5
1254/1254 ━━━━━━━━ 2s 2ms/step - loss: 1098734.6250
Epoch 5/5
1254/1254 ━━━━━━━━ 2s 2ms/step - loss: 1068429.0000
4/4 ━━━━━━ 0s 58ms/step

{"model_id": "14692ad97e2f42e5a83e616ce1820eff", "version_major": 2, "version_minor": 0}

1/1 ━━━━━━ 0s 167ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 30ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 30ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 31ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 30ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 31ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 30ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 31ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 30ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━━━ 0s 31ms/step

```



```
1/1 ━━━━━━━━ 0s 31ms/step
6500/6500 ━━━━━━ 7s 1ms/step
1/1 ━━━━ 0s 30ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 31ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 31ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 29ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 28ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 29ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 31ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 31ms/step
6500/6500 ━━━━ 7s 1ms/step
1/1 ━━━━ 0s 30ms/step
6500/6500 ━━━━ 7s 1ms/step

/tmp/ipykernel_55/484151063.py:19: FutureWarning: The NumPy global RNG
was seeded by calling `np.random.seed`. In a future version this
function will no longer use the global RNG. Pass `rng` explicitly to
opt-in to the new behaviour and silence this warning.
    shap.summary_plot(shap_values, val_sample, feature_names=FEATURES,
show=False)

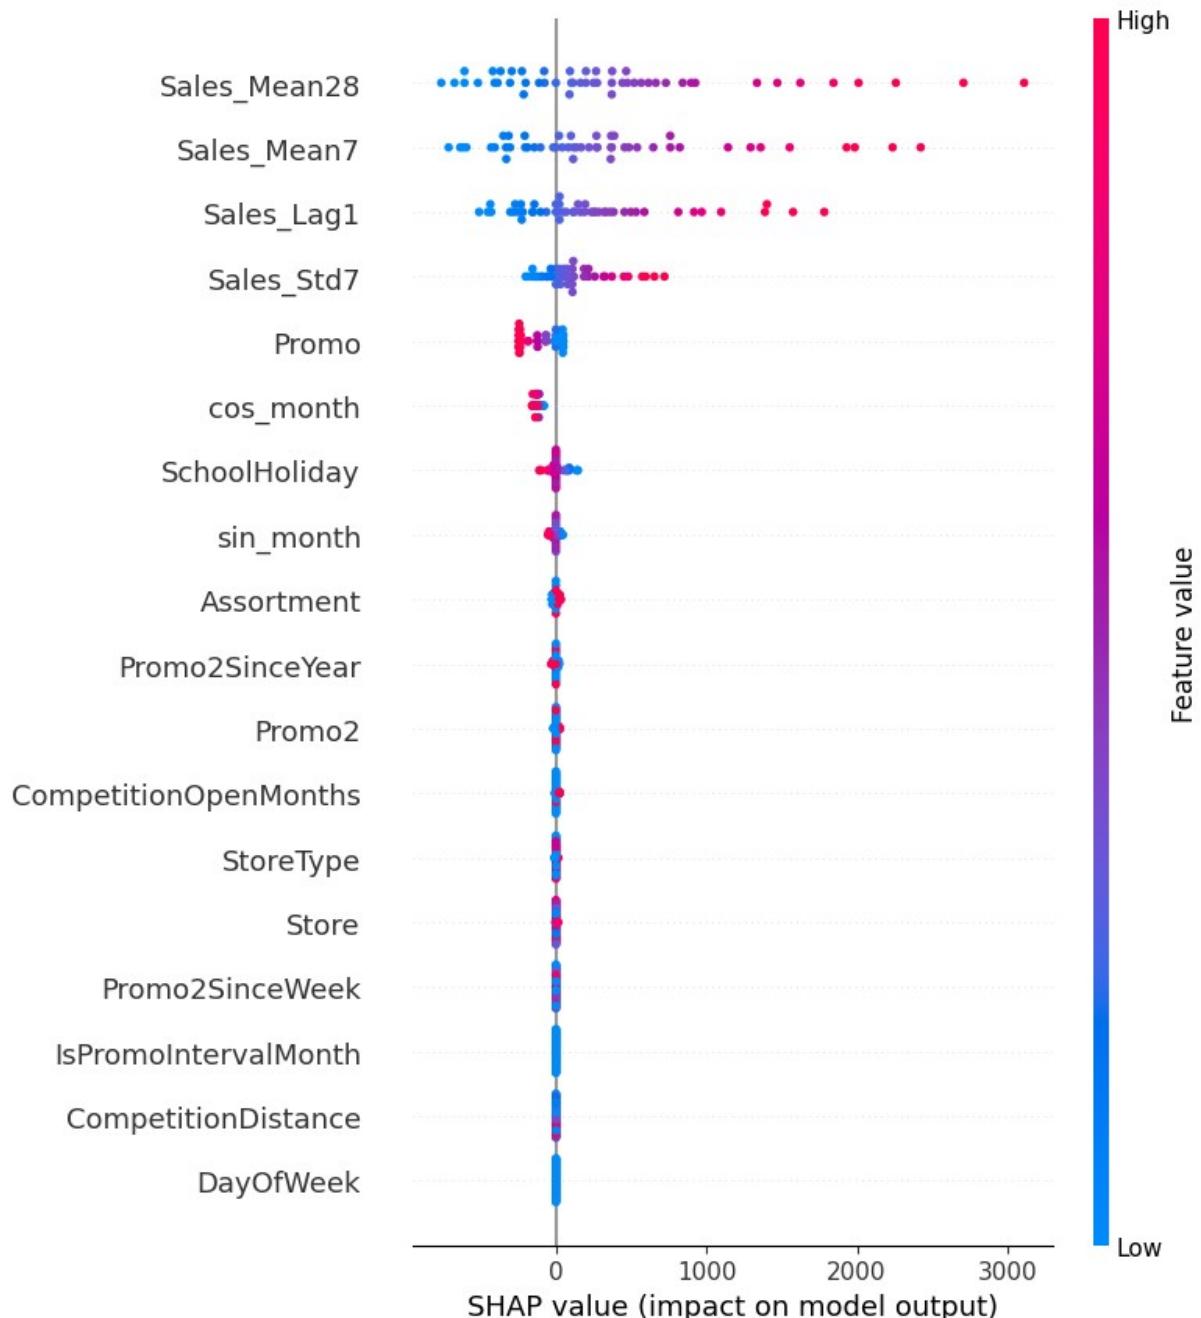
Saved SHAP: /kaggle/working/models_phase5/shap_summary.png

import numpy as np, pandas as pd

sv = np.array(shap_values)

mean_abs = np.mean(np.abs(sv), axis=0)
fi = pd.Series(mean_abs, index=FEATURES).sort_values(ascending=False)
print(fi.head(20))
from IPython.display import Image
Image(filename='/kaggle/working/models_phase5/shap_summary.png')
```

```
Sales_Mean28          652.893720
Sales_Mean7           559.414780
Sales_Lag1            395.885431
Sales_Std7             169.381155
Promo                 141.703138
cos_month              126.541633
SchoolHoliday          34.717644
sin_month              21.886080
Assortment              18.100117
Promo2SinceYear         11.516077
Promo2                  7.572928
CompetitionOpenMonths    6.602969
StoreType                1.989237
Store                   1.234588
DayOfWeek               0.000000
CompetitionDistance      0.000000
Promo2SinceWeek          0.000000
IsPromoIntervalMonth     0.000000
dtype: float64
```



Find best hyper parameters for our lstm,Tcm models

```
def optuna_objective(trial):
    units = trial.suggest_categorical('units', [64,128])
    dropout = trial.suggest_float('dropout', 0.1, 0.3)
    lr = trial.suggest_loguniform('lr', 1e-4, 1e-3)
    m = build_lstm_model(SEQ_LEN, n_features, HORIZON,
n_quantiles=len(QUANTILES), units=units, dropout=dropout, lr=lr)
    m.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=lr),
```

```

loss=multi_quantile_loss_wrapper(QUANTILES))
    history = m.fit(X_train_s, y_train_s, validation_data=(X_val_s,
y_val_s), batch_size=512, epochs=6, verbose=0)
    val_loss = min(history.history['val_loss'])
    K.clear_session()
    return val_loss

study = optuna.create_study(direction='minimize',
sampler=optuna.samplers.TPESampler(seed=RANDOM_SEED))
study.optimize(optuna_objective, n_trials=6)
print("Optuna best:", study.best_trial.params)
import joblib
joblib.dump(study, os.path.join(MODEL_DIR, "optuna_study.pkl"))

[I 2026-02-21 17:43:24,340] A new study created in memory with name:
no-name-230e3072-f6f3-4998-b55a-08ddef9fbaf
/tmp/ipykernel_55/574459363.py:4: FutureWarning: suggest_loguniform
has been deprecated in v3.0.0. This feature will be removed in v6.0.0.
See https://github.com/optuna/optuna/releases/tag/v3.0.0. Use
suggest_float(..., log=True) instead.
    lr = trial.suggest_loguniform('lr', 1e-4, 1e-3)
[I 2026-02-21 17:45:58,804] Trial 0 finished with value:
0.010437797755002975 and parameters: {'units': 128, 'dropout':
0.24639878836228102, 'lr': 0.00039687933304443713}. Best is trial 0
with value: 0.010437797755002975.
[I 2026-02-21 17:47:51,472] Trial 1 finished with value:
0.010964587330818176 and parameters: {'units': 64, 'dropout':
0.1116167224336399, 'lr': 0.0007348118405270452}. Best is trial 0
with value: 0.010437797755002975.
[I 2026-02-21 17:50:24,822] Trial 2 finished with value:
0.010731698013842106 and parameters: {'units': 128, 'dropout':
0.10411689885916049, 'lr': 0.0009330606024425672}. Best is trial 0
with value: 0.010437797755002975.
[I 2026-02-21 17:52:17,389] Trial 3 finished with value:
0.012731725350022316 and parameters: {'units': 64, 'dropout':
0.1363649934414201, 'lr': 0.00015254729458052615}. Best is trial 0
with value: 0.010437797755002975.
[I 2026-02-21 17:54:50,372] Trial 4 finished with value:
0.011255506426095963 and parameters: {'units': 128, 'dropout':
0.18638900372842315, 'lr': 0.00019553708662745247}. Best is trial 0
with value: 0.010437797755002975.
[I 2026-02-21 17:56:43,044] Trial 5 finished with value:
0.011262130923569202 and parameters: {'units': 64, 'dropout':
0.15842892970704364, 'lr': 0.0002324672848950435}. Best is trial 0
with value: 0.010437797755002975.

Optuna best: {'units': 128, 'dropout': 0.24639878836228102, 'lr':
0.00039687933304443713}

['/kaggle/working/models_phase5/optuna_study.pkl']

```

**Run this if you want to download a zip of models**

```
import os
import zipfile

directory_to_zip = '/kaggle/working/models_phase5'
output_filename = 'rossman_project_backup.zip'

with zipfile.ZipFile(output_filename, 'w', zipfile.ZIP_DEFLATED) as
zip_ref:
    for root, dirs, files in os.walk(directory_to_zip):
        for file in files:
            file_path = os.path.join(root, file)
            zip_ref.write(file_path, os.path.relpath(file_path,
directory_to_zip))

print(f"Successfully created {output_filename}!")
```

Successfully created rossman\_project\_backup.zip!